

UNIVERSITY OF SALERNO

DEPARTMENT OF INFORMATION ENGINEERING AND ELECTRICAL AND APPLIED MATHEMATICS



Master's Degree in Computer Engineering

Project report

Parallel version of the Tarjan's and Kosaraju's algorithm to find strongly connected components in a graph

Group members:

Amato Emilio	0622701903	e.amato16@studenti.unisa.it
Bove Antonio	0622701898	a.bove57@studenti.unisa.it
De Gregorio Marco	0622701865	m.deggregorio19@studenti.unisa.it



ACADEMIC YEAR 2022/2023

Indice

1 Problem Description	5
1.1 Introduction to the problem and its solution	5
1.2 Execution Example	6
2 Experimental Setup	11
2.1 Hardware	11
2.1.1 CPU	11
2.1.2 RAM	11
2.2 Software	11
3 Performance, Speed-Up & Efficiency	12
3.1 OpenMP & MPI	12
3.2 Graph Density	13
3.3 Analisys	13
3.4 Meaning of the graphs representing the analyzes carried out	14
4 Tarjan's Algorithm	15
4.1 Case Study 1	16
4.1.1 500 vertices - 25% of density	16
4.1.2 500 vertices - 50% of density	17
4.1.3 500 vertices - 75% of density	18
4.2 Case Study 2	19
4.2.1 1000 vertices - 25% of density	19
4.2.2 1000 vertices - 50% of density	20
4.2.3 1000 vertices - 75% of density	21
4.2.4 1000 vertices - Special Cases - 0% of density	22
4.2.5 1000 vertices - Special Cases - 100% of density	23
4.3 Case Study 3	24
4.3.1 2500 vertices - 25% of density	24
4.3.2 2500 vertices - 50% of density	25

4.3.3	2500 vertices - 75% of density	26
4.4	Case Study 4	27
4.4.1	10000 vertices - 25% of density	27
4.4.2	10000 vertices - 50% of density	28
4.4.3	10000 vertices - 75% of density	29
4.4.4	10000 vertices - Special Cases - 0% of density	30
4.4.5	10000 vertices - Special Cases - 100% of density	31
5	Kosaraju's algorithm	32
5.1	Kosaraju and Tarjan algorithm	32
5.2	Kosaraju's parallel version analysis	32
5.3	Case Study 1	33
5.3.1	500 vertices - 25% of density	33
5.3.2	500 vertices - 50% of density	34
5.3.3	500 vertices - 75% of density	35
5.4	Case Study 2	36
5.4.1	1000 vertices - 25% of density	36
5.4.2	1000 vertices - 50% of density	37
5.4.3	1000 vertices - 75% of density	38
5.4.4	1000 vertices - Special Cases - 0%	39
5.4.5	1000 vertices - Special Cases - 100%	40
5.5	Case Study 3	41
5.5.1	2500 vertices - 25% of density	41
5.5.2	2500 vertices - 50% of density	42
5.5.3	2500 vertices - 75% of density	43
5.6	Case Study 4	44
5.6.1	10000 vertices - 25% of density	44
5.6.2	10000 vertices - 50% of density	45
5.6.3	10000 vertices - 75% of density	46
5.6.4	10000 vertices - Special Cases - 0%	47

5.6.5	10000 vertices - Special Cases - 100%	48
6	Cluster	49
6.1	Tarjan	50
6.1.1	500 vertices - 25% of density	50
6.1.2	500 vertices - 50% of density	51
6.1.3	500 vertices - 75% of density	52
6.1.4	1000 vertices - 25% of density	53
6.1.5	1000 vertices - 50% of density	54
6.1.6	1000 vertices - 75% of density	55
6.1.7	2500 vertices - 25% of density	56
6.1.8	2500 vertices - 50% of density	57
6.1.9	2500 vertices - 75% of density	58
6.2	Kosaraju	59
6.2.1	500 vertices - 25% of density	59
6.2.2	500 vertices - 50% of density	60
6.2.3	500 vertices - 75% of density	61
6.2.4	1000 vertices - 25% of density	62
6.2.5	1000 vertices - 50% of density	63
6.2.6	1000 vertices - 75% of density	64
6.2.7	2500 vertices - 25% of density	65
6.2.8	2500 vertices - 50% of density	66
6.2.9	2500 vertices - 75% of density	67
7	Final Considerations	68
7.1	Tarjan Implementation	68
7.2	Kosaraju Implementation	69
7.3	Cluster execution	69
8	Test Case	69
9	Api Documentation	71

9.1	include/tarjan.h File Reference	71
9.2	include/kosaraju.h File Reference	72
9.3	include/Utils.h File Reference	74

10 How To Run

79

1 Problem Description

Provide a parallel version of the Tarjan's algorithm to find Strongly Connected Components in a graph. The implementation MUST use an hybrid message passing / shared memory paradigm and has to be implemented by using MPI and OpenMP. Students MUST provide parallel processes on different nodes, and each process has to be parallelized by using OpenMP (i.e.: MPI will spawn OPENMP-compiled processes). Students can choose the graph allocation method they prefer. They can eventually produce the graph directly in distributed memory (without store anything). Good graph dimensions are greater than 4GB of data.

Optional: use the same parallelization method to parallelize one of the other proposed algorithm in the Wikipedia page. Analyze differences.

1.1 Introduction to the problem and its solution

Tarjan's and Kosaraju's, the analyzed algorithm in our solution, are two graph analysis algorithms used to find strongly connected components in an oriented graph. Tarjan's algorithm was first proposed by Robert Tarjan in 1972 and has been used in a variety of applications, such as creating topological maps and analyzing social networks. Kosaraju's algorithm, proposed by Sharir and Kosaraju in 1978, uses a dfs technique similar to Tarjan's algorithm and, like the latter, finds strongly connected components in an oriented graph. Parallelization of Tarjan's and Kosaraju's algorithms is an important problem because both are sequential algorithms and thus can be slow for large graphs. Parallelization of the algorithms allows multiple processors to be used simultaneously to run the algorithm, thus improving performance. There are several ways to parallelize Tarjan and Kosaraju algorithms, including data-based parallelization and task-based parallelization. Data-based parallelization consists of executing the algorithm on different portions of the data in parallel, while task-based parallelization consists of executing different stages of the algorithm in parallel. The following project report will examine in detail the data-based parallelization approach of the Tarjan and Kosaraju algorithms and evaluate their performance in terms of speed and scalability. In particular, we will explore the data-driven approach using OpenMP and MPI. In addition, we will compare the results obtained with those of the original Tarjan and Kosaraju algorithms to evaluate the effectiveness of parallelization. In summary, the goal of this project is to find the best strategy for parallelization of the Tarjan and Kosaraju algorithms and to evaluate the performance of the different approaches through a series of tests on large graphs. We believe that the results obtained in this project can help to develop more efficient and scalable Tarjan and Kosaraju algorithms for the analysis of large-dimensional oriented graphs.

Following an analysis carried out in order to understand what might be a possible solution to the problem under examination, the following solution idea has been reached. Assuming the use of 4 MPI processes, the first step involves reading the entire initial graph by these. At this point, each process, based on its rank, will calculate its own competence interval that we will indicate with $[start; stop]$; in particular, a mechanism of graph division is used in order to distribute the load equally among all processes. Subsequently, each process will perform a specific algorithm for calculating the strongly connected components on its own portion of the graph, storing them in a dynamic array of arrays. At this point, a reduction process begins, starting from the 4 initial processes that will lead to having a single final process that will have the strongly connected components of the initial graph. In particular, a cycle will begin, whose exit condition is given by the fact that only one process remains, in which:

- if the process has an even rank and the process with $rank + 1$ exists, then the latter will send it, in addition to the strongly connected components found, a series of information in order to allow the process with $rank + 1$ to build a new graph on which to continue execution;
- if the process has an odd rank, on the other hand, it will receive the aforementioned information, store it in appropriate data structures, take care of creating a new graph that will replace the previous one and reapply the algorithm for calculating strongly connected components on this, replacing the previous array.

Subsequently, the even-rank processes that had sent their information, following a split of the communicator, are killed and there will be a new assignment of ranks to the remaining processes (in the case under examination, the processes with rank 0 and 2 after communicating with the processes with rank 1 and 3 are killed and the process with rank 1 is assigned rank 0 and the process with rank 3 is assigned rank 1). The cycle begins again until only one process remains, which will be the owner of the total strongly connected components.

1.2 Execution Example

Let's now provide an example of the execution of the previously proposed solution, illustrating all the steps. In particular, we consider having an MPI process number of 4 ($size = 4$) and the following graph made up of 10 nodes:

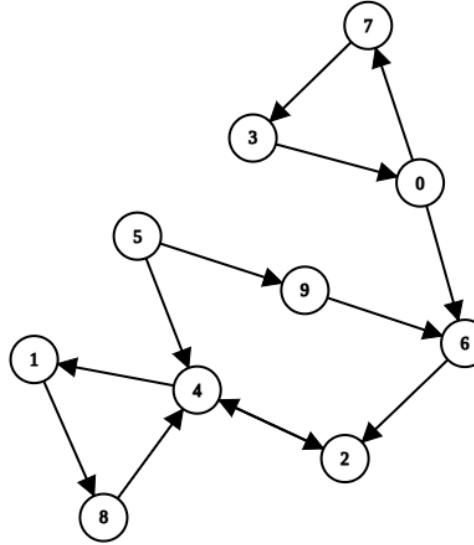
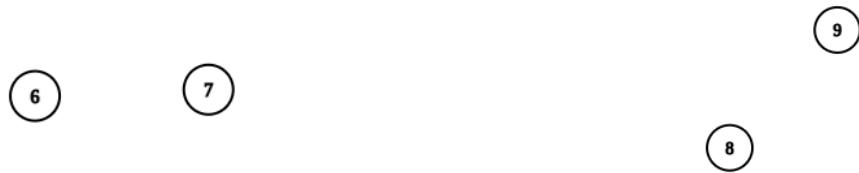


Figura 1: Initial graph

At this point, as previously stated, each process will calculate its own competence range $[start; stop]$ and will execute the algorithm to calculate the strongly connected components, storing them in a dynamic array of arrays that we will call *sccs*.





Process 2

start → 6

stop → 8

sccs: $\begin{bmatrix} [0] \rightarrow [6] \\ [1] \rightarrow [7] \end{bmatrix}$

Process 3

start → 8

stop → 10

sccs: $\begin{bmatrix} [0] \rightarrow [8] \\ [1] \rightarrow [9] \end{bmatrix}$

Subsequently, each process instantiates two hash tables:

- *sccsHT*: which has as key the name of a macronode and as value the array of nodes that belong to it;
- *auxiliaryGraphHT*: which has as key the name of a node and as value the macronode it belongs to.

At this point, the first iteration of the reduction process begins and thus *size* = 4 and *numIteration* = 1. In particular, process 0 sends to process 1 only the strongly connected components found (*sccs*), as initially the hash tables will be empty. Process 1, as a result, will concatenate its *sccs* array with the one received from process 0, populate its hash tables, and create a temporary array (*nodes*) in which it will insert the nodes of the new graph that needs to be constructed, in the following way:

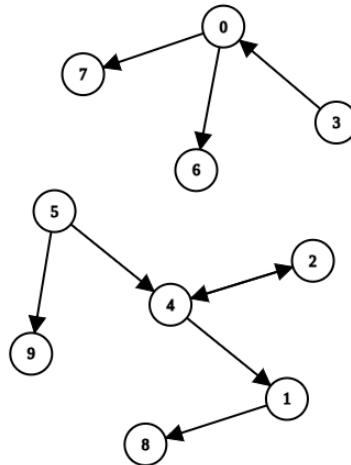
- for each element of the *sccs* array that represents a macronode, it is assigned a new name and added to the temporary array *nodes*; in addition, the hash tables *sccsHT* and *auxiliaryGraphHT* are also populated
- for all other elements, on the other hand, the node in question is only added to the temporary array *nodes*

Therefore, we will have:

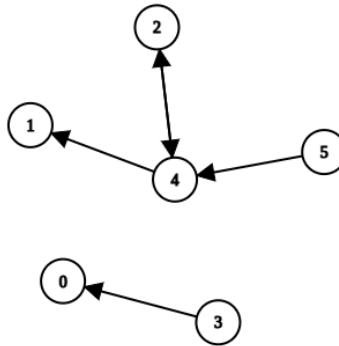
Process 0 sends to process 1

sccs: $\begin{bmatrix} [0] \rightarrow [0] \\ [1] \rightarrow [1] \\ [2] \rightarrow [2] \\ [3] \rightarrow [3] \\ [4] \rightarrow [4] \\ [5] \rightarrow [5] \end{bmatrix}$	sccsHT: empty auxiliaryGraphHT: empty nodes: [0 1 2 3 4 5]
---	---

At this point, based on the previous information, process 1 will construct a new graph, obtaining:



and, as a result, will calculate the strongly connected components this time considering *start* as the value 0 and *stop* as the number of nodes in the graph. So, again considering only the nodes that are part of the graph, process 1 will calculate the strongly connected components on the following graph:



overwriting *sccs* in the following way:

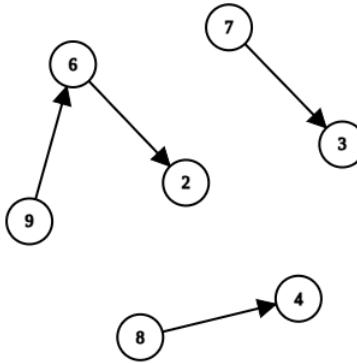
$$\text{sccs: } \begin{bmatrix} [0] \rightarrow [0] \\ [1] \rightarrow [3] \\ [2] \rightarrow [1] \\ [3] \rightarrow [5] \\ [4] \rightarrow [2, 4] \end{bmatrix}$$

What has been said for processes 0 and 1 also applies to processes 2 and 3. Therefore:

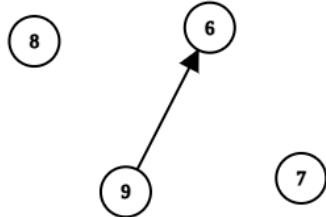
Process 2 sends to process 3

$$\text{sccs: } \begin{bmatrix} [0] \rightarrow [6] \\ [1] \rightarrow [7] \\ [2] \rightarrow [8] \\ [3] \rightarrow [9] \end{bmatrix} \quad \text{sccsHT: empty} \quad \text{auxiliaryGraphHT: empty} \quad \text{nodes: [6 7 8 9]}$$

At this point, based on the previous information, process 3 will build a new graph, obtaining:



and, as a result, it will calculate the strongly connected components, this time considering 0 as the value of *start* and the number of nodes in the graph as the value of *stop*. Therefore, taking into account only the nodes that are part of the graph, process 3 will calculate the strongly connected components on the following graph:



overwriting *sccs* in the following way:

$$\text{sccs: } \begin{bmatrix} [0] \rightarrow [7] \\ [1] \rightarrow [8] \\ [2] \rightarrow [9] \\ [3] \rightarrow [6] \end{bmatrix}$$

At this point, processes 0 and 2 are killed and processes 3 and 4, following a split operation of the initial communicator, respectively obtain ranks 0 and 1:

P1 becomes P0
P3 becomes P1

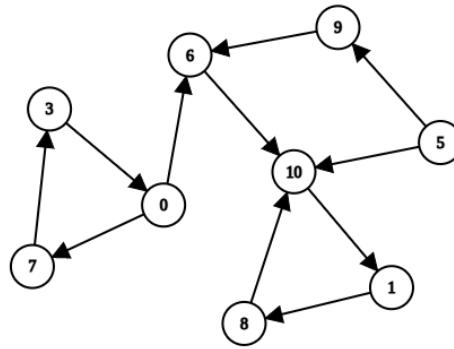
So, the first iteration ends and we have *size* = 2 and *numIteration* = 2. Since *size* is still greater than 1, the previously described process is repeated again.

Process 0 sends to process 1

$$\begin{aligned}
 \text{sccs: } & \begin{bmatrix} [0] \rightarrow [0] \\ [1] \rightarrow [3] \\ [2] \rightarrow [1] \\ [3] \rightarrow [5] \\ [4] \rightarrow [2 4] \\ [5] \rightarrow [7] \\ [6] \rightarrow [8] \\ [7] \rightarrow [9] \\ [8] \rightarrow [6] \end{bmatrix} & \text{sccsHT: } 10 \rightarrow [2 4] & \text{auxiliaryGraphHT: } \begin{bmatrix} 2 \rightarrow 10 \\ 4 \rightarrow 10 \end{bmatrix} & \text{nodes: } [0 3 1 5 10 7 8 9 6]
 \end{aligned}$$

We note that in this case, since there is an array of size greater than 1 in *sccs*, a macronode is created and assigned an ID (greater than the number of nodes in the graph minus one, to avoid having nodes with the same name), for example 10.

Starting from the previous information, process 1 will construct a new graph, obtaining:



and, as a result, will calculate the strongly connected components by again examining only the nodes that are part of the graph, overwriting *sccs* in the following way:

$$\text{sccs: } \begin{bmatrix} [0] \rightarrow [3 0 7] \\ [1] \rightarrow [10 1 8] \\ [2] \rightarrow [5] \\ [3] \rightarrow [6] \\ [4] \rightarrow [9] \end{bmatrix}$$

At this point, process 0 is killed and process 1, following the split of the communicator again, obtains rank 0:

P1 becomes P0

The second iteration ends this way and we have *size* = 1 and *numIteration* = 3. Since *size* = 1 the loop ends and the process with rank 0 (the last remaining one) has in *sccs* the strongly connected components. In particular, by expanding the macronodes through the *sccsHT*, we find that the final strongly connected components are 5:

$$\begin{bmatrix} SCC\ 1 \rightarrow [3 0 7] \\ SCC\ 2 \rightarrow [2 4 1 8] \\ SCC\ 3 \rightarrow [5] \\ SCC\ 4 \rightarrow [6] \\ SCC\ 5 \rightarrow [9] \end{bmatrix}$$

2 Experimental Setup

2.1 Hardware

2.1.1 CPU

Model Name: AMD Ryzen 9 3900X
Socket: AMD Socket AM4
Foundry: TSMC
Process Size: 7 nm
Transistors: 7,600 million
Die Size: 2x 74 mm²
I/O Process Size: 12 nm
I/O Die Size: 124 mm²
Package: µOPGA-1331
tCaseMax: 95°C
Frequency: 3.8 GHz
Turbo Clock: up to 4.6 GHz
Base Clock: 100 MHz
Multiplier: 38.0x
Multiplier Unlocked: Yes
TDP: 105 W
FP32: 2,649.6 GFLOPS
Number of Cores: 12
Number of Threads: 24
SMP # CPUs: 1
Integrated Graphics: N/A
Cache L1: 64K (per core)
Cache L2: 512K (per core)
Cache L3: 64MB
Memory Support: DDR4-3200 MHz
Dual-channel
ECC Memory: No
PCI-Express: Gen 4, 24 Lanes (CPU only)
Generation: Ryzen 9 (Zen 2 (Matisse))

2.1.2 RAM

Memory Series: VENGEANCE LPX
SPD Latency: CL16 (16-18-18-36)
SPD Speed: 2666MHz
SPD Voltage: 1.2V
Speed Rating: PC4-25600 (3200MHz)
Heat Spreader: Anodized Aluminum
Package Memory Format: DIMM senza buffer
Performance Profile: XMP 2.0
Package Memory Pin: 288
Densità: 16GB (2 x 8GB)
Tipo: DRAM kit di memoria
Tecnologia: DDR4 SDRAM
Data Integrity Check: Non ECC

2.2 Software

- Ubuntu 22.10
- GCC 11.3.0
- Visual Studio 1.74.2
- CMake 3.22.1
- Python 3.10.6

3 Performance, Speed-Up & Efficiency

In the field of high performance computing, the relationship between speed-up and efficiency is a topic of great importance. Speed-up is a measure of the increase in speed obtained by using multiple processors compared to a single processor, while efficiency is a measure of the amount of work actually performed by the processors compared to the maximum theoretical amount of work.

Speed-up is typically expressed as the ratio of the execution time on a single processor to the execution time on multiple processors:

$$S_n = \frac{T_1}{T_p(n)}$$

where T_1 is the time taken by the best known serial algorithm, while $T_p(n)$ is the execution time of the parallel algorithm on n processors.

For example, if the execution time of a certain application is 100 seconds on a single processor and 50 seconds on 4 processors, the speed-up is 2. In theory, the use of multiple processors should allow for linear speed-up relative to the number of processors used. However, in practice this is not always possible due to issues of communication and synchronization between processors.

Efficiency, particularly relative efficiency, is given by the ratio of the time to execute an algorithm on a single processor to n multiplied by the time to execute a parallel algorithm on n processors:

$$E_r = \frac{T_1}{n \cdot T_n}$$

Using the previous example, relative efficiency would be $100/(4 \cdot 50)$ on 4 processors. In theory, efficiency should always be greater than or equal to 1, as more processors are used to perform the work. However, in practice there may be losses in efficiency due to issues of communication and synchronization between processors.

3.1 OpenMP & MPI

MPI, or Message Passing Interface, is a widely-used library for parallel computing on distributed memory systems. It provides a set of functions for inter-process communication, allowing multiple processes to work together and solve a problem faster by dividing the workload among them. MPI is a standard, meaning that it is not tied to any particular system or architecture. This makes it portable and allows it to be used on a wide range of computers, from small clusters to large supercomputers. One of the main advantages of MPI is that it allows for a high degree of flexibility in designing parallel algorithms. The library provides a variety of communication operations, including point-to-point communication (sending messages between two processes), collective communication (sending messages between all processes), and non-blocking communication (allowing processes to continue executing while waiting for a message to arrive). Despite its advantages, MPI also has some limitations. One limitation is that it does not provide any support for shared memory programming, which can be useful for certain types of algorithms.

On the other hand, OpenMP is a useful solution for parallelizing existing code that has been written to run on a single core. However, it is not always possible to gain benefits from using OpenMP. There are some reasons why OpenMP might not be suitable for a certain code or why significant benefits may not be gained from parallelization. One of the reasons why OpenMP might not be suitable for a certain code is that the code might already be optimized for execution on a single core. In this case, parallelizing the code may not lead to a significant increase in performance. Furthermore, OpenMP is mainly designed to parallelize code that is heavily dependent on loops, such as the code used to solve scientific computing problems. If the code is not heavily dependent on loops, it may not be possible to gain benefits from parallelization. Another reason why OpenMP might not be suitable for a certain code is that it might require a high level of communication between cores. If the code requires a high level of communication between cores, it might be necessary to use another parallel library, such as MPI, which is designed to handle a high level of communication between cores.

3.2 Graph Density

The density of a graph is a measure of how "populated" the graph is with edges compared to the maximum possible number of edges. In other words, it indicates how "complete" the graph is in relation to its size.

The density of a graph is calculated as the ratio of the number of edges present in the graph to the maximum number of edges possible for that graph. In the case of an undirected graph with n vertices, the maximum number of possible edges is:

$$\frac{n(n - 1)}{2}$$

In the case of a directed graph with n vertices, the maximum number of possible edges is $n(n - 1)$.

The density of a graph can vary from 0 to 1. A graph with density 0 is an empty graph, meaning a graph without edges. A graph with density 1 is a complete graph, meaning a graph in which every vertex is connected to all other vertices.

Graphs with densities greater than 0.5 are considered dense graphs. Graphs with densities less than 0.5 are considered sparse graphs.

The density of a graph can be used to determine the properties of the graph itself and to compare graphs of different sizes. For example, dense graphs are often more similar to each other than sparse graphs, regardless of their sizes. Furthermore, the density of a graph can be used to determine the complexity of algorithms that operate on that graph. For example, algorithms that explore a complete graph have a higher complexity than algorithms that explore a sparse graph.

3.3 Analisys

For the analysis of Performance, Speed-up and Efficiency, the measurements that have been conducted are based on the following criterion: given the sequential version of the assigned algorithm and given its parallel implementation, for all three levels of optimization each of the latter was tested, for each optimization level set, on three types of graph. Given a graph with a certain number of vertices, three possible scenarios were considered: one where the graph is sparse, one where it is moderately dense, and one where it is very dense. In addition, two special cases were also considered, one where the graph has a density of 0 and one where it has a density of 1.

For each graph tested in the different optimization levels set, the variability between the different versions of the parallelized algorithm was managed by spreading the number of MPI processes and the number of OpenMP threads with different combinations.

The comparison between the two versions was carried out, once the number of measurements for each version was fixed, on the following input graphs and, for each of it, on how many levels of density the graph has, represented by min number - max number of edges per vertex:

- graph with a number of vertices equal to 500:
 1. **low density 25%**: 0 - 250
 2. **medium density 50%**: 250 - 500
 3. **high density 75%**: 500 - 750
- graph with a number of vertices equal to 1000:
 1. **low density 25%**: 0 - 250
 2. **medium density 50%**: 250 - 500
 3. **high density 75%**: 500 - 750

- graph with a number of vertices equal to 10000:
 1. **low density 25%**: 0 - 2500
 2. **medium density 50%**: 2500 - 5000
 3. **high density 75%**: 5000 - 7500
- special case graph with a number of vertices equal to 1000
 1. **zero density 0%**: 0 - 0
 2. **maximum density 100%**: 999 - 999
- special case graph with a number of vertices equal to 10000
 1. **zero density 0%**: 0 - 0
 2. **maximum density 100%**: 9999 - 9999

For each density level of each graph on which the execution of the algorithm version has been defined, as regards the management of MPI processes and OpenMP threads, the latter has been structured by assigning the following resources to each of prefixed density level:

- 0 MPI process - No OpenMP threads (*Sequential version*)
- 1 MPI process - {1, 2, 4, 8, 16} OpenMP threads (*Parallel version*)
- 2 MPI processes - {1, 2, 4, 8, 16} OpenMP threads (*Parallel version*)
- 4 MPI processes - {1, 2, 4, 8, 16} OpenMP threads (*Parallel version*)
- 8 MPI processes - {1, 2, 4, 8, 16} OpenMP threads (*Parallel version*)

Therefore, from the considerations made, we can define the case studies analysed:

1. **Case Study 1**: the sequential program and the parallel program are compiled with a graph of 500 vertices.
2. **Case Study 2**: the sequential program and the parallel program are compiled with a graph of 1000 vertices.
3. **Case Study 3**: the sequential program and the parallel program are compiled with a graph of 2500 vertices.
4. **Case Study 4**: the sequential program and the parallel program are compiled with a graph of 10000 vertices.

3.4 Meaning of the graphs representing the analyzes carried out

x and *y axis* represents the number of MPI processes with which the evaluation of the speedup was computed.

In particular, the meaning of the colors are the following ones:

- **red** representes the evaulation of Speedup with 1 OpenMP thread.
- **green** representes the evaulation of Speedup with 2 OpenMP threads.
- **blue** representes the evaulation of Speedup with 4 OpenMP threads.
- **yellow** representes the evaulation of Speedup with 8 OpenMP threads.
- **cyan** representes the evaulation of Speedup with 16 OpenMP threads.

4 Tarjan's Algorithm

The Tarjan algorithm is a well-known algorithm used for finding strongly connected components in a graph. It was first proposed by Robert Tarjan in 1972 and since then it has become a widely used method in graph theory and computer science. The algorithm is based on depth-first search (DFS) and uses a shared stack to keep track of visited vertices. Now, we will discuss the Tarjan algorithm, including its working principle and time complexity.

The algorithm proceeds as follows:

1. initialize an empty stack and an array of flags to mark visited vertices
2. for each unvisited vertex in the graph, perform a DFS starting from that vertex
3. during the DFS, for each vertex v:
 - if the vertex has not been visited yet, mark it as visited and push it onto the stack
 - if the vertex has already been visited and it is still on the stack, then it is a vertex of a strongly connected component. In this case, keep removing vertices from the top of the stack until v is encountered. The removed vertices form a strongly connected component

The Tarjan algorithm runs in time $O(V + E)$ where V is the number of vertices and E the number of edges in the graph. The time complexity of the algorithm is determined by the number of vertices and edges in the graph, and it is considered very efficient. Furthermore, it's widely used in different applications like compilers, operating systems, and databases, and it's an algorithm that still being used and studied today.

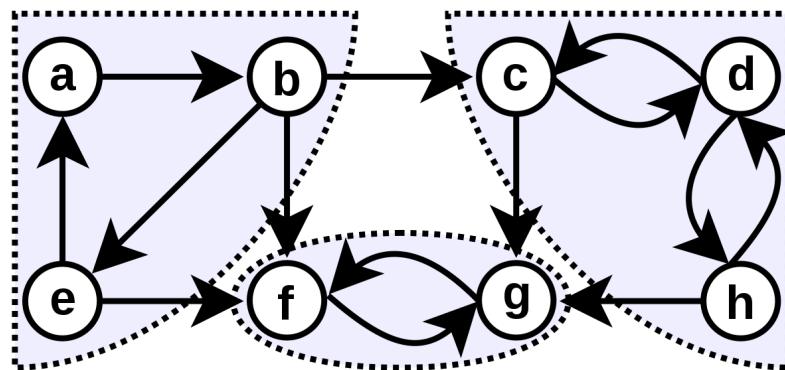


Figura 2: Strongly connected components on a graph

We will then show the results obtained by implementing the solution described in paragraphs 1.1 and 1.2 and using Tarjan as the algorithm for calculating the strongly connected components.

4.1 Case Study 1

4.1.1 500 vertices - 25% of density

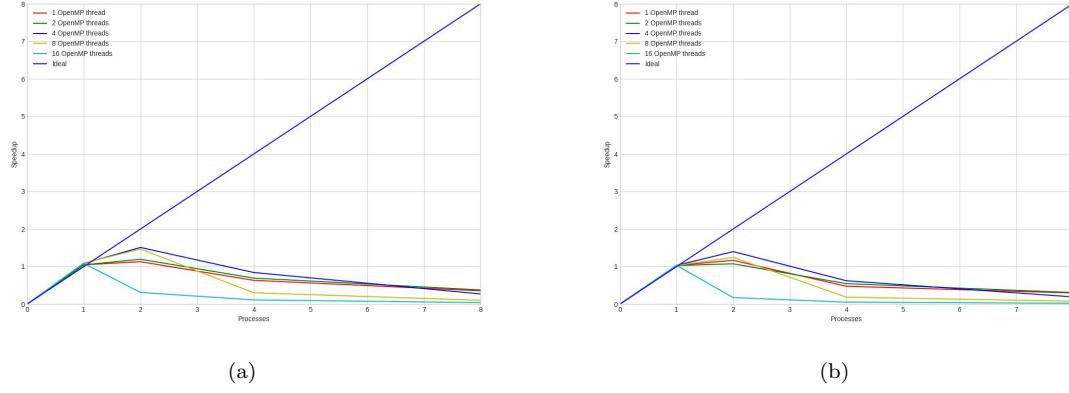


Figura 3: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

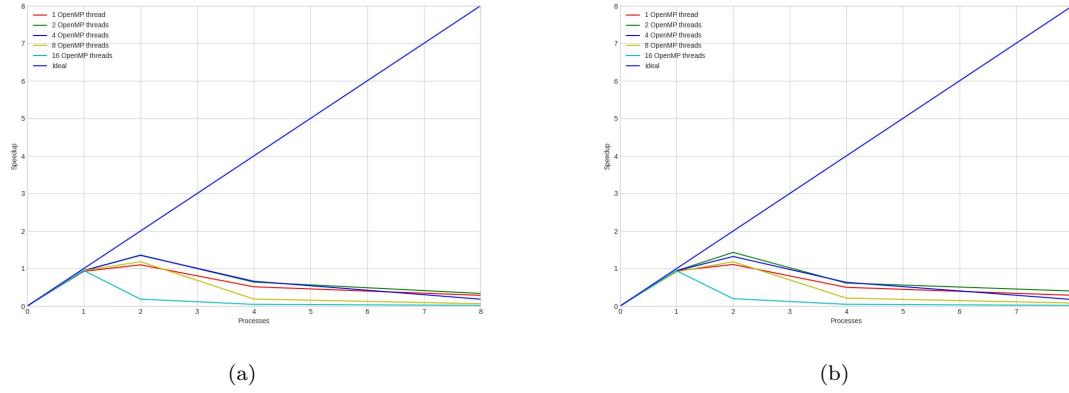


Figura 4: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Tarjan	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.002801	0.0	0.002801	1.0	1.0
Parallel	1	1	0.002678000000000003	0.0	0.002678000000000003	1.0459297983569829	1.0459297983569829
Parallel	1	2	0.002678000000000003	0.0	0.002678000000000003	1.0459297983569829	1.0459297983569829
Parallel	1	4	0.002582	0.0	0.002582	1.084817970565453	1.084817970565453
Parallel	1	8	0.0025568	0.0	0.0025568	1.0955100125156445	1.0955100125156445
Parallel	1	16	0.002582	0.0	0.002582	1.084817970565453	1.084817970565453
Parallel	2	1	0.000924	0.001556399999999999	0.002480400000000002	1.1292533462344783	0.5646266731172391
Parallel	2	2	0.000823800000000001	0.0015204	0.0023444	1.1947619860092136	0.5973809930046068
Parallel	2	4	0.0009254	0.000931000000000001	0.0018564	1.50883430295195	0.754417151475975
Parallel	2	8	0.000822	0.0010866	0.0019088	1.4674140821458508	0.7337070410729254
Parallel	2	16	0.000826599999999999	0.0081986	0.0090248	0.3103669887421328	0.1551834943710664
Parallel	4	1	0.0004456	0.003982	0.004428	0.632565494232159	0.1581413730803975
Parallel	4	2	0.0005086	0.0035364	0.00404479999999994	0.6924940664556963	0.17312351661392408
Parallel	4	4	0.0005088	0.00281879999999998	0.0033276	0.84174780622671	0.2104369515566775
Parallel	4	8	0.0005154	0.0087472	0.0092626	0.3023988944788721	0.07559972361971802
Parallel	4	16	0.000482799999999997	0.0246219999999998	0.0251048	0.11157228896466015	0.027893072241165037
Parallel	8	1	0.000399000000000005	0.0075142	0.0079134	0.35395657997826474	0.04424457249728309
Parallel	8	2	0.000467200000000001	0.006951400000000001	0.007418800000000001	0.37755432145360435	0.04719429018170054
Parallel	8	4	0.000452400000000005	0.009898	0.010350400000000001	0.2706175606739836	0.03382719508424795
Parallel	8	8	0.0004408	0.0272402	0.027681	0.10118854087641342	0.012648567609551678
Parallel	8	16	0.0004058	0.071011	0.07141700000000001	0.03922035369729896	0.00490254421216237

Tabella 1: Measurement with O0 optimization

4.1.2 500 vertices - 50% of density

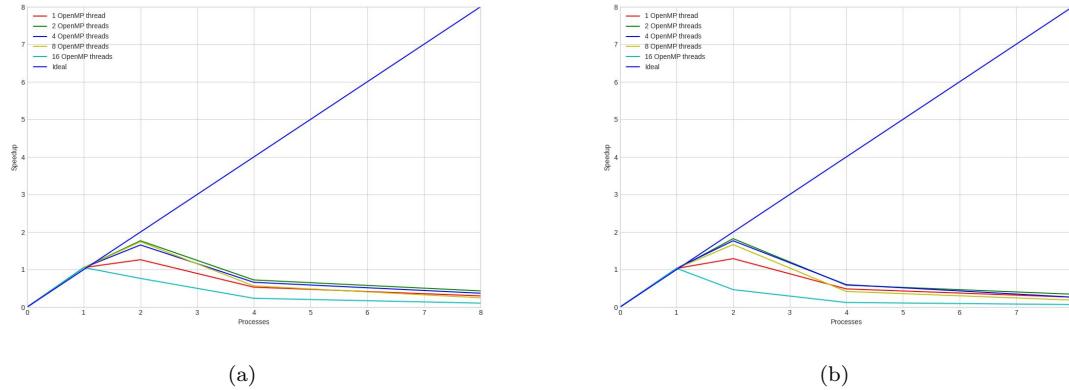


Figura 5: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

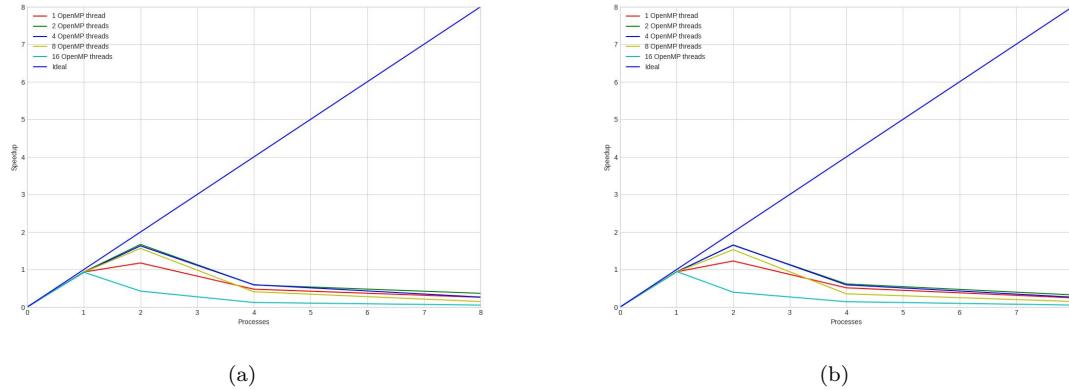


Figura 6: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Tarjan	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.007753	0.0	0.007753	1.0	1.0
Parallel	1	1	0.0073644	0.0	0.0073644	1.0527673673347455	1.0527673673347455
Parallel	1	2	0.0073706	0.0	0.0073706	1.051881800667517	1.051881800667517
Parallel	1	4	0.00737739999999999	0.0	0.00737739999999999	1.050912245565472	1.050912245565472
Parallel	1	8	0.0073690000000000005	0.0	0.0073690000000000005	1.0521101913421087	1.0521101913421087
Parallel	1	16	0.0073606	0.0	0.0073606	1.0533108713963535	1.0533108713963535
Parallel	2	1	0.0023594	0.0037822000000000003	0.0061414	1.2624157358257075	0.6312078679128538
Parallel	2	2	0.002359399999999998	0.0020282	0.0043874	1.7671058029812647	0.8835529014906324
Parallel	2	4	0.0025060000000000004	0.002177	0.0046832	1.655491971301674	0.827745985650837
Parallel	2	8	0.002215999999999997	0.002233	0.0044449	1.7426387952348843	0.8713193976174421
Parallel	2	16	0.0023628000000000004	0.00777739999999999	0.01014019999999999	0.7645805802646892	0.3822902901323446
Parallel	4	1	0.001233400000000002	0.01337659999999999	0.01460980000000001	0.530671193308601	0.13266779832715025
Parallel	4	2	0.001071999999999998	0.009644	0.0107164	0.7234705684744878	0.18086764211862194
Parallel	4	4	0.0013166	0.0104204	0.01173700000000001	0.6605606202607139	0.16514015506517848
Parallel	4	8	0.0011576	0.0125832	0.01374080000000001	0.5642320680018631	0.14105801700046577
Parallel	4	16	0.0013014	0.0317456	0.033047	0.23460525917632463	0.05865131479408116
Parallel	8	1	0.0007768	0.0251106	0.0258878	0.29948469935645367	0.03743558741955671
Parallel	8	2	0.000697599999999999	0.0172658	0.0179634	0.431599804995045	0.05394976062438065
Parallel	8	4	0.000696000000000001	0.0202517999999997	0.0209478	0.3701104650607701	0.046263808132596264
Parallel	8	8	0.0007078	0.03053080000000004	0.0312388	0.24818494948589576	0.03102311868573697
Parallel	8	16	0.000947199999999999	0.07242180000000001	0.0733689999999999	0.1056713325791547	0.013208916572394338

Tabella 2: Measurement with O0 optimization

4.1.3 500 vertices - 75% of density

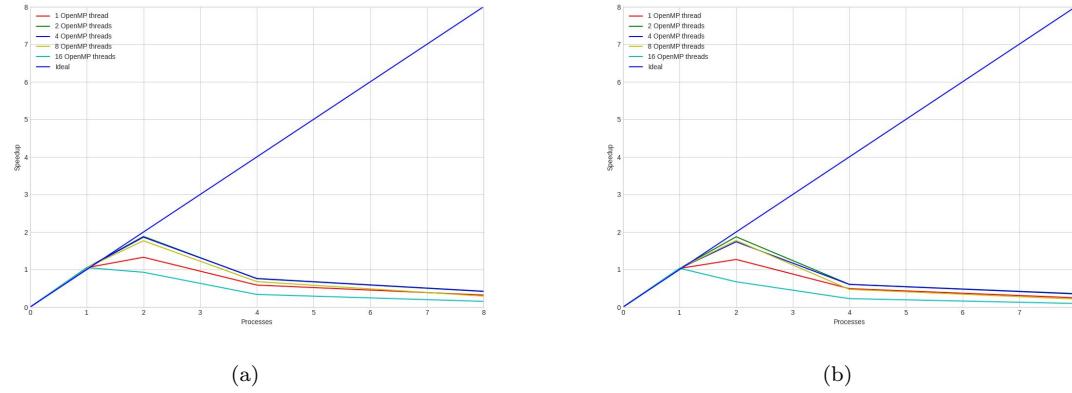


Figura 7: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

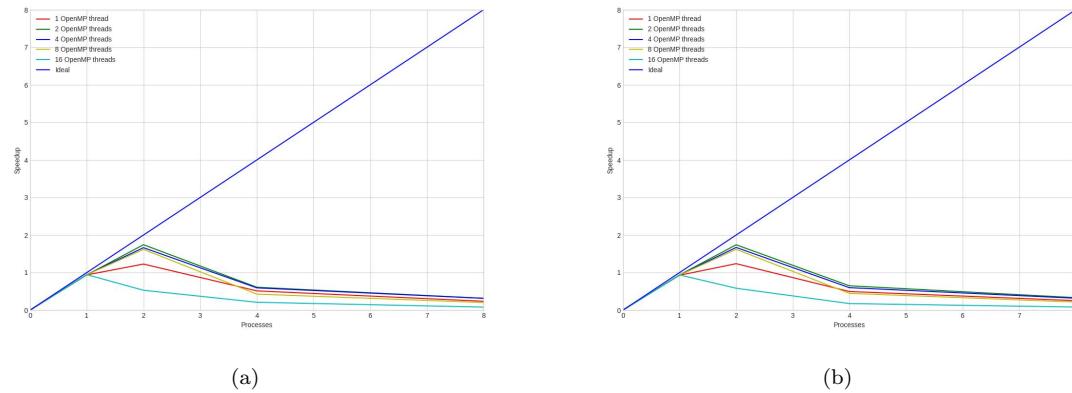


Figura 8: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Tarjan	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.012997	0.0	0.012997	1.0	1.0
Parallel	1	1	0.0123714	0.0	0.0123714	1.050568246116042	1.050568246116042
Parallel	1	2	0.01240159999999999	0.0	0.01240159999999999	1.0480099342020386	1.0480099342020386
Parallel	1	4	0.0124882	0.0	0.0124882	1.0407424608830735	1.0407424608830735
Parallel	1	8	0.0123866	0.0	0.0123866	1.0492790596289538	1.0492790596289538
Parallel	1	16	0.0123656000000000001	0.0	0.0123656000000000001	1.0510610079575595	1.0510610079575595
Parallel	2	1	0.0037416	0.006045	0.0097866	1.328043817464697	0.6640201908732348
Parallel	2	2	0.0037504	0.003151800000000006	0.006902	1.8830773688785858	0.9415386844392929
Parallel	2	4	0.003766	0.003208000000000003	0.0069742	1.863582919904792	0.931791459952396
Parallel	2	8	0.003791799999999993	0.003565	0.007357	1.7666168275112137	0.8833084137556069
Parallel	2	16	0.003742199999999998	0.01026419999999998	0.0140062	0.9279461952563864	0.4639730976281932
Parallel	4	1	0.0017048	0.020432200000000005	0.02213739999999995	0.5871059835391691	0.14677649588479227
Parallel	4	2	0.0018276	0.01535859999999997	0.0171864	0.7562374901084578	0.18905937252711444
Parallel	4	4	0.001580399999999998	0.0154654	0.0170458	0.7624752138356663	0.19061880345891657
Parallel	4	8	0.001847000000000001	0.0173048	0.01915159999999998	0.678637816161574	0.1696594540403935
Parallel	4	16	0.001757600000000002	0.03673340000000006	0.0384908	0.33766510438858116	0.08441627609714529
Parallel	8	1	0.001035800000000001	0.0395396	0.0405754	0.3203172365521967	0.040039654569024585
Parallel	8	2	0.0012642	0.0299956	0.03125980000000004	0.4157736133948393	0.051971701674354914
Parallel	8	4	0.0011126	0.0296396	0.03075240000000003	0.4226336806233009	0.052829210077912615
Parallel	8	8	0.001271600000000001	0.04311800000000004	0.0443894	0.2927951267644978	0.03659939084556223
Parallel	8	16	0.001272000000000001	0.0836908	0.0849626	0.15297319055678615	0.01912164881959827

Tabella 3: Measurement with $O0$ optimization

4.2 Case Study 2

4.2.1 1000 vertices - 25% of density

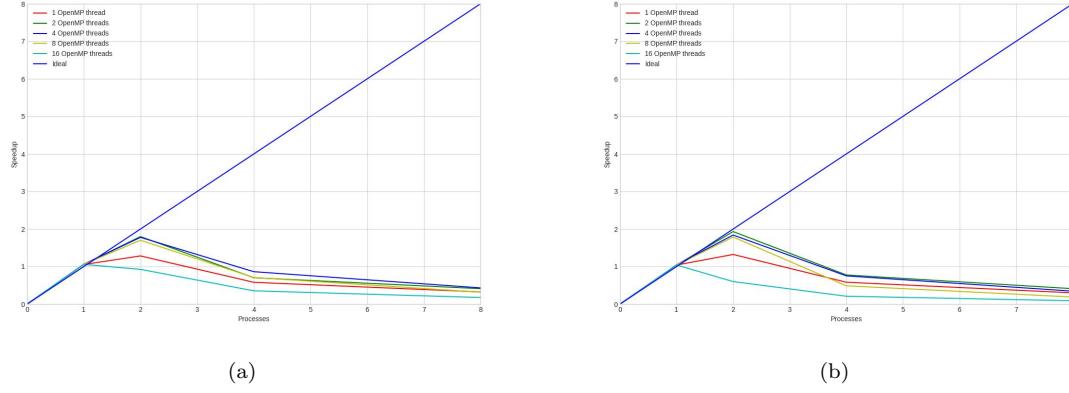


Figura 9: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

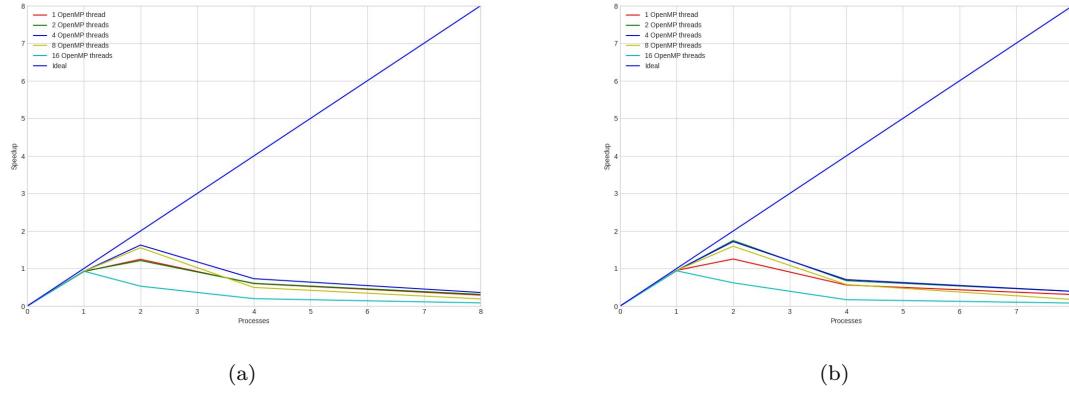


Figura 10: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Tarjan	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.0062416	0.0	0.0062416	1.0	1.0
Parallel	1	1	0.0059964000000000001	0.0	0.0059964000000000001	1.0408912013874991	1.0408912013874991
Parallel	1	2	0.0060056	0.0	0.0060056	1.0392966564539763	1.0392966564539763
Parallel	1	4	0.0059784	0.0	0.0059784	1.0440251572327044	1.0440251572327044
Parallel	1	8	0.006006	0.0	0.006006	1.0392274392274392	1.0392274392274392
Parallel	1	16	0.0059976000000000005	0.0	0.0059976000000000005	1.0406829398426036	1.0406829398426036
Parallel	2	1	0.0017104000000000002	0.003009999999999997	0.0047204	1.3222608253537838	0.6611304126768919
Parallel	2	2	0.001604	0.0016262	0.003230000000000002	1.9323839009287924	0.9661919504643962
Parallel	2	4	0.001678	0.001714200000000001	0.0033922	1.839985849890926	0.919992924945463
Parallel	2	8	0.0016158	0.0018766000000000002	0.0034926	1.787092710301781	0.8935463551508905
Parallel	2	16	0.001684799999999998	0.008705599999999999	0.010390199999999999	0.600719909145156	0.3003595954572578
Parallel	4	1	0.0006928	0.010013000000000001	0.010706	0.583000186811134	0.1457500467027835
Parallel	4	2	0.0007056	0.0073324	0.008038	0.7765115700422991	0.19412789251057477
Parallel	4	4	0.0007256000000000001	0.007575	0.0083004	0.751963760782613	0.1879909401956324
Parallel	4	8	0.000841	0.01196940000000002	0.01281019999999999	0.4872367332282088	0.1218091833070522
Parallel	4	16	0.000759	0.02909260000000003	0.0298518	0.20908621925646023	0.05227155481411506
Parallel	8	1	0.000602	0.020151400000000003	0.020753599999999997	0.3007478220646057	0.03759347775807571
Parallel	8	2	0.0006404	0.0145428	0.01518319999999999	0.41108593708836083	0.051385742136045104
Parallel	8	4	0.0006164	0.0174754	0.0180918	0.3449960755701478	0.04312450944626847
Parallel	8	8	0.0005352	0.0325566	0.03309199999999996	0.1886135621902575	0.023576695273782186
Parallel	8	16	0.0005718	0.0693241999999999	0.069896	0.08929838617374385	0.01116229827171798

Tabella 4: Measurement with O1 optimization

4.2.2 1000 vertices - 50% of density

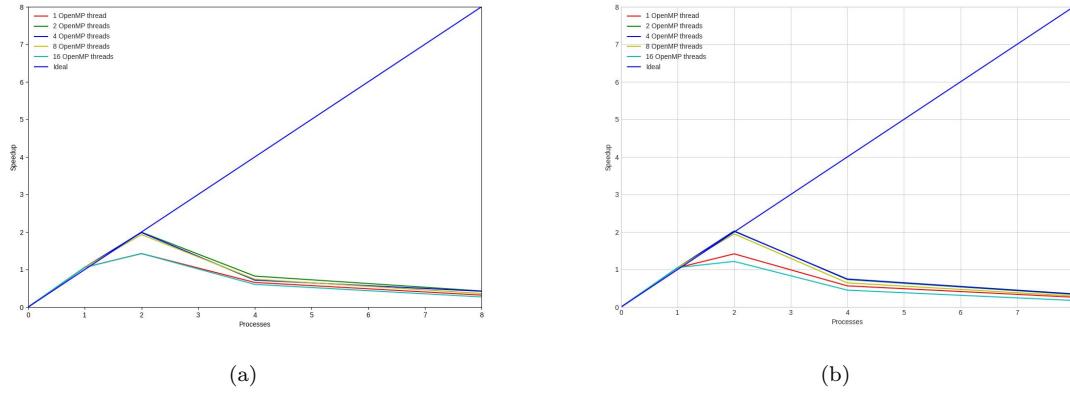


Figura 11: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

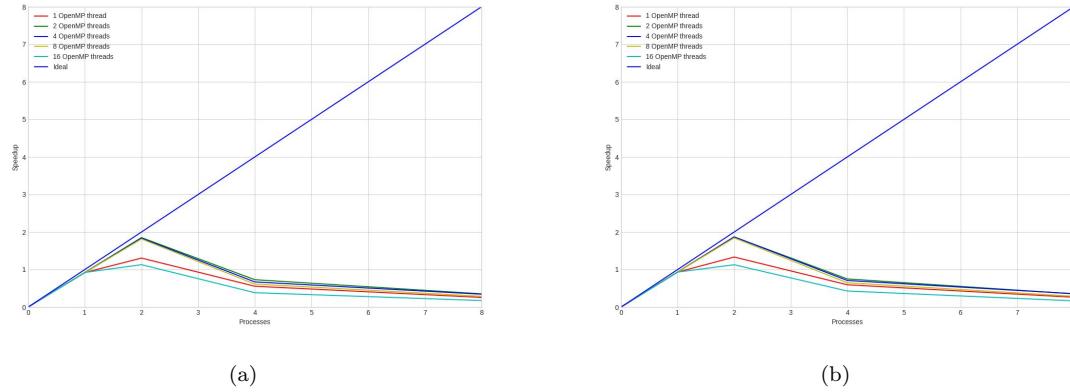


Figura 12: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Tarjan	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.0182018	0.0	0.0182018	1.0	1.0
Parallel	1	1	0.0173662	0.0	0.0173662	1.0481164561043868	1.0481164561043868
Parallel	1	2	0.017338799999999998	0.0	0.01733879999999998	1.049772763974439	1.049772763974439
Parallel	1	4	0.0173888	0.0	0.0173888	1.0467542326094958	1.0467542326094958
Parallel	1	8	0.0173158	0.0	0.0173158	1.051167142147634	1.051167142147634
Parallel	1	16	0.01728	0.0	0.01728	1.0533449074074075	1.0533449074074075
Parallel	2	1	0.0047852	0.00804079999999999	0.012826	1.4191330110712614	0.7095665055356307
Parallel	2	2	0.0047798	0.00418619999999999	0.0089664	2.030000892219843	1.0150004461099216
Parallel	2	4	0.0047786	0.0042382	0.009017	2.018609293556615	1.0093046467783076
Parallel	2	8	0.00489219999999999	0.0044402	0.00933219999999999	1.9504296950343973	0.9752148475171987
Parallel	2	16	0.004816200000000005	0.010138600000000001	0.014954600000000002	1.2171372019311784	0.6085686009655892
Parallel	4	1	0.001921399999999998	0.030237000000000003	0.0321578	0.5660150880968225	0.14150377202420564
Parallel	4	2	0.001769199999999999	0.0230006	0.02476959999999996	0.734843253019832	0.1837110813254958
Parallel	4	4	0.0018144	0.0225084	0.024323	0.7483369650125395	0.1870842412531349
Parallel	4	8	0.002053	0.0261378	0.0281908	0.645664543042411	0.16141613576060276
Parallel	4	16	0.002034	0.038254400000000001	0.0402888	0.4517831258314966	0.11294578145787415
Parallel	8	1	0.001029400000000002	0.0680304	0.0690596	0.2635665425226906	0.03294581781533632
Parallel	8	2	0.0011766	0.0524874	0.0536638	0.3391820929565182	0.042397761619564774
Parallel	8	4	0.0010342	0.0506008	0.051635	0.3525089571027404	0.04406361963784255
Parallel	8	8	0.0010764	0.06043860000000001	0.061515	0.2958920588474356	0.03698650735592945
Parallel	8	16	0.001166199999999998	0.1013552	0.1025215999999999	0.17754112304138836	0.022192640380173545

Tabella 5: Measurement with $O1$ optimization

4.2.3 1000 vertices - 75% of density

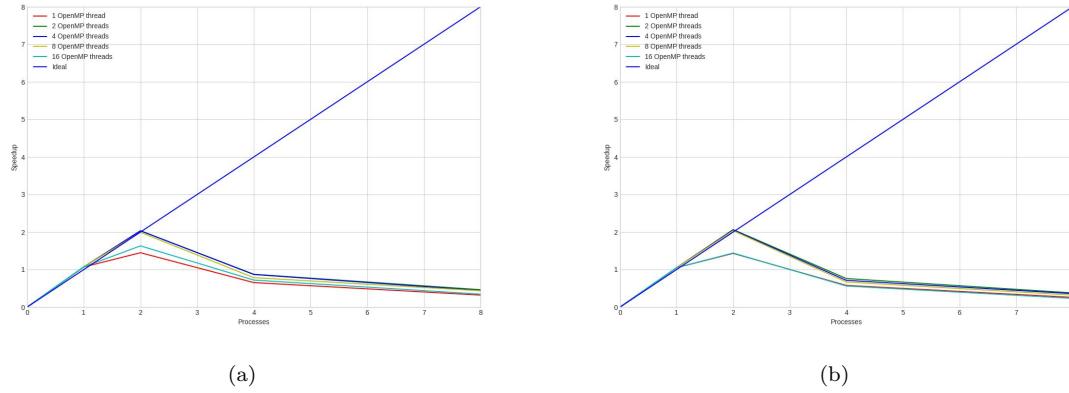


Figura 13: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

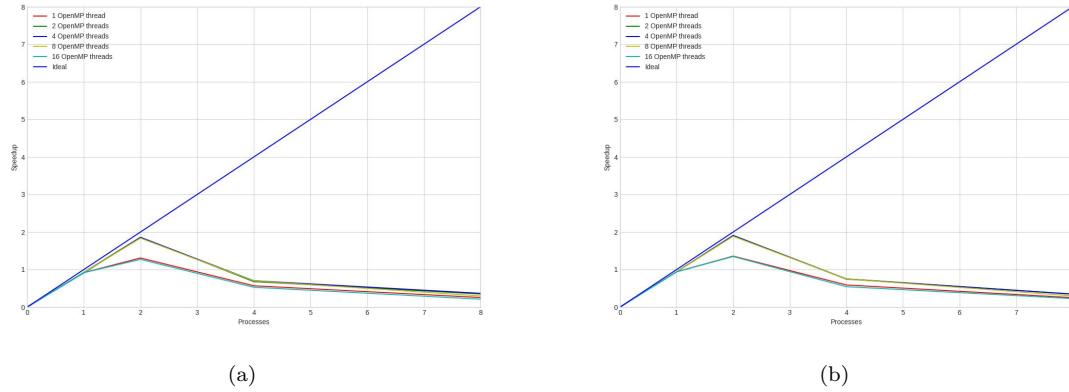


Figura 14: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Tarjan	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.030330600000000003	0.0	0.030330600000000003	1.0	1.0
Parallel	1	1	0.028916800000000003	0.0	0.028916800000000003	1.0488919935815857	1.0488919935815857
Parallel	1	2	0.0289202	0.0	0.0289202	1.048768680714518	1.048768680714518
Parallel	1	4	0.0288968	0.0	0.0288968	1.04961179507765566	1.04961179507765566
Parallel	1	8	0.028938399999999996	0.0	0.028938399999999996	1.0481090868880105	1.0481090868880105
Parallel	1	16	0.028862	0.0	0.028862	1.0508835146559492	1.0508835146559492
Parallel	2	1	0.007817000000000001	0.01338160000000002	0.0211986	1.4307831649259857	0.7153915824629928
Parallel	2	2	0.007829	0.0068596	0.01468839999999999	2.0649355954358546	1.0324677977179273
Parallel	2	4	0.0078662	0.006890999999999996	0.014757600000000001	2.0552528866482356	1.0276264433241178
Parallel	2	8	0.0078308	0.00706379999999999	0.0148946	2.03634874384005	1.018174371920025
Parallel	2	16	0.0078954	0.0131566	0.0210522	1.4407330350272183	0.7203665175136091
Parallel	4	1	0.002909	0.04959280000000006	0.0525018	0.577705907226038	0.1444264768065095
Parallel	4	2	0.0029782	0.03686120000000004	0.03983979999999995	0.7613140628216006	0.19032851570540016
Parallel	4	4	0.0029094	0.0396074	0.0425168	0.7133791818763408	0.1783447954690852
Parallel	4	8	0.003094999999999997	0.04226220000000001	0.045357	0.6687082478999934	0.16717706197499835
Parallel	4	16	0.003022399999999997	0.05100200000000006	0.0540238	0.5614303325571323	0.14035758313928307
Parallel	8	1	0.001638	0.1155638000000002	0.1172018	0.25878954077497107	0.032348692596871384
Parallel	8	2	0.001442599999999998	0.0787549999999999	0.0801974	0.37819929324392065	0.04727491165549008
Parallel	8	4	0.001525599999999998	0.0831544	0.0846802	0.3581781809679241	0.04477227262099051
Parallel	8	8	0.001716399999999999	0.0928762	0.0945920000000001	0.32064656630581867	0.040080820788227334
Parallel	8	16	0.0016976	0.1290541999999998	0.1307512	0.2319718671798041	0.02899648339747551

Tabella 6: Measurement with $O1$ optimization

4.2.4 1000 vertices - Special Cases - 0% of density

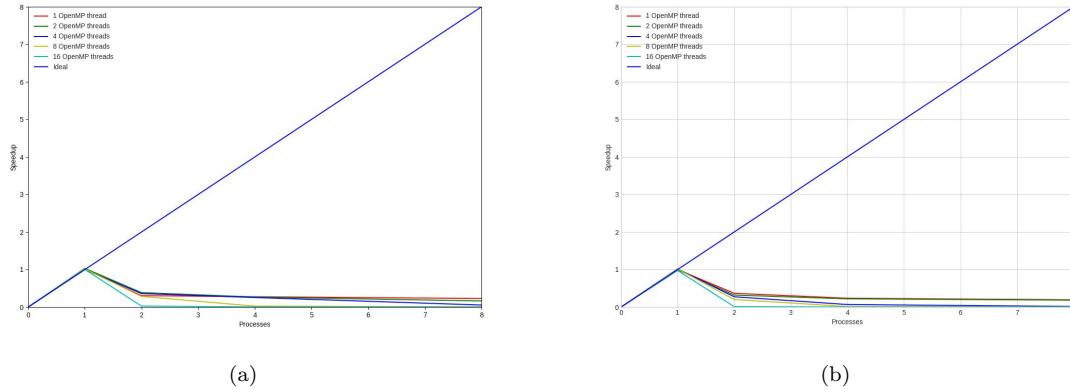


Figura 15: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

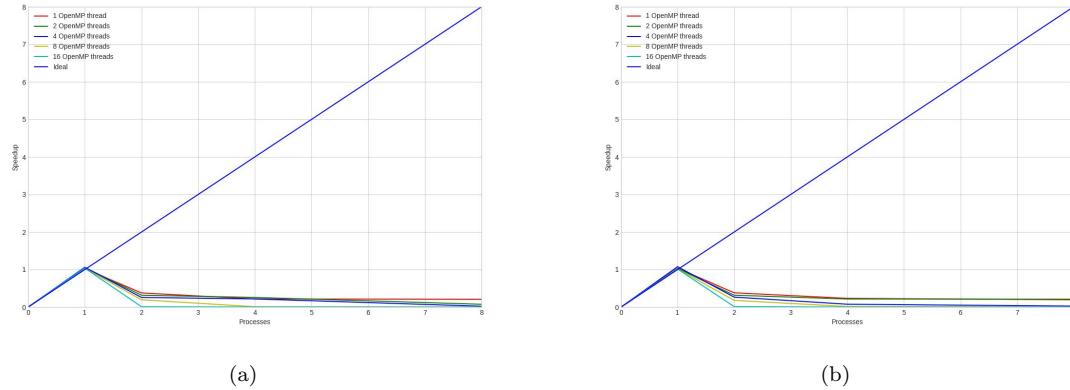


Figura 16: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Tarjan	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.0001796	0.0	0.0001796	1.0	1.0
Parallel	1	1	0.0001778	0.0	0.0001778	1.0101237345331833	1.0101237345331833
Parallel	1	2	0.000174	0.0	0.000174	1.032183908045977	1.032183908045977
Parallel	1	4	0.0001788	0.0	0.0001788	1.0044742729306488	1.0044742729306488
Parallel	1	8	0.0001762	0.0	0.0001762	1.0192962542565267	1.0192962542565267
Parallel	1	16	0.000179	0.0	0.000179	1.0033519553072627	1.0033519553072627
Parallel	2	1	0.0003109999999999997	0.0002768	0.000588	0.3054421768707483	0.15272108843537416
Parallel	2	2	0.0002923999999999995	0.0001721999999999998	0.0004648	0.3864027538726334	0.1932013769363167
Parallel	2	4	0.0002905999999999996	0.000207	0.0004971999999999999	0.3612228479485117	0.18061143769452586
Parallel	2	8	0.0003272000000000004	0.0003016	0.0006288	0.2856234096692112	0.1428117048346056
Parallel	2	16	0.0003516	0.00620779999999999	0.0065594	0.027380553099368845	0.013690276549684422
Parallel	4	1	0.0003886	0.000261999999999997	0.0006512	0.2757985257985258	0.06894963144963145
Parallel	4	2	0.0004142	0.0002536	0.000668	0.2688622754491018	0.06721556886227545
Parallel	4	4	0.0004482	0.0002556000000000003	0.000704	0.2551136363636364	0.0637784090909091
Parallel	4	8	0.0004188	0.006672	0.00709079999999995	0.025328594798894342	0.006332148699723586
Parallel	4	16	0.0004046	0.0321834	0.032588	0.005511231128022585	0.0013778077820056462
Parallel	8	1	0.0004486000000000006	0.0003353999999999997	0.0007842	0.22902320836521298	0.028627901045651622
Parallel	8	2	0.0004548	0.0006468	0.0011014	0.1630651897584892	0.02038314871981115
Parallel	8	4	0.0004827999999999997	0.002801	0.0032842	0.05468607271177151	0.006835759088971439
Parallel	8	8	0.0004474000000000003	0.02930400000000004	0.0297516	0.006036650129740921	0.0007545812662176152
Parallel	8	16	0.0004478000000000004	0.069749	0.07019700000000001	0.0025585138966052676	0.00031981423707565845

Tabella 7: Measurement with O0 optimization

4.2.5 1000 vertices - Special Cases - 100% of density

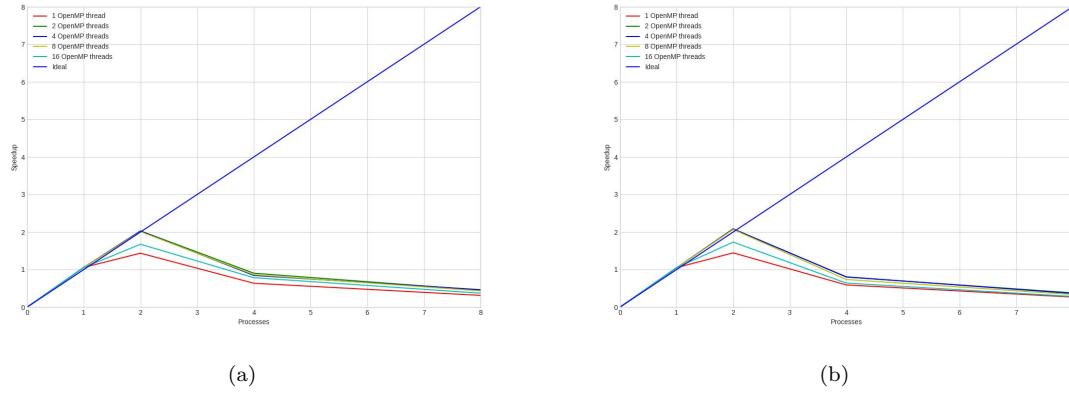


Figura 17: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

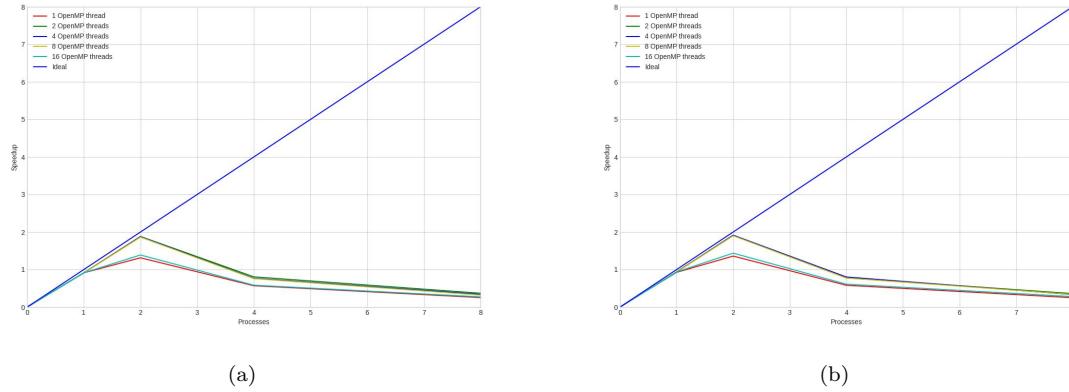


Figura 18: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Tarjan	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.04883020000000004	0.0	0.04883020000000004	1.0	1.0
Parallel	1	1	0.0464392	0.0	0.0464392	1.0514866750503886	1.0514866750503886
Parallel	1	2	0.04677560000000001	0.0	0.04677560000000001	1.0439246102668913	1.0439246102668913
Parallel	1	4	0.04661979999999996	0.0	0.04661979999999996	1.04741333081652	1.04741333081652
Parallel	1	8	0.0465936	0.0	0.0465936	1.0480023007451669	1.0480023007451669
Parallel	1	16	0.04650520000000003	0.0	0.04650520000000003	1.0499944092273552	1.0499944092273552
Parallel	2	1	0.01252800000000001	0.0212566	0.0337846	1.4453390006097455	0.7226695003048728
Parallel	2	2	0.01259859999999998	0.0108266	0.0234252	2.0845158205693015	1.0422579102846508
Parallel	2	4	0.01259640000000002	0.0108528	0.023449	2.0824001023497805	1.0412000511748902
Parallel	2	8	0.01260599999999997	0.01101539999999998	0.023621200000000002	2.06721927759809	1.033609638799045
Parallel	2	16	0.01271779999999998	0.01547680000000002	0.02819459999999997	1.7318990161236552	0.8659495080618276
Parallel	4	1	0.00431320000000005	0.0784266	0.0827398	0.5901657968716386	0.14754144921790965
Parallel	4	2	0.00443308	0.05672439999999994	0.0610554	0.7997687346246196	0.1999421836561549
Parallel	4	4	0.0044066	0.056159	0.0605658	0.8062338811672595	0.20155847029181487
Parallel	4	8	0.004451399999999995	0.0619065999999999	0.066358	0.7358600319479189	0.18396500798697973
Parallel	4	16	0.004565	0.0718104	0.0763752	0.6393462799442752	0.1598365699860688
Parallel	8	1	0.0021868	0.1804298	0.1826168	0.26739160909620585	0.03342395113702573
Parallel	8	2	0.0022412	0.12609720000000002	0.1283378	0.38048182219112375	0.04756022777389047
Parallel	8	4	0.002529600000000003	0.1323329999999998	0.1348626	0.362073695746634	0.04525921196832925
Parallel	8	8	0.002550399999999995	0.1429974	0.145548	0.33549207134416137	0.04193650891802017
Parallel	8	16	0.0054926	0.1672506	0.1727436	0.2826744377215712	0.0353343047151964

Tabella 8: Measurement with $O1$ optimization

4.3 Case Study 3

4.3.1 2500 vertices - 25% of density

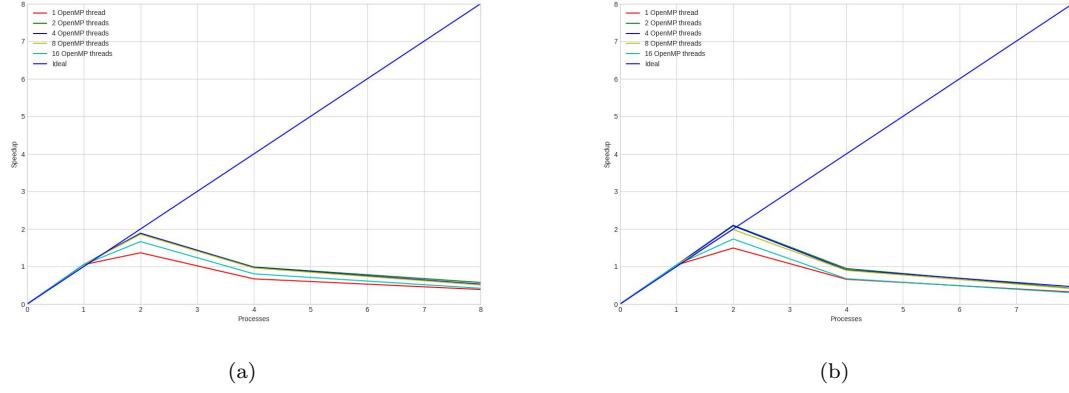


Figura 19: (a) Speedup with no optimization - (b) Speedup with *O1* optimization

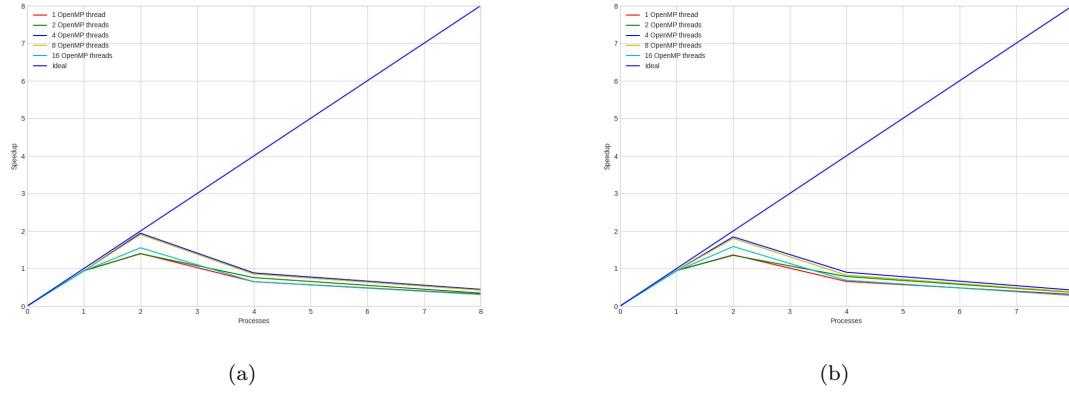


Figura 20: (a) Speedup with *O2* optimization - (b) Speedup with *O3* optimization

Version	Processes	OMP	Tarjan	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.0443628	0.0	0.0443628	1.0	1.0
Parallel	1	1	0.0425736	0.0	0.0425736	1.0420260443091494	1.0420260443091494
Parallel	1	2	0.042465800000000005	0.0	0.042465800000000005	1.044671241328316	1.044671241328316
Parallel	1	4	0.042559599999999996	0.0	0.042559599999999996	1.0423688192558203	1.0423688192558203
Parallel	1	8	0.0424726	0.0	0.0424726	1.0445039860992735	1.0445039860992735
Parallel	1	16	0.042426399999999996	0.0	0.042426399999999996	1.0456413930948656	1.0456413930948656
Parallel	2	1	0.0122476	0.0174326	0.0296802	1.4946934319849596	0.7473467159924798
Parallel	2	2	0.01205279999999999	0.0090528	0.0211056	2.101944507618831	1.0509722538094155
Parallel	2	4	0.012117	0.00915439999999998	0.02127140000000003	2.085560893970307	1.0427804469851536
Parallel	2	8	0.0121748	0.0101434	0.02231819999999996	1.987740946850553	0.9938704734252765
Parallel	2	16	0.012333	0.0132554	0.0255884	1.7337074611933532	0.8668537305966766
Parallel	4	1	0.004413400000000005	0.0627094	0.0671228	0.6609199854594863	0.16522999636487157
Parallel	4	2	0.004571	0.04223200000000006	0.046803	0.947862316518172	0.236965579129543
Parallel	4	4	0.004473599999999995	0.04411959999999995	0.0485932	0.9129425516327387	0.22823563790818469
Parallel	4	8	0.00496859999999999	0.0444124	0.04938140000000006	0.8983706415775979	0.22459266039439948
Parallel	4	16	0.0046974	0.06093859999999996	0.0656364	0.6758871601733185	0.16897179004332963
Parallel	8	1	0.00240679999999997	0.1355576	0.137964	0.32155344872575453	0.040194181090719316
Parallel	8	2	0.0023882	0.1030402	0.1054288000000002	0.4207844535838404	0.05259805669798005
Parallel	8	4	0.002622	0.0930526000000001	0.0956744	0.46368516551972105	0.05796064568996513
Parallel	8	8	0.0027224	0.105266	0.1079884000000001	0.41081079078864025	0.0513513488458003
Parallel	8	16	0.002955599999999996	0.1440043999999998	0.1469596	0.3018707182109913	0.037733839776373915

Tabella 9: Measurement with *O1* optimization

4.3.2 2500 vertices - 50% of density

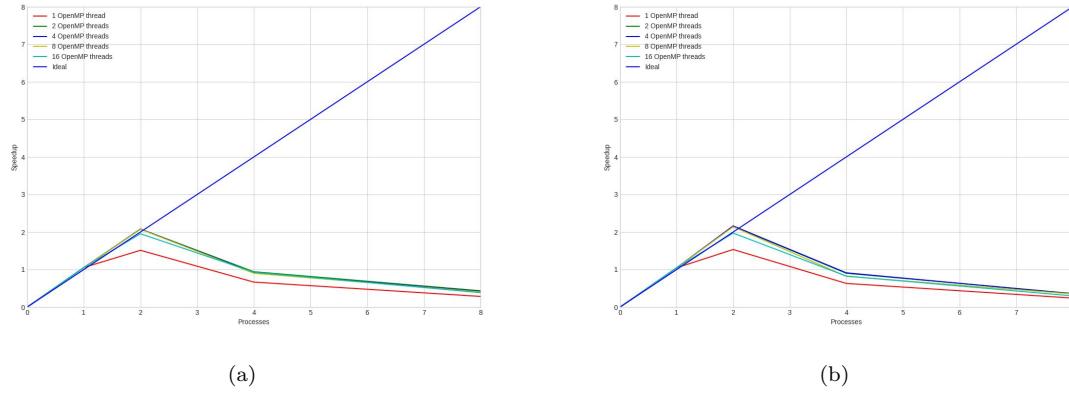


Figura 21: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

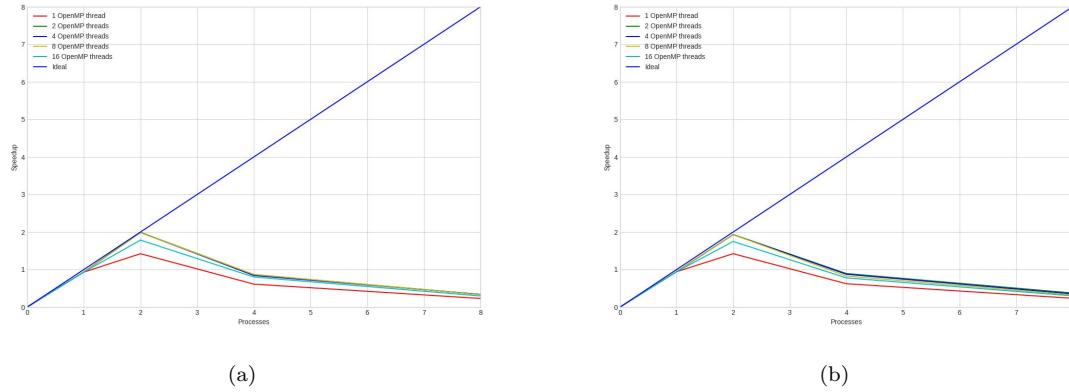


Figura 22: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Tarjan	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.13916080000000003	0.0	0.13916080000000003	1.0	1.0
Parallel	1	1	0.13302740000000002	0.0	0.13302740000000002	1.0461062908844345	1.0461062908844345
Parallel	1	2	0.13307000000000002	0.0	0.13307000000000002	1.0457713985120614	1.0457713985120614
Parallel	1	4	0.1330088	0.0	0.1330088	1.0462525787767427	1.0462525787767427
Parallel	1	8	0.1331484	0.0	0.1331484	1.045155630860003	1.045155630860003
Parallel	1	16	0.13308	0.0	0.13308	1.0456928163510673	1.0456928163510673
Parallel	2	1	0.03763180000000001	0.053057	0.09068860000000001	1.5344905533881879	0.7672452766940939
Parallel	2	2	0.0377744	0.026872	0.06464620000000001	2.1526524374209157	1.0763262187104579
Parallel	2	4	0.0374056	0.02693820000000002	0.06434360000000001	2.162776095835483	1.0813880479177416
Parallel	2	8	0.037787	0.027248	0.06503519999999999	2.1397766132801936	1.0698883066400968
Parallel	2	16	0.037883	0.0327006	0.07058339999999999	1.971579719877479	0.9857898599387395
Parallel	4	1	0.0126972	0.20786360000000004	0.220561	0.6309401934158805	0.15773504835397012
Parallel	4	2	0.01266020000000002	0.1416944	0.1543548	0.9015644476232683	0.22539111190581707
Parallel	4	4	0.0126626	0.1396828	0.152345	0.9134582690603565	0.22836456726508914
Parallel	4	8	0.01299520000000002	0.1568528	0.1698481999999998	0.819324549803884	0.204831137450971
Parallel	4	16	0.0131714	0.1551554	0.1683266	0.8267308910178192	0.2066827227544548
Parallel	8	1	0.0062284	0.5711672	0.5773954	0.24101473617559133	0.030126842021948916
Parallel	8	2	0.0062346	0.3775605999999997	0.383795	0.3625914876431429	0.045323935955392865
Parallel	8	4	0.0066832	0.384877199999999	0.39156040000000003	0.3554005972003298	0.044425074650041224
Parallel	8	8	0.0068606	0.39002220000000004	0.3968826	0.3506346713108613	0.04382933391385766
Parallel	8	16	0.007084200000000006	0.4594347999999999	0.4665187999999996	0.29829623157737706	0.03728702894717213

Tabella 10: Measurement with $O1$ optimization

4.3.3 2500 vertices - 75% of density

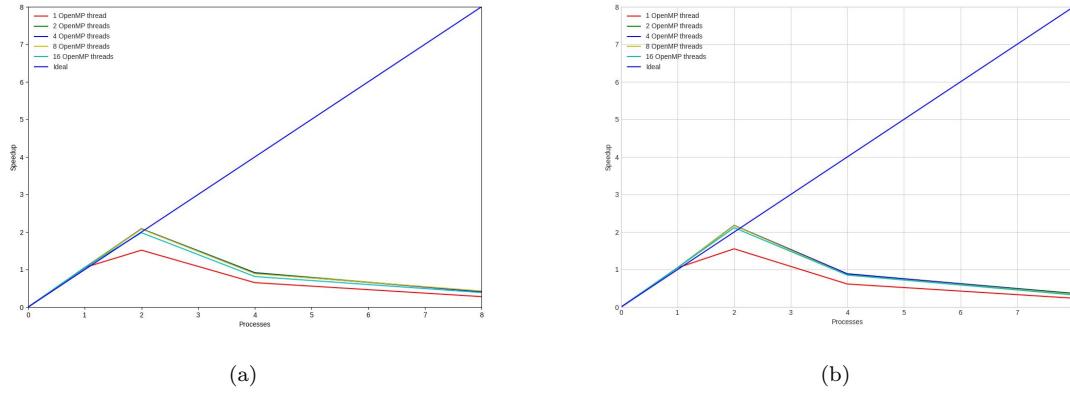


Figura 23: (a) Speedup with no optimization - (b) Speedup with *O1* optimization

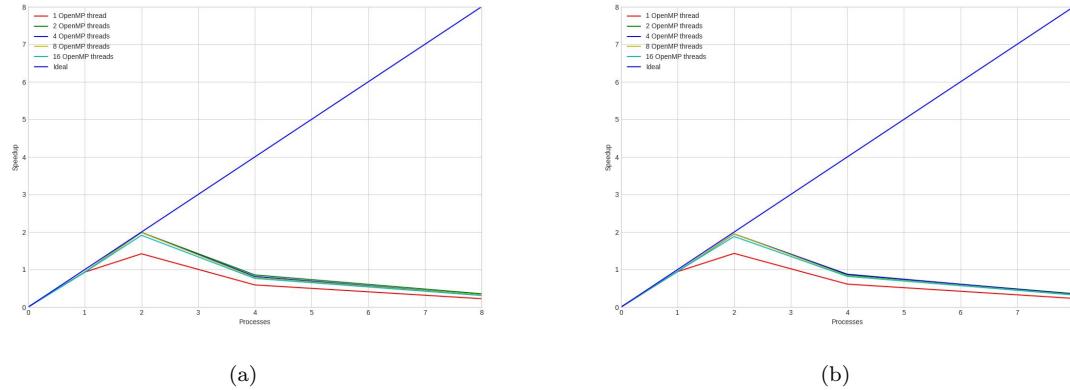


Figura 24: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Tarjan	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.42665120000000006	0.0	0.42665120000000006	1.0	1.0
Parallel	1	1	0.40412860000000006	0.0	0.40412860000000006	1.0557312696997936	1.0557312696997936
Parallel	1	2	0.404099	0.0	0.404099	1.0558086013575883	1.0558086013575883
Parallel	1	4	0.41049420000000003	0.0	0.41049420000000003	1.0393598740250167	1.0393598740250167
Parallel	1	8	0.40713039999999995	0.0	0.40713039999999995	1.04794729158029	1.04794729158029
Parallel	1	16	0.40456380000000003	0.0	0.40456380000000003	1.0545955915976666	1.0545955915976666
Parallel	2	1	0.1265274	0.1548254	0.28135260000000006	1.5164288511995268	0.7582144255997634
Parallel	2	2	0.12657639999999998	0.07782039999999998	0.2043968	2.0873673169051576	1.0436836584525788
Parallel	2	4	0.1266986	0.07795579999999999	0.20465460000000002	2.084737894970355	1.0423689474851776
Parallel	2	8	0.12664340000000002	0.0786228	0.20526619999999998	2.0785263233791054	1.0392631616895527
Parallel	2	16	0.12647919999999999	0.0887808	0.21252599999999995	1.9820273158041446	0.9910136579020723
Parallel	4	1	0.0522686	0.6028165999999999	0.6550848	0.6512915579784481	0.16282288949461202
Parallel	4	2	0.052274	0.41169399999999995	0.46396740000000003	0.9195715043772472	0.2298928760943118
Parallel	4	4	0.0522532	0.4202108	0.4724642	0.9030339229935307	0.22575848074838267
Parallel	4	8	0.05237419999999996	0.4262898	0.4786639999999999	0.891337556198085	0.22283438904952124
Parallel	4	16	0.05253539999999996	0.4710888	0.5236240000000001	0.814804516217744	0.203701129054436
Parallel	8	1	0.030617	1.4972408	1.5278576	0.2792480136892339	0.03490600171115424
Parallel	8	2	0.0305944	1.0106538	1.0412482	0.4097497599515659	0.051218719993945734
Parallel	8	4	0.03082080000000002	1.0092431999999998	1.0400639999999999	0.4102162943818843	0.051277036797735535
Parallel	8	8	0.0312626	0.9612740000000001	0.9925362	0.4298595859778213	0.053732448247227664
Parallel	8	16	0.03252240000000001	1.0702375999999998	1.1027602	0.3868939049486915	0.04836173811858644

Tabella 11: Measurement with O0 optimization

4.4 Case Study 4

4.4.1 10000 vertices - 25% of density

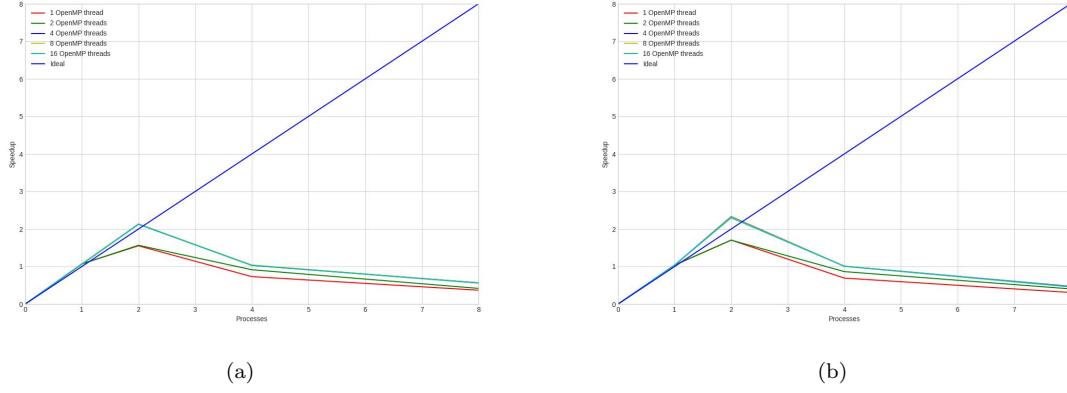


Figura 25: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

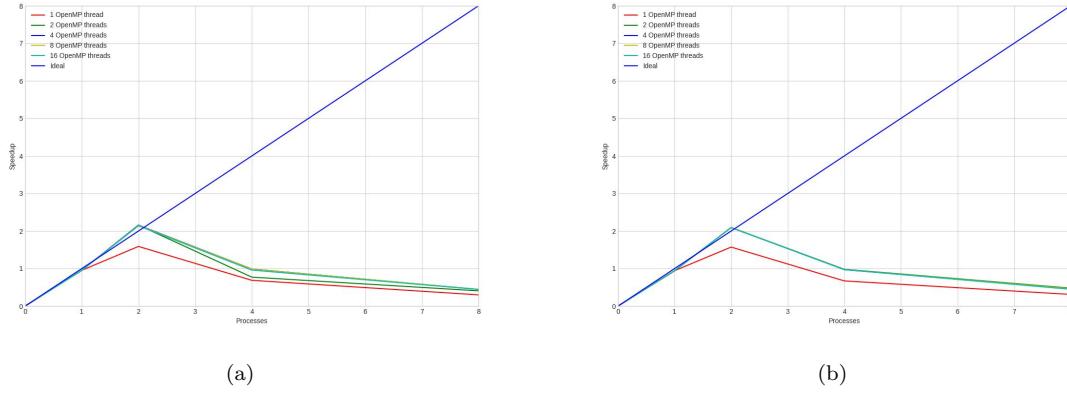


Figura 26: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Tarjan	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.890088199999999	0.0	0.890088199999999	1.0	1.0
Parallel	1	1	0.858418	0.0	0.858418	1.0368936811669838	1.0368936811669838
Parallel	1	2	0.8573674	0.0	0.8573674	1.038164268900357	1.038164268900357
Parallel	1	4	0.8581804	0.0	0.8581804	1.0371807605953247	1.0371807605953247
Parallel	1	8	0.8603012000000001	0.0	0.8603012000000001	1.0346239200875227	1.0346239200875227
Parallel	1	16	0.8575512	0.0	0.8575512	1.0379417578798793	1.0379417578798793
Parallel	2	1	0.2327393999999999	0.2892004	0.5219396	1.7053471321202682	0.8526735660601341
Parallel	2	2	0.2327043999999998	0.28963	0.5223342	1.7040588190472687	0.8520294095236344
Parallel	2	4	0.2337522	0.1481583999999997	0.3819106	2.3306192601095646	1.1653096300547823
Parallel	2	8	0.2327020000000002	0.1507102000000002	0.383412	2.3214928066935827	1.1607464033467914
Parallel	2	16	0.2334785999999998	0.1535436000000003	0.3870226	2.299835203422287	1.1499176017111143
Parallel	4	1	0.0888838	1.196400999999998	1.2852846000000002	0.6925222631625709	0.17313056579064273
Parallel	4	2	0.089077	0.9389668	1.028044	0.8658074946208527	0.21645187365521318
Parallel	4	4	0.0891376	0.791598799999999	0.8807362	1.0106183894791652	0.2526545973697913
Parallel	4	8	0.0900496000000001	0.789336199999999	0.8793858	1.0121703125067518	0.25304257812668796
Parallel	4	16	0.0900765999999999	0.794416799999999	0.884493399999999	1.0063254287708647	0.2515813571927162
Parallel	8	1	0.0486286	2.827530600000002	2.8761592	0.3094711168978407	0.038683889612230085
Parallel	8	2	0.0483292	2.1480826	2.196411800000004	0.40524650250012306	0.05065581281251538
Parallel	8	4	0.0515746000000005	1.8299406	1.8815154	0.473069845721167	0.059133730715145875
Parallel	8	8	0.0531627999999996	1.8656082	1.9187708	0.4638845869449336	0.0579855733681167
Parallel	8	16	0.0538454	1.8911652	1.945010599999998	0.4576240059648	0.05720330007456

Tabella 12: Measurement with O1 optimization

4.4.2 10000 vertices - 50% of density

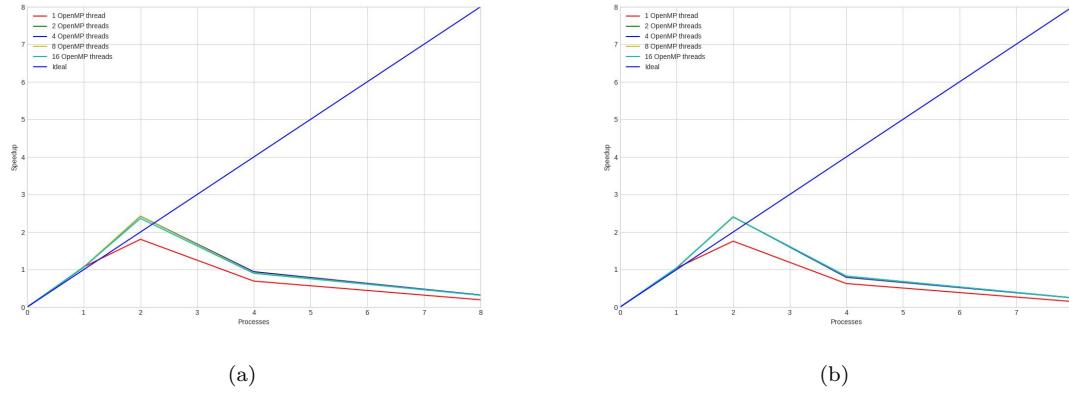


Figura 27: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

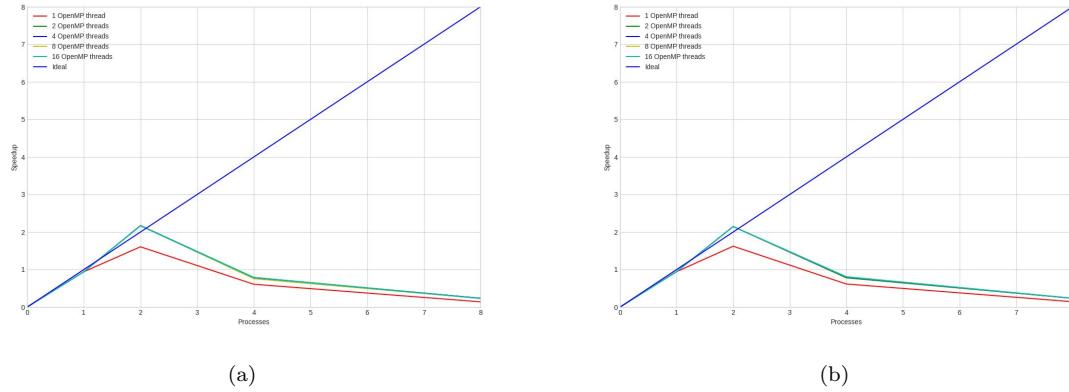


Figura 28: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Tarjan	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	5.228125799999999	0.0	5.228125799999999	1.0	1.0
Parallel	1	1	4.8749970000000005	0.0	4.8749970000000005	1.0724367214995205	1.0724367214995205
Parallel	1	2	4.9004150000000001	0.0	4.9004150000000001	1.0668740912759427	1.0668740912759427
Parallel	1	4	4.9207066000000001	0.0	4.9207066000000001	1.0624746047650957	1.0624746047650957
Parallel	1	8	4.8850076	0.0	4.8850076	1.070239030948488	1.070239030948488
Parallel	1	16	4.909088	0.0	4.909088	1.0649892199936117	1.0649892199936117
Parallel	2	1	1.4243508000000003	1.470082	2.8944327999999997	1.806269539234077	0.9031347696170385
Parallel	2	2	1.425722	0.736356	2.1620777999999996	2.418102530815496	1.209051265407748
Parallel	2	4	1.4233513999999998	0.7365878	2.1599394000000003	2.4204965194856847	1.2102482597428423
Parallel	2	8	1.4242178	0.7375038	2.161722	2.4185005287451387	1.2092502643725693
Parallel	2	16	1.4690277999999999	0.7448782	2.2139062000000003	2.3614938157723206	1.1807469078861603
Parallel	4	1	0.5742687999999999	6.970635400000001	7.544904	0.692934701356041	0.17323367533901024
Parallel	4	2	0.5746612	5.109519199999999	5.684180400000001	0.9197677469912809	0.22994193674782024
Parallel	4	4	0.5742339999999999	4.964993	5.539227	0.9438367122343964	0.2359591780585991
Parallel	4	8	0.5741222	5.1240908	5.698213200000005	0.9175026655724287	0.22937566639310716
Parallel	4	16	0.5747658	5.2438034	5.818569200000001	0.8985242970041498	0.22463107425103745
Parallel	8	1	0.3206894	26.510825600000004	26.83151479999997	0.19485019161124661	0.024356273951405827
Parallel	8	2	0.3200786	16.0177628	16.3378412	0.32000101702543166	0.04000012712817896
Parallel	8	4	0.3212148	15.810775800000002	16.1319906	0.32408435695468346	0.04051054461933543
Parallel	8	8	0.3217104	16.1981328	16.519843399999996	0.3164755060571579	0.03955943825714474
Parallel	8	16	0.3266297999999997	16.293736399999997	16.620366	0.31456141218550776	0.03932017652318847

Tabella 13: Measurement with $O0$ optimization

4.4.3 10000 vertices - 75% of density

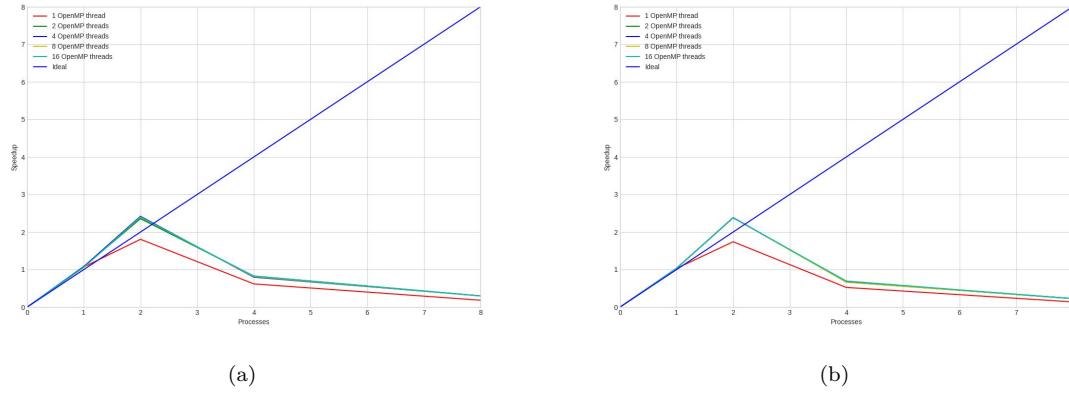


Figura 29: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

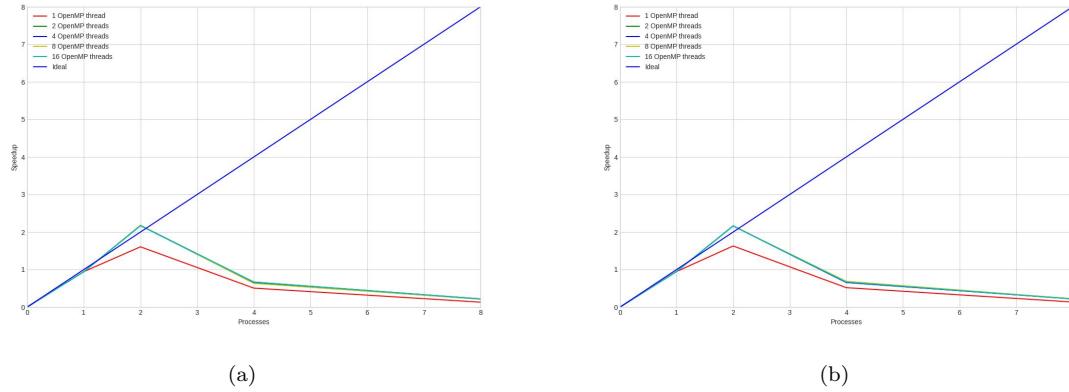


Figura 30: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Tarjan	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	8.7069948	0.0	8.7069948	1.0	1.0
Parallel	1	1	8.149559400000001	0.0	8.149559400000001	1.0684006794281418	1.0684006794281418
Parallel	1	2	8.1039906	0.0	8.1039906	1.0744083044716268	1.0744083044716268
Parallel	1	4	8.049288	0.0	8.049288	1.081709935089911	1.081709935089911
Parallel	1	8	8.2708258	0.0	8.2708258	1.052735846522121	1.052735846522121
Parallel	1	16	8.187621799999999	0.0	8.187621799999999	1.0634339265646102	1.0634339265646102
Parallel	2	1	2.3697942	2.4507336	4.8205278	1.806232670206777	0.9031163351033885
Parallel	2	2	2.4678876	1.2275548	3.6954424	2.356144097929926	1.178072048964963
Parallel	2	4	2.370035	1.2281488	3.5981838	2.4198304711393566	1.2099152355696783
Parallel	2	8	2.3960332	1.2287218000000002	3.6247548000000003	2.402092080821577	1.2010460404107886
Parallel	2	16	2.4062744	1.2383694	3.6446438	2.3889837465049397	1.1944918732524699
Parallel	4	1	0.9543166	13.135392	14.0897088	0.6179683997443581	0.15449209993608953
Parallel	4	2	0.9545667999999999	9.563544799999999	10.518111600000001	0.8278096992239558	0.20695242480598894
Parallel	4	4	0.9542000000000002	9.979966000000002	10.934196600000002	0.7963086012190415	0.19907715030476036
Parallel	4	8	0.9551701999999999	9.796744	10.7519142	0.8098088059519671	0.20245220148799178
Parallel	4	16	0.9568540000000001	9.5875586	10.544412600000001	0.8257448878660154	0.20643622196650385
Parallel	8	1	0.5322442000000001	47.077196400000005	47.609440400000004	0.18288378789682225	0.02286047348710278
Parallel	8	2	0.5302334	28.635598599999998	29.1658318	0.2985340812395414	0.037316760154942674
Parallel	8	4	0.5317035999999999	28.653636	29.185339399999997	0.29833453984091757	0.037291817480114696
Parallel	8	8	0.534897	28.8670814	29.401978200000002	0.2961363599677793	0.037017044995972415
Parallel	8	16	0.5385164	28.762609799999996	29.301126000000004	0.29715563831915537	0.03714445478989442

Tabella 14: Measurement with O0 optimization

4.4.4 10000 vertices - Special Cases - 0% of density

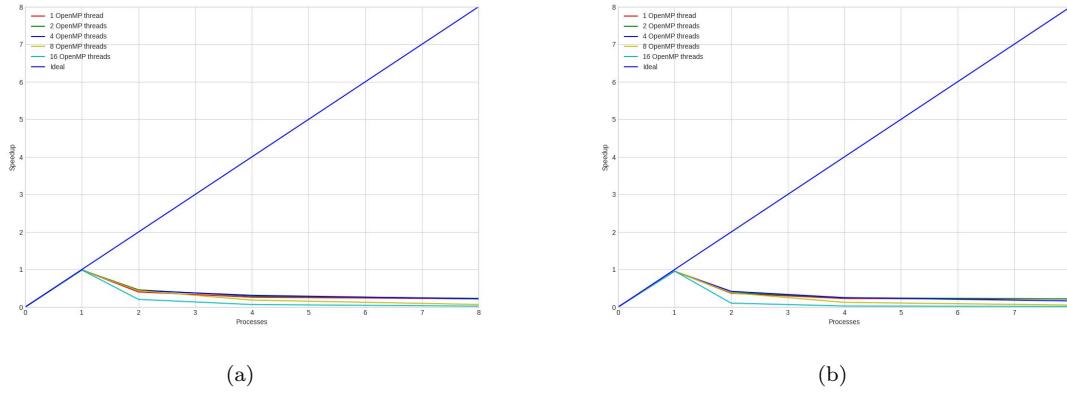


Figura 31: (a) Speedup with no optimization - (b) Speedup with *O1* optimization

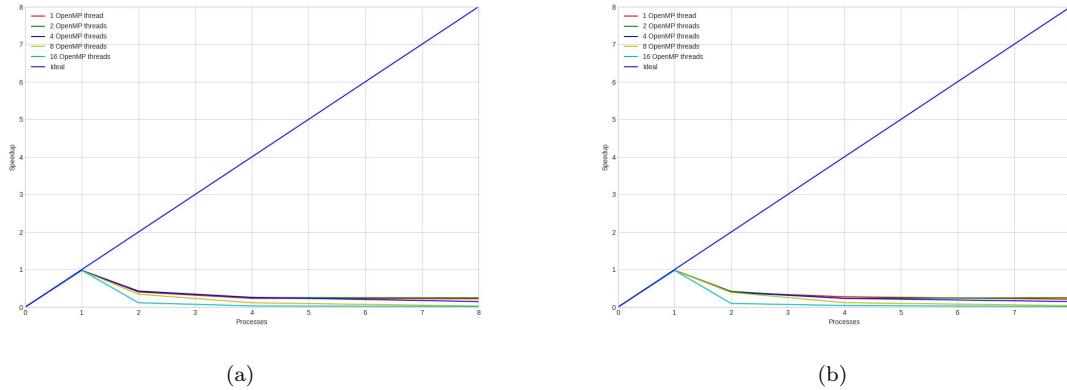


Figura 32: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Tarjan	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.0019538	0.0	0.0019538	1.0	1.0
Parallel	1	1	0.0019708000000000004	0.0	0.001970800000000004	0.9913740612949054	0.9913740612949054
Parallel	1	2	0.0019758	0.0	0.0019758	0.988865269764146	0.988865269764146
Parallel	1	4	0.0019784	0.0	0.0019784	0.9875657096643752	0.9875657096643752
Parallel	1	8	0.0019844	0.0	0.0019844	0.9845797218302762	0.9845797218302762
Parallel	1	16	0.0019768	0.0	0.0019768	0.9883650343990287	0.9883650343990287
Parallel	2	1	0.003304400000000003	0.0015406	0.00448444	0.4033110395508216	0.2016555197754108
Parallel	2	2	0.0031212	0.0011252	0.0042458	0.4601724056714871	0.23008620283574355
Parallel	2	4	0.0033998	0.0010236	0.004423199999999995	0.4417160441130946	0.2208582022065473
Parallel	2	8	0.0034608	0.0009824	0.00444300000000001	0.43974791807337377	0.21987395903668688
Parallel	2	16	0.0035496	0.0059520000000001	0.00950120000000001	0.20563718267166248	0.10281859133583124
Parallel	4	1	0.004377799999999999	0.0030162	0.00739420000000001	0.2642341294528143	0.06605853236320358
Parallel	4	2	0.004398	0.002373799999999997	0.0067716	0.288528560458385	0.07213214011459625
Parallel	4	4	0.003975	0.002245400000000003	0.00622039999999999	0.3140955655584853	0.07852388913896213
Parallel	4	8	0.0046246	0.0058328	0.01045760000000001	0.18683063035495714	0.046707657588739286
Parallel	4	16	0.0051588	0.0235674	0.0287262	0.0680145651008487	0.017003641275212176
Parallel	8	1	0.0050378	0.003803399999999997	0.00884099999999998	0.22099310032801722	0.027624137541002153
Parallel	8	2	0.004713400000000001	0.003633199999999996	0.00834679999999998	0.23407773038769353	0.02925971629846169
Parallel	8	4	0.004716	0.0042224	0.00893819999999999	0.218589726813005	0.027323734085162563
Parallel	8	8	0.0044928	0.0235066	0.0279994	0.0697800667151534	0.008722508339464417
Parallel	8	16	0.0048454	0.0709326	0.075778	0.025783208846894876	0.0032229011058618594

Tabella 15: Measurement with O0 optimization

4.4.5 10000 vertices - Special Cases - 100% of density

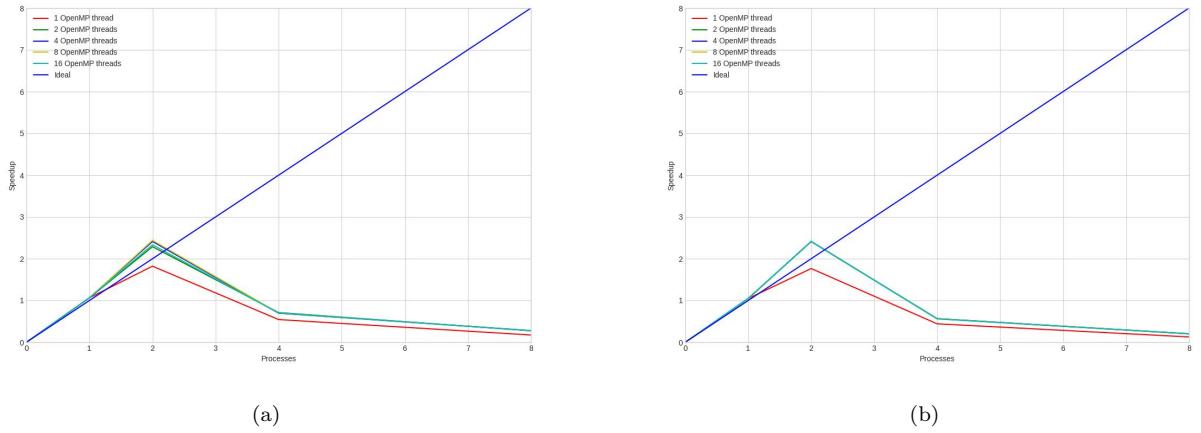


Figura 33: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

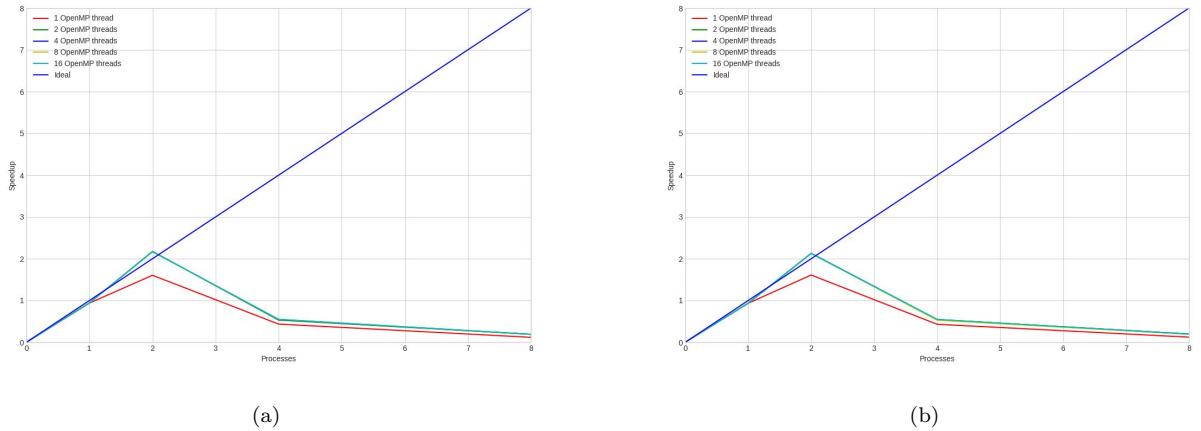


Figura 34: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Tarjan	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	7.154188599999995	0.0	7.154188599999995	1.0	1.0
Parallel	1	1	6.8340534	0.0	6.8340534	1.0468441174310987	1.0468441174310987
Parallel	1	2	6.8174104	0.0	6.8174104	1.0493997251507698	1.0493997251507698
Parallel	1	4	6.846560800000001	0.0	6.846560800000001	1.0449317268897982	1.0449317268897982
Parallel	1	8	6.881190599999999	0.0	6.881190599999999	1.0396730763423412	1.0396730763423412
Parallel	1	16	6.819358800000001	0.0	6.819358800000001	1.0490998948464185	1.0490998948464185
Parallel	2	1	1.869538	2.183862399999998	4.053400399999999	1.7649844313431264	0.8824922156715632
Parallel	2	2	1.8661496	1.098324599999998	2.9644742	2.413307762975302	1.206653881487651
Parallel	2	4	1.868193399999998	1.0978906	2.9660842	2.4119978117950933	1.2059989058975467
Parallel	2	8	1.8709226	1.0960744	2.9669972000000002	2.4112555935878214	1.2056277976939107
Parallel	2	16	1.8695856	1.0993586	2.9689438	2.409674645912799	1.2048373229563996
Parallel	4	1	0.6471376	15.5539132	16.201051200000002	0.4415879260970423	0.11039698152426057
Parallel	4	2	0.6464604	12.0673058	12.713766	0.5627119926542615	0.14067799816356538
Parallel	4	4	0.6509688	11.957168	12.6081368	0.5674263147271689	0.14185657868179222
Parallel	4	8	0.6478364000000001	11.98399579999999	12.631832	0.566361918049575	0.14159047951239376
Parallel	4	16	0.6472222	12.15159939999999	12.7988216	0.5589724447756971	0.13974311119392427
Parallel	8	1	0.2934982	55.7621624	56.05566040000001	0.12762651530549088	0.01595331441318636
Parallel	8	2	0.2933952	34.64820659999995	34.941602	0.20474701188571717	0.025593376485714647
Parallel	8	4	0.298612	35.0996636	35.3982758	0.2021055669609761	0.02526319587012201
Parallel	8	8	0.2993162	34.5301598	34.829476	0.20540615081317903	0.02567576885164738
Parallel	8	16	0.308752399999999	35.0826502	35.391402400000004	0.20214481808723123	0.025268102260903903

Tabella 16: Measurement with $O1$ optimization

5 Kosaraju's algorithm

Using the same parallelization method, we decide to parallelize one of the other proposed algorithm for finding the strongly connected components, and in particular the selected algorithm is the Kosaraju's algorithm.

Kosaraju's algorithm is a linear time algorithm, that makes use of the fact that the transpose graph (the same graph with the direction of every edge reversed) has exactly the same strongly connected components as the original graph. In this case, where strongly connected components are to be represented by appointing a separate root vertex for each component, and assigning to each vertex the root vertex of its component, then Kosaraju's algorithm can be stated as follows:

1. perform a depth-first search (DFS) on the graph, marking the finish times of each vertex as it is visited. This is done by starting at an arbitrary vertex and visiting all its unvisited neighbors recursively. The DFS should be done on the original graph, not its transpose
2. create a new graph that is the transpose of the original graph, where all edges are reversed. Perform another DFS on this new graph, but this time visit the vertices in decreasing order of their finish times from the first DFS. When visiting a vertex, add it to the current strongly connected component
3. repeat step 2 for all vertices that have not been visited yet
4. the strongly connected components are the subgraphs of connected vertices that have been found

Kosaraju's algorithm is efficient because it takes advantage of the fact that the strongly connected components in a graph are the strongly connected components in the graph's transpose. The first pass of DFS finds all the strongly connected components in the transpose of the original graph. The second pass of DFS uses the order of the stack from the first pass to find the strongly connected components in the original graph.

5.1 Kosaraju and Tarjan algorithm

The main difference between the two algorithms is how they find the strongly connected components.

Kosaraju's algorithm works as mentioned above. Tarjan's algorithm, on the other hand, uses a single depth-first search (DFS) to find the strongly connected components. It uses a stack to keep track of the nodes being visited and a low-link value for each node to determine whether a component has been found. The low-link value of a node is the smallest index of any node known to be reachable from that node, including the node itself.

Both algorithms are linear time complexity and have $O(V+E)$ space complexity, where V is the number of nodes and E is the number of edges in the graph.

Another difference is that Kosaraju's algorithm is a two-pass algorithm and Tarjan's algorithm is a single-pass algorithm, which means Kosaraju's algorithm needs more memory space than Tarjan's algorithm. Furthermore, Kosaraju's algorithm is also more robust to graphs with negative cycles, while Tarjan's algorithm can fail in those cases.

5.2 Kosaraju's parallel version analysis

Also for Kosaraju, as done for the Tarjan algorithm, the analysis was conducted following the same structuring as the previously proposed case studies. So, subsequent analyses will be shown referring to the structuring given previously to those of Tarjan's algorithm. Therefore, we will also have here the 4 proposed case studies, which will now be presented.

5.3 Case Study 1

5.3.1 500 vertices - 25% of density

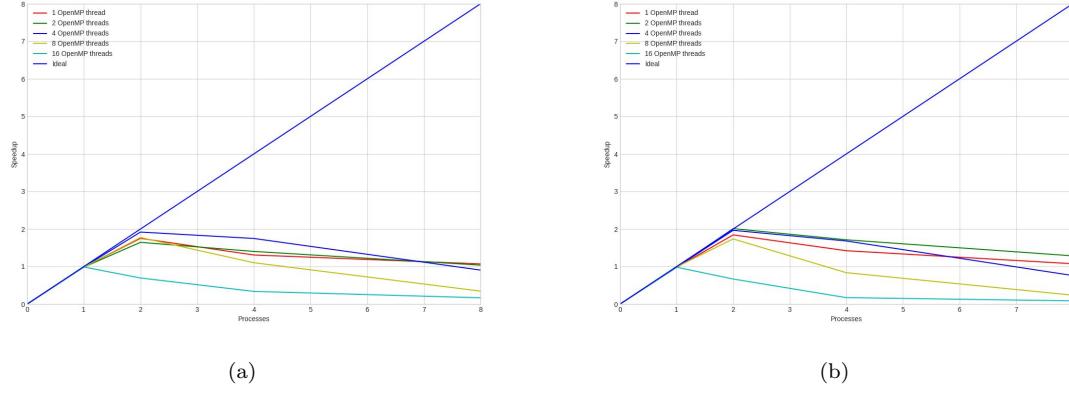


Figura 35: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

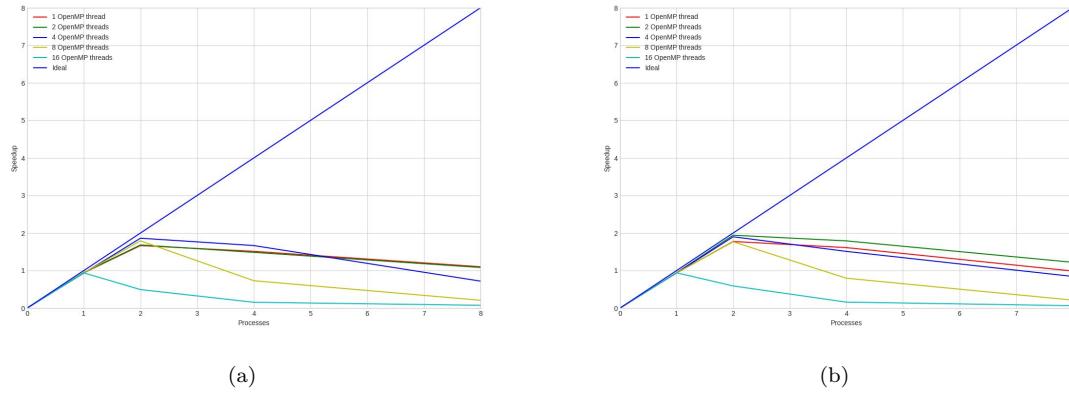


Figura 36: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Kosaraju	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.0060886000000000004	0.0	0.0060886000000000004	1.0	1.0
Parallel	1	1	0.0062234000000000005	0.0	0.0062234000000000005	0.9783398142494456	0.9783398142494456
Parallel	1	2	0.0061952	0.0	0.0061952	0.9827931301652894	0.9827931301652894
Parallel	1	4	0.006185	0.0	0.006185	0.9844139046079224	0.9844139046079224
Parallel	1	8	0.0061470000000000006	0.0	0.0061470000000000006	0.9904994306165609	0.9904994306165609
Parallel	1	16	0.006196599999999999	0.0	0.006196599999999999	0.9825710873704937	0.9825710873704937
Parallel	2	1	0.0025136	0.0007882	0.003302199999999995	1.84380110229544	0.92190055114772
Parallel	2	2	0.002518999999999995	0.0005088	0.0030278	2.010899002576128	1.005449501288064
Parallel	2	4	0.0025296	0.000568	0.003097600000000003	1.9655862603305785	0.9827931301652892
Parallel	2	8	0.002648	0.0008596	0.003508	1.7356328392246296	0.8678164196123148
Parallel	2	16	0.002639	0.00647279999999999	0.0091118	0.6682104523804299	0.33410522619021493
Parallel	4	1	0.0015974	0.002684200000000003	0.0042816	1.422038490284006	0.3555096225710015
Parallel	4	2	0.001430399999999999	0.002124	0.0035546	1.712879086254431	0.42821977156360774
Parallel	4	4	0.001599199999999998	0.002022	0.003621400000000003	1.6812834815264814	0.42032087038162036
Parallel	4	8	0.001669199999999998	0.005604599999999995	0.007273799999999995	0.8370590337925157	0.20926475844812892
Parallel	4	16	0.001569600000000002	0.03377599999999994	0.0353458	0.17225809007010737	0.04306452251752684
Parallel	8	1	0.0009876	0.0046728	0.0056602	1.0756863715063074	0.13446079643828843
Parallel	8	2	0.0009914	0.00374339999999996	0.0047348	1.285925487876996	0.1607406859846245
Parallel	8	4	0.001052399999999998	0.00694039999999999	0.00799259999999999	0.7617796461727099	0.09522245577158873
Parallel	8	8	0.001022199999999998	0.0244956	0.0255178	0.23860207384649149	0.029825259230811436
Parallel	8	16	0.0009768	0.0692841999999999	0.0702609999999999	0.08665689358249956	0.010832111697812445

Tabella 17: Measurement with $O1$ optimization

5.3.2 500 vertices - 50% of density

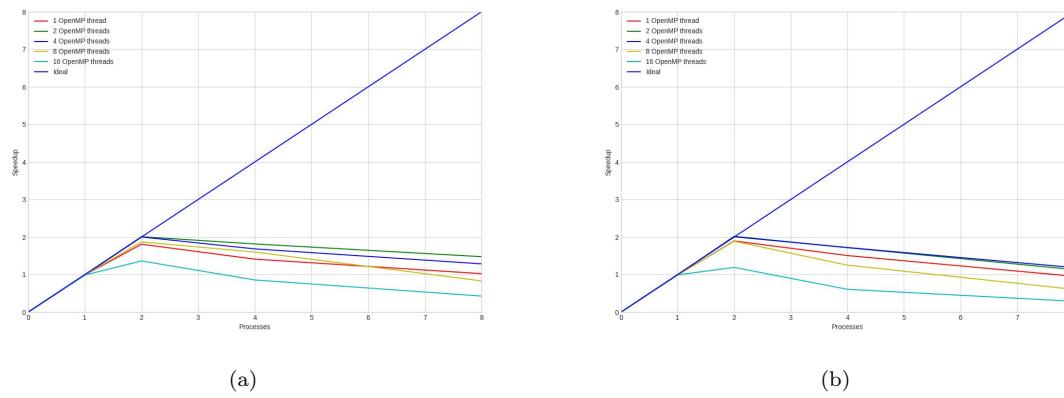


Figura 37: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

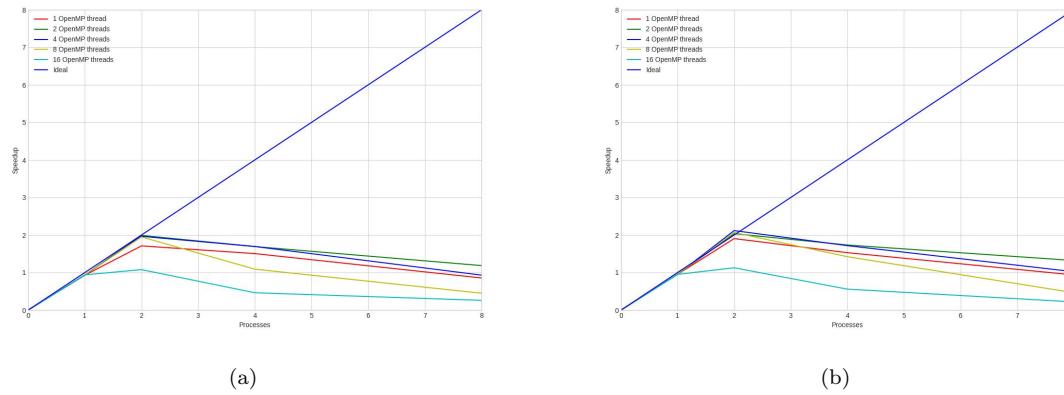


Figura 38: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Kosaraju	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.0314408	0.0	0.0314408	1.0	1.0
Parallel	1	1	0.03183719999999996	0.0	0.03183719999999996	0.9875491563328433	0.9875491563328433
Parallel	1	2	0.03184299999999996	0.0	0.03184299999999996	0.9873692805326132	0.9873692805326132
Parallel	1	4	0.0318494	0.0	0.0318494	0.987170872920683	0.987170872920683
Parallel	1	8	0.031863800000000005	0.0	0.031863800000000005	0.9867247472052924	0.9867247472052924
Parallel	1	16	0.031850600000000001	0.0	0.031850600000000001	0.9871336803702282	0.9871336803702282
Parallel	2	1	0.0136764	0.003710200000000003	0.0173866	1.8083351546593354	0.9041675773296677
Parallel	2	2	0.013685600000000001	0.001970200000000005	0.0156556	2.0082781879966274	1.0041390939983137
Parallel	2	4	0.0136778	0.002031200000000004	0.0157088	2.0014768792014666	1.0007384396007333
Parallel	2	8	0.01452539999999999	0.002289800000000003	0.01681520000000002	1.8697844807079305	0.9348922403539652
Parallel	2	16	0.01431659999999999	0.0087574	0.02307399999999997	1.362607263586721	0.6813036317933605
Parallel	4	1	0.0090662	0.0132152	0.0222812	1.4110909645800045	0.3527727411450011
Parallel	4	2	0.00827339999999998	0.0090424	0.01731599999999998	1.8157080157080159	0.45392700392700397
Parallel	4	4	0.009501800000000001	0.00917059999999998	0.0186724	1.683811400784045	0.42095285019601125
Parallel	4	8	0.008631400000000001	0.0110262	0.019657600000000004	1.59942210646264	0.39985552661566
Parallel	4	16	0.0086538	0.0280568	0.0367108	0.8564455146714317	0.21411137866785793
Parallel	8	1	0.005132400000000005	0.0255768	0.03070899999999997	1.0238301475137581	0.12797876843921976
Parallel	8	2	0.0047246	0.01656319999999997	0.021287400000000005	1.4769675958548245	0.18462094948185306
Parallel	8	4	0.0055	0.018961600000000002	0.0244618	1.2853019810480013	0.16066274763100016
Parallel	8	8	0.005852	0.031959400000000006	0.03781140000000001	0.8315164209735686	0.1039395262169607
Parallel	8	16	0.0048398	0.0688247999999999	0.0736646	0.4268101639050507	0.05335127048813134

Tabella 18: Measurement with O0 optimization

5.3.3 500 vertices - 75% of density

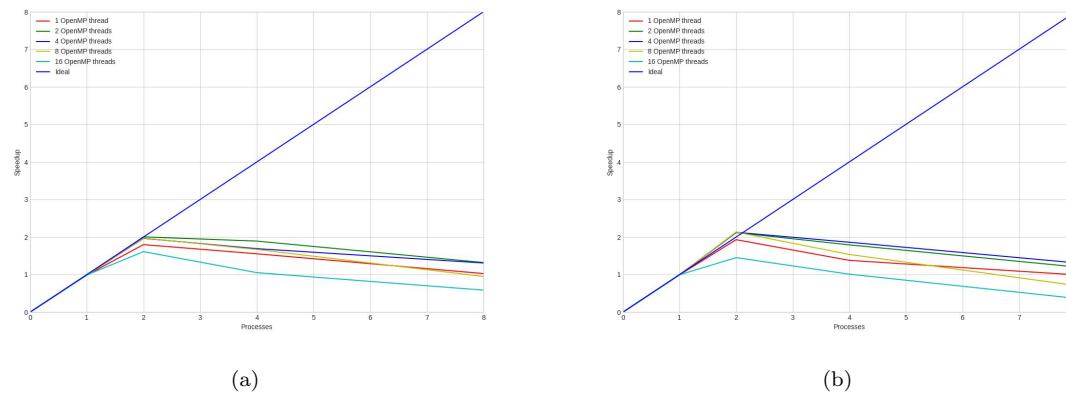


Figura 39: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

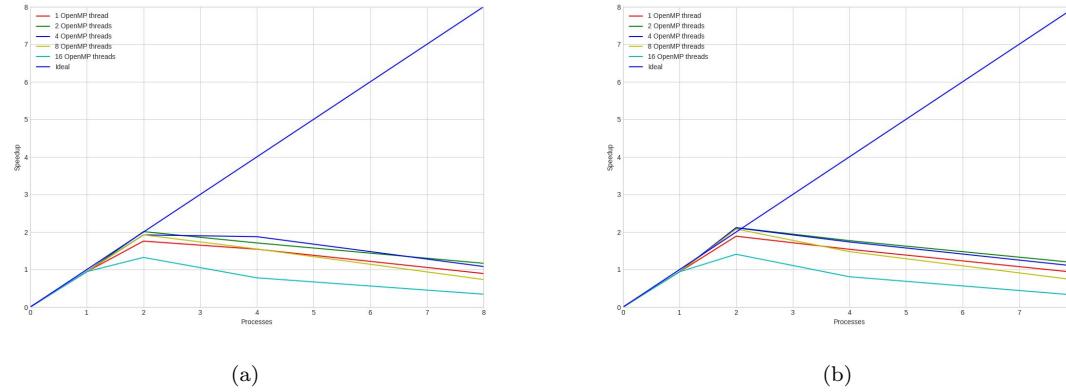


Figura 40: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Kosaraju	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.0285124	0.0	0.0285124	1.0	1.0
Parallel	1	1	0.028688200000000004	0.0	0.028688200000000004	0.9938720449522799	0.9938720449522799
Parallel	1	2	0.028771800000000004	0.0	0.028771800000000004	0.9909842276117586	0.9909842276117586
Parallel	1	4	0.0286378	0.0	0.0286378	0.9956211720174035	0.9956211720174035
Parallel	1	8	0.02872899999999998	0.0	0.02872899999999998	0.9924605799018414	0.9924605799018414
Parallel	1	16	0.0286668	0.0	0.0286668	0.9946139785396347	0.9946139785396347
Parallel	2	1	0.011406	0.003373	0.014779200000000001	1.9292248565551584	0.9646124282775792
Parallel	2	2	0.0116682	0.0018	0.0134682	2.1170163793231462	1.0585081896615731
Parallel	2	4	0.01156559999999999	0.0018482	0.01341399999999999	2.1255702996868946	1.0627851498434473
Parallel	2	8	0.0114146	0.001961	0.01337559999999998	2.131672597864769	1.0658362989323844
Parallel	2	16	0.01130859999999999	0.0083376	0.019646000000000004	1.4513081543316704	0.7256540771658352
Parallel	4	1	0.007363	0.01332319999999999	0.020686	1.3783428405685005	0.34458571014212513
Parallel	4	2	0.006093400000000006	0.009825200000000001	0.01591859999999998	1.7911374115814207	0.4477843528953552
Parallel	4	4	0.005730600000000001	0.00958919999999999	0.015319600000000003	1.8611713099558733	0.46529282748896833
Parallel	4	8	0.006486000000000004	0.012080200000000003	0.01856639999999997	1.5356988969320926	0.38392472423302315
Parallel	4	16	0.006234200000000005	0.0219888	0.0282228	1.0102612072508752	0.2525653018127188
Parallel	8	1	0.003354200000000003	0.02535860000000002	0.028713	0.9930136175251628	0.12412670219064535
Parallel	8	2	0.004	0.0196726	0.0236724	1.2044575117013907	0.15055718896267384
Parallel	8	4	0.003368600000000002	0.01836519999999998	0.02173379999999998	1.3118920759370198	0.16398650949212748
Parallel	8	8	0.0036606	0.0365864	0.04024679999999999	0.7084389317908506	0.08855486647385633
Parallel	8	16	0.004942	0.0728564	0.07779820000000001	0.3664917697324616	0.0458114712165577

Tabella 19: Measurement with $O1$ optimization

5.4 Case Study 2

5.4.1 1000 vertices - 25% of density

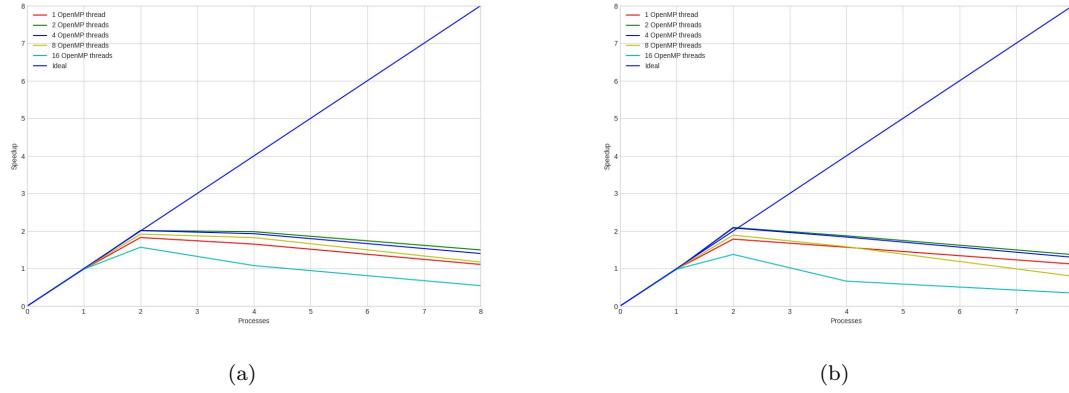


Figura 41: (a) Speedup with no optimization - (b) Speedup with *O1* optimization

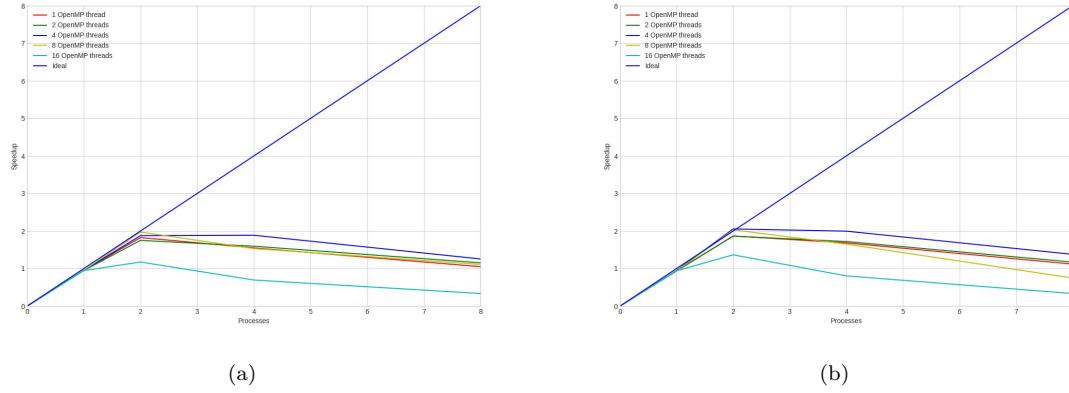


Figura 42: (a) Speedup with *O2* optimization - (b) Speedup with *O3* optimization

Version	Processes	OMP	Kosaraju	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.02550999999999998	0.0	0.02550999999999998	1.0	1.0
Parallel	1	1	0.0258718	0.0	0.0258718	0.9860156618403048	0.9860156618403048
Parallel	1	2	0.0258802	0.0	0.0258802	0.9856956283181737	0.9856956283181737
Parallel	1	4	0.0259158	0.0	0.0259158	0.9843415985614953	0.9843415985614953
Parallel	1	8	0.0259386	0.0	0.0259386	0.9834763634120577	0.9834763634120577
Parallel	1	16	0.025985600000000005	0.0	0.025985600000000005	0.9816975555692381	0.9816975555692381
Parallel	2	1	0.0113484	0.002946	0.01429419999999998	1.7846399238852122	0.8923199619426061
Parallel	2	2	0.0105888	0.0015926	0.01218139999999998	2.094176367248428	1.047088183624214
Parallel	2	4	0.0106054	0.001649	0.0122546	2.0816672922820816	1.0408336461410408
Parallel	2	8	0.0111108	0.002374000000000002	0.01348439999999999	1.8918157278039809	0.9459078639019904
Parallel	2	16	0.01113399999999998	0.0073924	0.018525999999999997	1.376983698587715	0.688491849298858
Parallel	4	1	0.0061832	0.0100936	0.0162768	1.5672613781578686	0.39181534453946715
Parallel	4	2	0.00616139999999999	0.0074576	0.013619000000000001	1.8731184374770538	0.46827960936926344
Parallel	4	4	0.0064834	0.0073812	0.0138646	1.8399376830200653	0.45998442075501633
Parallel	4	8	0.0060988	0.01002939999999999	0.016128200000000002	1.5817016158033752	0.3954254039508438
Parallel	4	16	0.0057426	0.03263319999999994	0.038376	0.6647383781530123	0.16618459453825307
Parallel	8	1	0.00360359999999998	0.019167	0.02277099999999996	1.1202845724825437	0.14003557156031796
Parallel	8	2	0.0039162	0.01467419999999998	0.018590600000000002	1.372198853183867	0.17152485664798336
Parallel	8	4	0.00368519999999996	0.0159148	0.01960019999999998	1.3015173314558015	0.1626896664319752
Parallel	8	8	0.0042654	0.02777880000000003	0.03204419999999995	0.7960879035831758	0.09951098794789698
Parallel	8	16	0.0047254	0.0682765999999999	0.0730018	0.34944343838097136	0.04368042979762142

Tabella 20: Measurement with *O1* optimization

5.4.2 1000 vertices - 50% of density

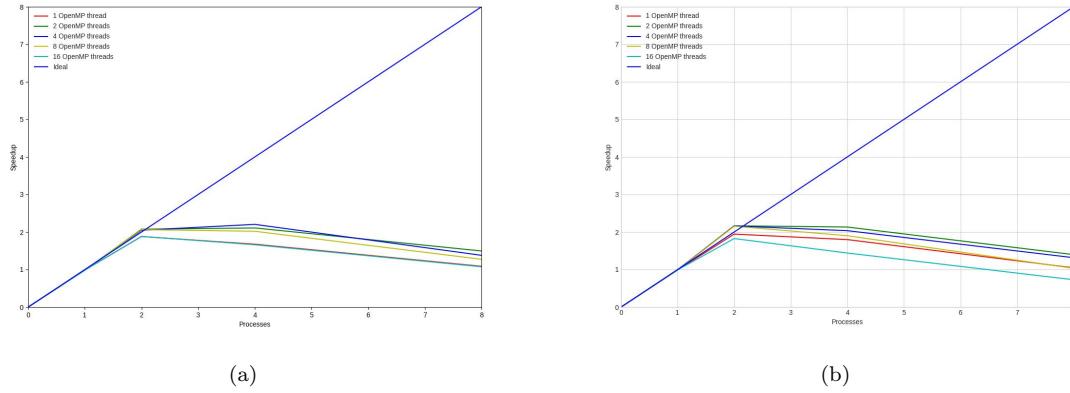


Figura 43: (a) Speedup with no optimization - (b) Speedup with *O1* optimization

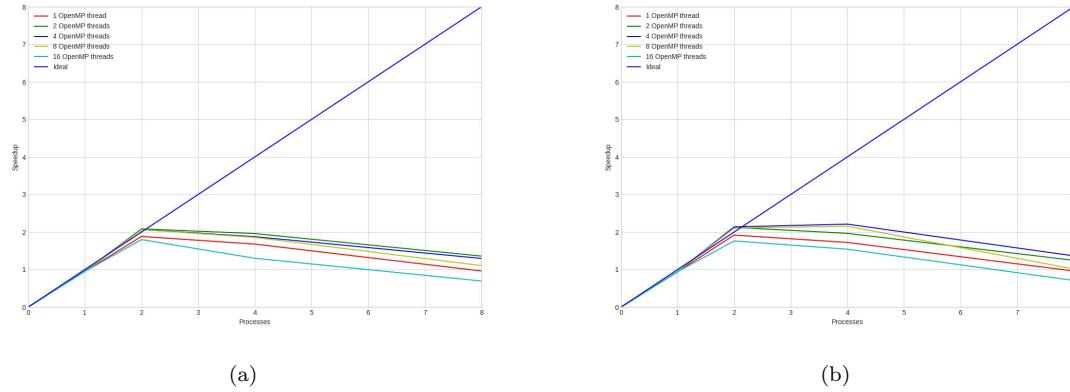


Figura 44: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Kosaraju	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.07762240000000001	0.0	0.07762240000000001	1.0	1.0
Parallel	1	1	0.07864420000000001	0.0	0.07864420000000001	0.9870073063239246	0.9870073063239246
Parallel	1	2	0.07837000000000001	0.0	0.07837000000000001	0.9904606354472375	0.9904606354472375
Parallel	1	4	0.0782384	0.0	0.0782384	0.9921266283564082	0.9921266283564082
Parallel	1	8	0.07872660000000001	0.0	0.07872660000000001	0.9859742450455119	0.9859742450455119
Parallel	1	16	0.078431	0.0	0.078431	0.98969030102893	0.98969030102893
Parallel	2	1	0.03186940000000006	0.00804779999999999	0.0399172	1.944585291553516	0.972292645776758
Parallel	2	2	0.031666	0.0041818	0.0358478	2.165332321648749	1.0826661608243744
Parallel	2	4	0.03168220000000001	0.004228200000000005	0.03591039999999995	2.1615576546070225	1.0807788273035113
Parallel	2	8	0.03170920000000001	0.00432899999999995	0.03603860000000004	2.1538683522667363	1.0769341761333682
Parallel	2	16	0.0317238	0.010787	0.04251120000000006	1.8259282259733904	0.9129641129866952
Parallel	4	1	0.01534100000000002	0.02782759999999997	0.04316820000000004	1.7981384445031297	0.4495346111257824
Parallel	4	2	0.0142398	0.0221474	0.03638699999999996	2.1332453898370303	0.5333113474592576
Parallel	4	4	0.0154258	0.0227044	0.03813019999999996	2.035719718228596	0.508929929557149
Parallel	4	8	0.0150528	0.025766	0.04081860000000003	1.9016428784916681	0.47541071962291703
Parallel	4	16	0.0154336	0.0384456	0.0538794	1.4406693467262073	0.3601673366815518
Parallel	8	1	0.00803020000000001	0.0661454	0.0741756	1.0464681108073277	0.1308051385091596
Parallel	8	2	0.009586	0.0459496	0.0555358	1.397700222199014	0.1747125277487675
Parallel	8	4	0.010313	0.04868680000000006	0.0589815999999999	1.3160443256880117	0.16450554071100146
Parallel	8	8	0.0103296	0.06484340000000001	0.075173	1.0325835073763188	0.12907293842203985
Parallel	8	16	0.01022019999999999	0.0963814	0.1066018	0.7281528079263203	0.09101910099079004

Tabella 21: Measurement with O1 optimization

5.4.3 1000 vertices - 75% of density

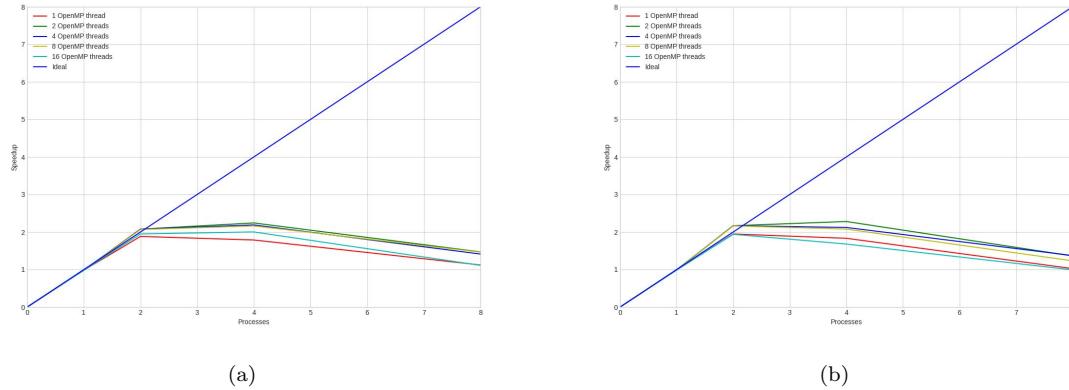


Figura 45: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

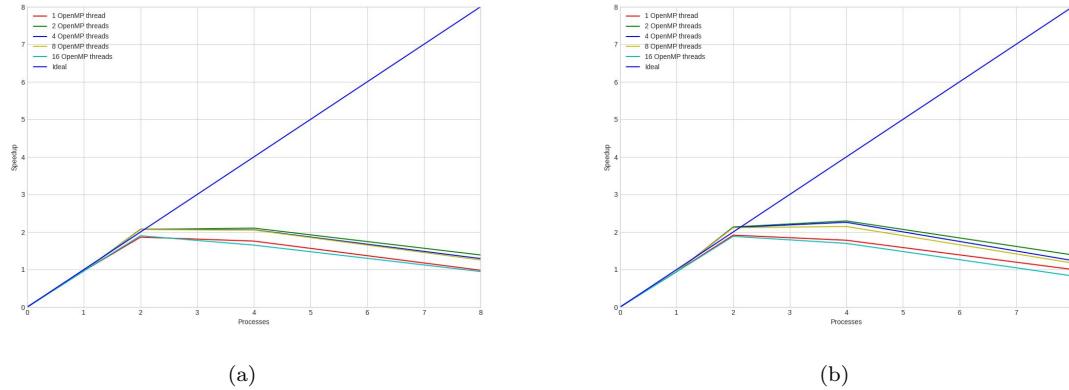


Figura 46: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Kosaraju	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.2386202	0.0	0.2386202	1.0	1.0
Parallel	1	1	0.2427738	0.0	0.2427738	0.982891069794187	0.982891069794187
Parallel	1	2	0.24257420000000002	0.0	0.24257420000000002	0.9836998328758787	0.9836998328758787
Parallel	1	4	0.24308760000000001	0.0	0.24308760000000001	0.9816222629208565	0.9816222629208565
Parallel	1	8	0.2428168	0.0	0.2428168	0.9827170113435314	0.9827170113435314
Parallel	1	16	0.24290859999999997	0.0	0.24290859999999997	0.982345623003879	0.982345623003879
Parallel	2	1	0.1029105999999999	0.0238158	0.1267262	1.8829586936245226	0.9414793468122613
Parallel	2	2	0.10277500000000002	0.0121438	0.114919	2.076420783343487	1.0382103916671743
Parallel	2	4	0.1027398	0.0121748	0.11491480000000001	2.0764966740576494	1.0382483370288247
Parallel	2	8	0.10272060000000001	0.01270339999999998	0.1154238	2.067339664783173	1.0336698323915865
Parallel	2	16	0.1026314	0.0197984	0.1224296	1.9490401014133838	0.9745200507066919
Parallel	4	1	0.052272	0.0812218	0.133494	1.7874975654336525	0.4468743913584131
Parallel	4	2	0.05230040000000004	0.0541194	0.10642	2.242249577147153	0.5605623942867882
Parallel	4	4	0.05218120000000004	0.05682280000000001	0.1090039999999999	2.1890958129976883	0.5472739532494221
Parallel	4	8	0.0524529999999999	0.0579074	0.1103601999999999	2.162194341800758	0.5405485854501895
Parallel	4	16	0.05283780000000004	0.066348	0.1191857999999998	2.0020858189482307	0.5005214547370577
Parallel	8	1	0.0329756	0.17973060000000002	0.2127062	1.1218300171786248	0.1402287521473281
Parallel	8	2	0.0377181999999999	0.1245855999999999	0.1623038	1.4702071054405381	0.18377588818006727
Parallel	8	4	0.0385352	0.1303324	0.1688674	1.4130625567753161	0.17663281959691451
Parallel	8	8	0.03250459999999995	0.1304674	0.1629719999999998	1.4641791227941	0.1830223903492625
Parallel	8	16	0.0330037999999999	0.1816046	0.2146081999999997	1.1118876165961973	0.13898595207452466

Tabella 22: Measurement with O0 optimization

5.4.4 1000 vertices - Special Cases - 0%

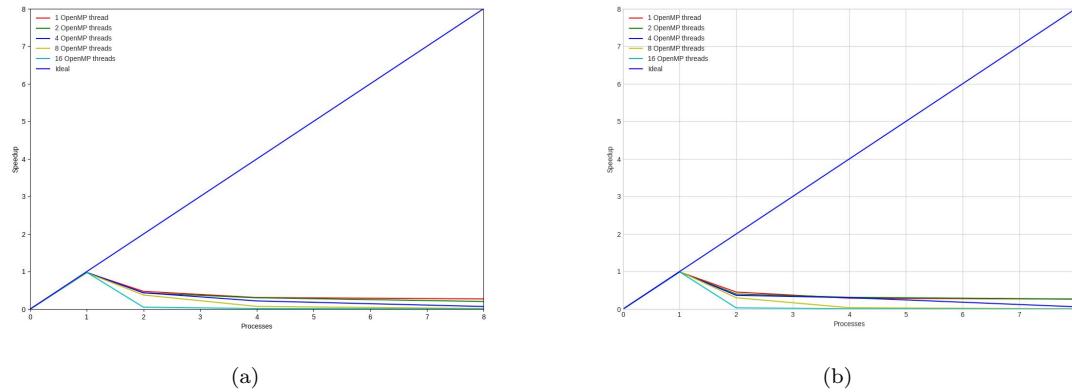


Figura 47: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

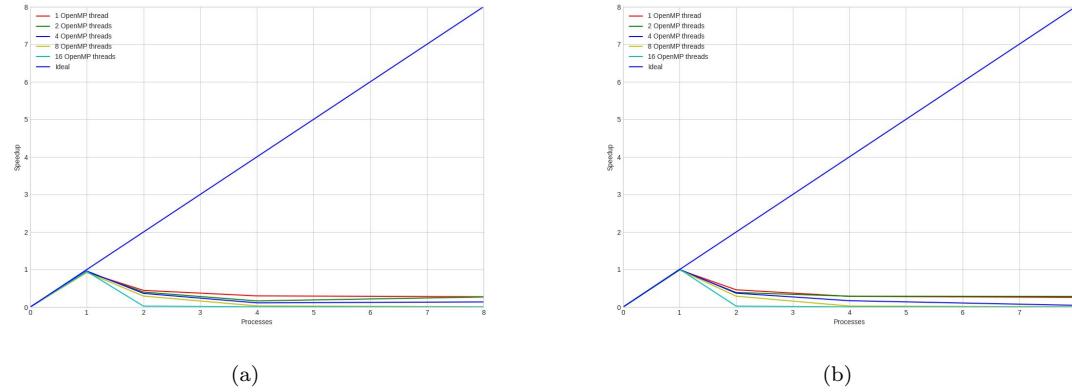


Figura 48: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Kosaraju	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.000192	0.0	0.000192	1.0	1.0
Parallel	1	1	0.0001936	0.0	0.0001936	0.9917355371900827	0.9917355371900827
Parallel	1	2	0.0001944	0.0	0.0001944	0.9876543209876543	0.9876543209876543
Parallel	1	4	0.0001909999999999998	0.0	0.0001909999999999998	1.005235602094241	1.00523560209424
Parallel	1	8	0.0001934	0.0	0.0001934	0.9927611168562565	0.9927611168562565
Parallel	1	16	0.00019020000000000002	0.0	0.00019020000000000002	1.0094637223974763	1.0094637223974763
Parallel	2	1	0.0003372	7.92e-05	0.0004168	0.46065259117082535	0.23032629558541268
Parallel	2	2	0.000363	0.00012880000000000001	0.0004921999999999999	0.39008533116619265	0.19504266558309633
Parallel	2	4	0.0003464	0.00017040000000000002	0.0005170000000000001	0.37137330754352027	0.18568665377176014
Parallel	2	8	0.0003788	0.0002805999999999994	0.0006594	0.29117379435850776	0.14558689717925388
Parallel	2	16	0.0003683999999999996	0.0071778	0.00754619999999999	0.025443269460125633	0.012721634730062817
Parallel	4	1	0.0005068	0.00016340000000000001	0.0006704	0.2863961813842482	0.07159904534606205
Parallel	4	2	0.000492199999999999	0.0001686	0.0006606	0.29064486830154407	0.07266121707538602
Parallel	4	4	0.0004654000000000004	0.0006612	0.0011268	0.1703940362087327	0.042598509052183174
Parallel	4	8	0.0004562000000000003	0.00647699999999999	0.006933200000000005	0.027692840246927825	0.006923210061731956
Parallel	4	16	0.0005022	0.034721	0.0352234	0.005450921830374126	0.0013627304575935316
Parallel	8	1	0.0005362	0.0002056	0.000741799999999999	0.2588298732812079	0.032353734160150985
Parallel	8	2	0.0004816	0.0001944	0.000676000000000001	0.28402366863905326	0.03550295857988166
Parallel	8	4	0.0005158000000000001	0.003543599999999996	0.0040592	0.047299960583366184	0.005912495072920773
Parallel	8	8	0.000475	0.0271712	0.02764639999999998	0.006944846345274611	0.0008681057931593264
Parallel	8	16	0.0004778	0.0603056	0.06078359999999999	0.003158746767220106	0.00039484334590251326

Tabella 23: Measurement with $O1$ optimization

5.4.5 1000 vertices - Special Cases - 100%

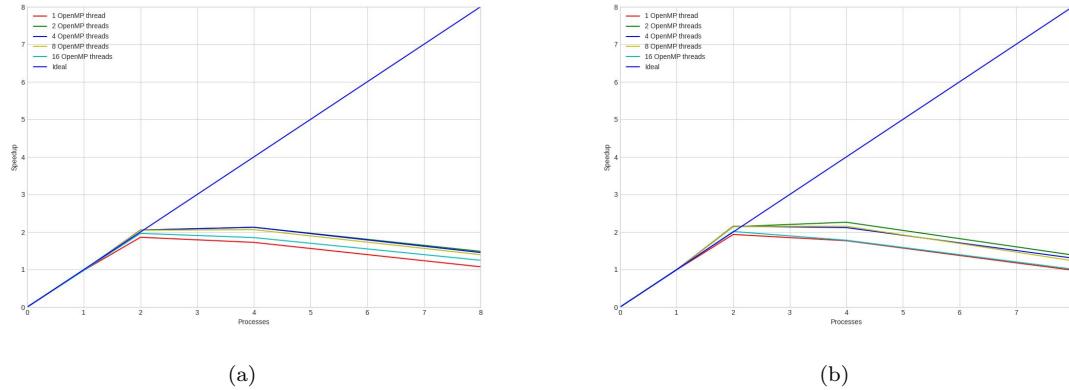


Figura 49: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

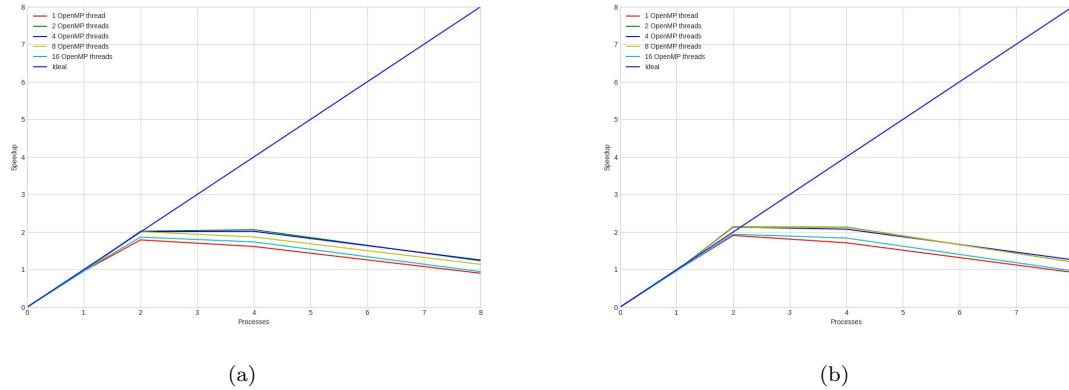


Figura 50: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Kosaraju	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.20019	0.0	0.20019	1.0	1.0
Parallel	1	1	0.20231719999999997	0.0	0.20231719999999997	0.9894858173205245	0.9894858173205245
Parallel	1	2	0.201967	0.0	0.201967	0.9912015329236955	0.9912015329236955
Parallel	1	4	0.2016072	0.0	0.2016072	0.9929704891491972	0.9929704891491972
Parallel	1	8	0.2019534	0.0	0.2019534	0.9912682826830348	0.9912682826830348
Parallel	1	16	0.20206280000000004	0.0	0.20206280000000004	0.9907315943360181	0.9907315943360181
Parallel	2	1	0.0822112000000001	0.0212174	0.1034284	1.9355418821136168	0.9677709410568084
Parallel	2	2	0.0826932	0.01076899999999999	0.0934622000000001	2.141935456259322	1.070967728129661
Parallel	2	4	0.0820808	0.0108172	0.0928976	2.1549534110676705	1.0774767055338352
Parallel	2	8	0.0822054	0.0109280000000002	0.0931333999999999	2.149497387618191	1.0747486938090955
Parallel	2	16	0.0826302000000001	0.0167954	0.0994261999999999	2.0134531944296374	1.0067265972148187
Parallel	4	1	0.03675	0.076424	0.1131737999999999	1.7688723008328784	0.4422180752082196
Parallel	4	2	0.036967	0.0515588	0.0885256	2.261379759075341	0.5653449397688353
Parallel	4	4	0.0369992	0.05752740000000006	0.0945268	2.1178120913857237	0.5294530228464309
Parallel	4	8	0.0372234	0.0557446	0.09296800000000001	2.15332157301437	0.5383303932535926
Parallel	4	16	0.0373812	0.0750035999999999	0.1123847999999999	1.7812907083520193	0.4453226770880048
Parallel	8	1	0.0229383999999998	0.1806706000000001	0.2036089999999998	0.9832080114336794	0.12290100142920993
Parallel	8	2	0.0211616	0.1231074	0.1442686	1.3876200365152225	0.1734525045644028
Parallel	8	4	0.0221332	0.1311654	0.1532986	1.3058827673573015	0.16323534591966268
Parallel	8	8	0.0223742	0.1395198000000003	0.161894	1.236549841254154	0.15456873015676925
Parallel	8	16	0.0239474	0.1737928000000003	0.1977404	1.0123879591626193	0.1265484948953274

Tabella 24: Measurement with O1 optimization

5.5 Case Study 3

5.5.1 2500 vertices - 25% of density

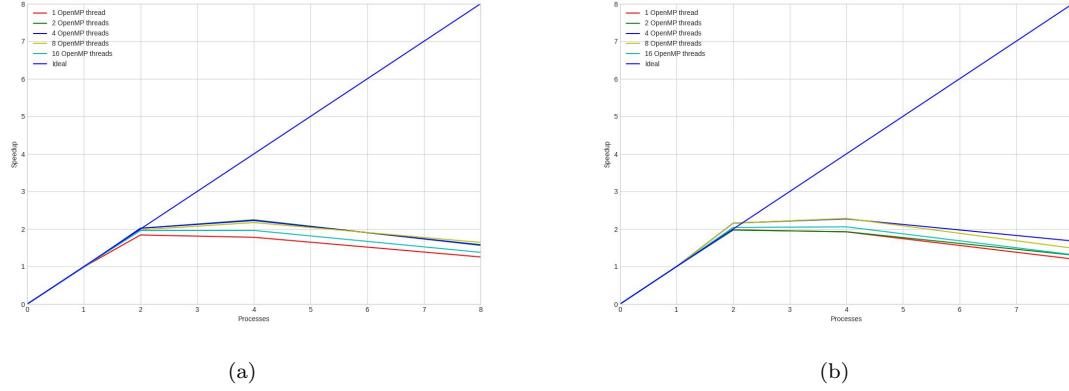


Figura 51: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

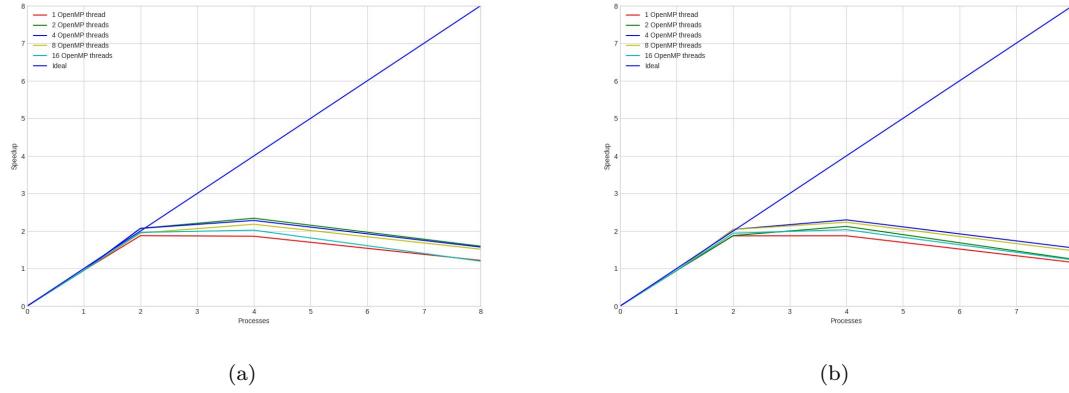


Figura 52: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Kosaraju	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.18749880000000002	0.0	0.18749880000000002	1.0	1.0
Parallel	1	1	0.18718059999999997	0.0	0.18718059999999997	1.0016999624961136	1.0016999624961136
Parallel	1	2	0.187158	0.0	0.187158	1.0018209213605618	1.0018209213605618
Parallel	1	4	0.1870102	0.0	0.1870102	1.002612691714142	1.002612691714142
Parallel	1	8	0.18710180000000004	0.0	0.18710180000000004	1.0021218395547236	1.0021218395547236
Parallel	1	16	0.18725679999999997	0.0	0.18725679999999997	1.0012923429215925	1.0012923429215925
Parallel	2	1	0.07759	0.0171842	0.0947742	1.9783738612407176	0.9891869306203588
Parallel	2	2	0.0778402	0.0172546	0.09509480000000001	1.971704025877335	0.9858520129386675
Parallel	2	4	0.07800860000000001	0.0089186	0.08692720000000001	2.156963528101676	1.078481764050838
Parallel	2	8	0.0777244	0.0091698	0.0868944	2.157777716400597	1.078888582002985
Parallel	2	16	0.0777668	0.01411359999999999	0.0918804	2.0406833231026424	1.0203416615513212
Parallel	4	1	0.0376676	0.05976240000000001	0.0974299999999999	1.9244462691162891	0.4811115672790723
Parallel	4	2	0.03736100000000005	0.0598828	0.09724400000000001	1.9281271852247954	0.48203179630619886
Parallel	4	4	0.0376352	0.0449994	0.08263480000000001	2.2690053101114787	0.5672513275278697
Parallel	4	8	0.037739	0.0443926	0.0821316	2.282906944464737	0.5707267361161843
Parallel	4	16	0.0389298	0.0521706	0.0911006	2.0581510988950678	0.5145377747237669
Parallel	8	1	0.0234846	0.1325586	0.1560434	1.20158109851802	0.15019763732397526
Parallel	8	2	0.0217792	0.1217695999999999	0.1435488	1.3061676586638133	0.16327095733297667
Parallel	8	4	0.02288	0.0885823999999999	0.11146260000000001	1.6821678302856744	0.2102709787857093
Parallel	8	8	0.02435559999999998	0.1014478	0.1258032	1.4904135983822353	0.1863016997977794
Parallel	8	16	0.0236308	0.1187376	0.1423682	1.3169991613295666	0.16462489516619583

Tabella 25: Measurement with O1 optimization

5.5.2 2500 vertices - 50% of density

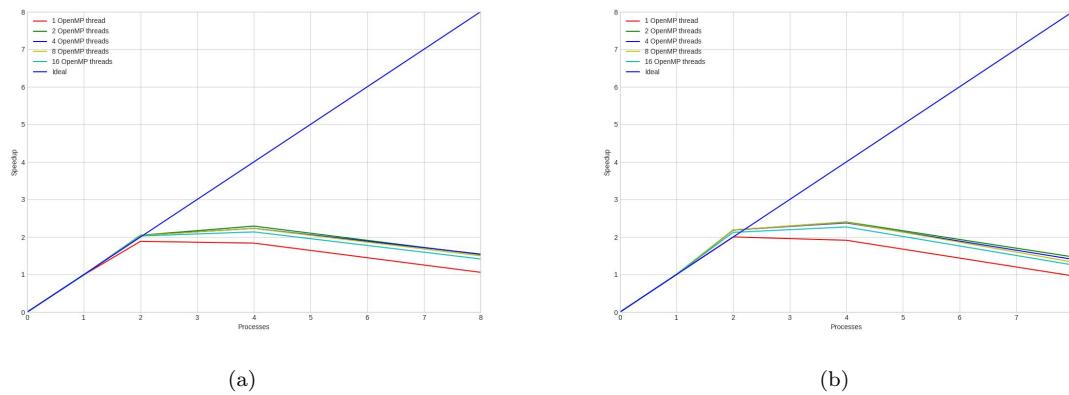


Figura 53: (c) Speedup with no optimization - (d) Speedup with $O1$ optimization

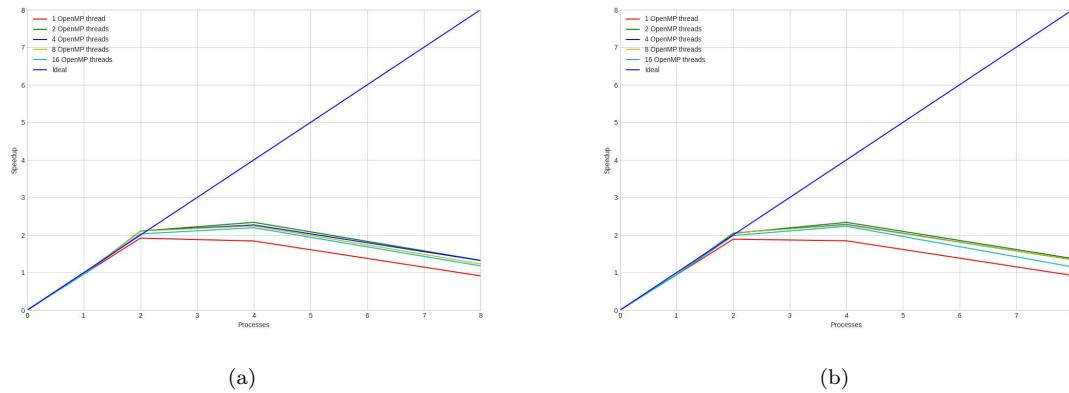


Figura 54: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Kosaraju	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.6142852000000001	0.0	0.6142852000000001	1.0	1.0
Parallel	1	1	0.6102293999999999	0.0	0.6102293999999999	1.0066463529944643	1.0066463529944643
Parallel	1	2	0.610177	0.0	0.610177	1.0067328004824831	1.0067328004824831
Parallel	1	4	0.6103080000000001	0.0	0.6103080000000001	1.0065167095958107	1.0065167095958107
Parallel	1	8	0.6089992	0.0	0.6089992	1.0086798143577203	1.0086798143577203
Parallel	1	16	0.6101135999999999	0.0	0.6101135999999999	1.006837415196121	1.006837415196121
Parallel	2	1	0.2541324	0.0526359999999995	0.3067686	2.0024383199584315	1.0012191599792157
Parallel	2	2	0.2541982	0.026652600000000005	0.2808502	2.187234334631776	1.0936171667315888
Parallel	2	4	0.2542297999999995	0.0266568	0.2808864	2.1869524476799165	1.0934762238399582
Parallel	2	8	0.253667	0.0267882	0.2804554	2.1903133261117453	1.0951566630558727
Parallel	2	16	0.2542532	0.0350288	0.2892820000000004	2.12348227680948	1.06174113840474
Parallel	4	1	0.1193104	0.2016354000000002	0.3209458	1.913984230359145	0.47849605758978625
Parallel	4	2	0.1192223999999999	0.1367314	0.2559538	2.3999846847360735	0.5999961711840184
Parallel	4	4	0.1189713999999999	0.1394936	0.2584652	2.376665021055059	0.5941662552637648
Parallel	4	8	0.1196246000000001	0.137184	0.2568086	2.3919962181951853	0.5979990545487963
Parallel	4	16	0.1200026	0.1507972000000002	0.2707998	2.2684108333905715	0.5671027083476429
Parallel	8	1	0.0652122	0.570825	0.6360372	0.9658007424722959	0.12072509280903698
Parallel	8	2	0.065706	0.3516758000000004	0.4173822000000004	1.4717570610342272	0.1839696326292784
Parallel	8	4	0.0661853999999999	0.3704338	0.4366194	1.406912290200573	0.17586403627507163
Parallel	8	8	0.0668569999999999	0.3931410000000001	0.4599978	1.3354089954343262	0.16692612442929078
Parallel	8	16	0.066733	0.4220766	0.4888096	1.2566962678310738	0.1570870334788422

Tabella 26: Measurement with O1 optimization

5.5.3 2500 vertices - 75% of density

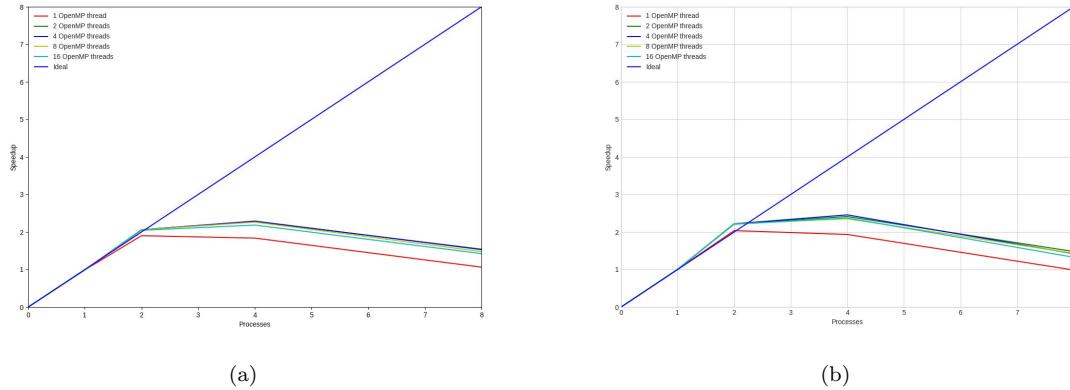


Figura 55: (a) Speedup with no optimization - (b) Speedup with *O1* optimization

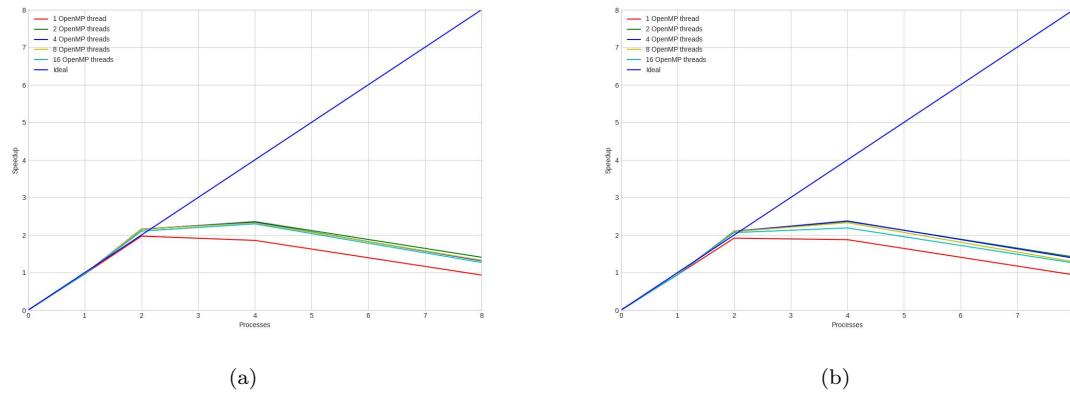


Figura 56: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Kosaraju	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	1.051844	0.0	1.051844	1.0	1.0
Parallel	1	1	1.0457075999999998	0.0	1.0457075999999998	1.0058681795943725	1.0058681795943725
Parallel	1	2	1.0439520000000002	0.0	1.0439520000000002	1.007559734547182	1.007559734547182
Parallel	1	4	1.0454519999999998	0.0	1.0454519999999998	1.0061071728098048	1.0061071728098048
Parallel	1	8	1.0454478	0.0	1.0454478	1.0061181438231541	1.0061181438231541
Parallel	1	16	1.0455804	0.0	1.0455804	1.0059905484073726	1.0059905484073726
Parallel	2	1	0.4298963999999996	0.0863763999999999	0.5162726000000001	2.0373810270000767	1.0186905135000384
Parallel	2	2	0.4301187999999997	0.0434293999999999	0.4735479999999997	2.221198273459079	1.1105991367295396
Parallel	2	4	0.4304928	0.04354379999999994	0.4740366	2.218908835309341	1.1094544176546706
Parallel	2	8	0.4307708	0.04363799999999996	0.4744089999999997	2.2171670436269126	1.1085835218134563
Parallel	2	16	0.4299009999999999	0.0468074	0.47671740000000007	2.2064308959563883	1.1032154479781942
Parallel	4	1	0.2009918	0.3427253999999996	0.5437166	1.934545770829877	0.4836361442707469
Parallel	4	2	0.2010844	0.2345645999999998	0.4356487999999995	2.4144310738374584	0.6306077684593646
Parallel	4	4	0.2011178	0.22696	0.42807780000000006	2.4571327922167416	0.6142831980541854
Parallel	4	8	0.2015065999999995	0.2452874	0.44679440000000004	2.354201395541215	0.5885503488853038
Parallel	4	16	0.2013352	0.24146600000000001	0.4428012	2.3754316835636398	0.5938579208909099
Parallel	8	1	0.10956740000000001	0.9555200000000001	1.0650870000000001	0.9875627392879764	0.12344578424109955
Parallel	8	2	0.1100904	0.6003244	0.710415	1.4806049984868	0.18507562481085
Parallel	8	4	0.1101835999999998	0.6297222	0.7390960000000002	1.4215913913388996	0.17769892391736244
Parallel	8	8	0.11137920000000001	0.6226374	0.7340166	1.432997564360261	0.179112469554503263
Parallel	8	16	0.1109378	0.68105	0.791987999999999	1.3281059814037588	0.16601324767546985

Tabella 27: Measurement with O1 optimization

5.6 Case Study 4

5.6.1 10000 vertices - 25% of density

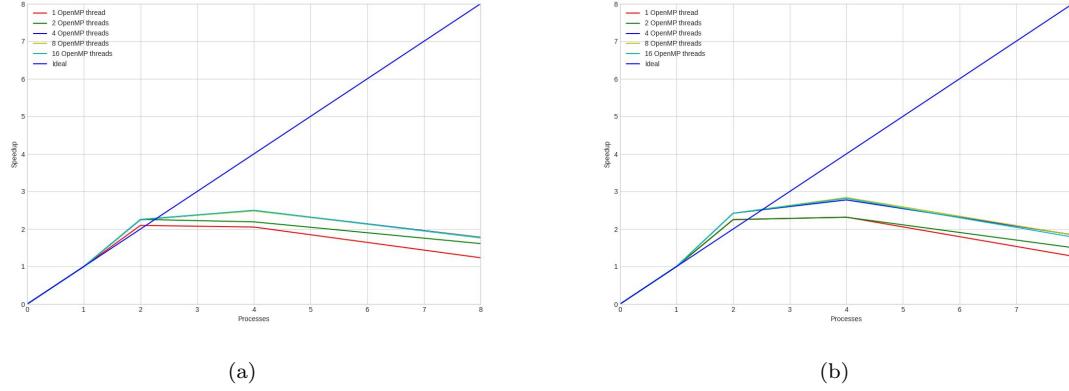


Figura 57: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

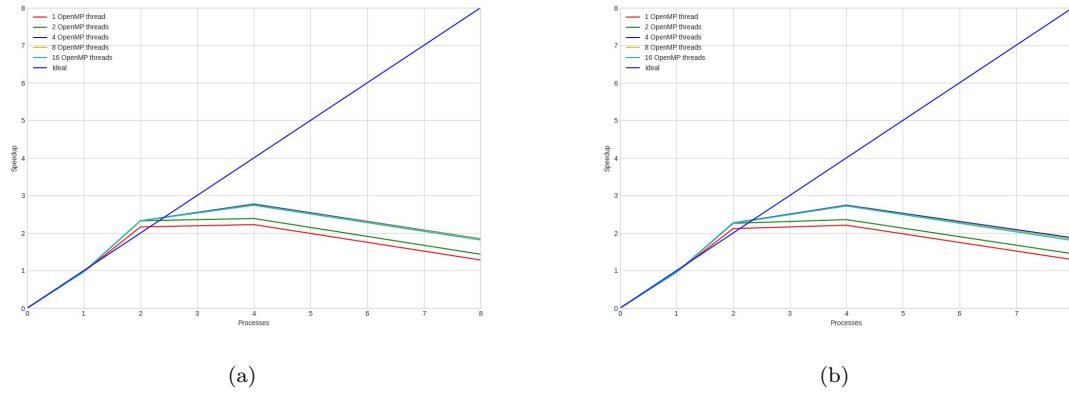


Figura 58: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Kosaraju	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	4.6047184	0.0	4.6047184	1.0	1.0
Parallel	1	1	4.5881896	0.0	4.5881896	1.0036024666461039	1.0036024666461039
Parallel	1	2	4.5959942	0.0	4.5959942	1.0018982182353495	1.0018982182353495
Parallel	1	4	4.5764958	0.0	4.5764958	1.0061668580576433	1.0061668580576433
Parallel	1	8	4.593551400000001	0.0	4.593551400000001	1.0024310166639256	1.0024310166639256
Parallel	1	16	4.5711714	0.0	4.5711714	1.0073388191044423	1.0073388191044423
Parallel	2	1	1.7510961999999999	0.2956798	2.046776	2.2497422287539037	1.1248711143769519
Parallel	2	2	1.7502232	0.2943407999999999	2.044564	2.252176209695564	1.126088104847782
Parallel	2	4	1.7537002000000002	0.1488809999999999	1.902581199999998	2.4202480293613755	1.2101240146806878
Parallel	2	8	1.7525688	0.14921040000000002	1.901779	2.4212689276724584	1.2106344638362292
Parallel	2	16	1.751189	0.1518658	1.9030550000000002	2.419645464792137	1.2098227323960684
Parallel	4	1	0.8126104	1.1748825999999999	1.9874934	2.31684714022195	0.5792117850554875
Parallel	4	2	0.8128476000000001	1.177126	1.9899738	2.3139593094140234	0.5784898273535058
Parallel	4	4	0.8135258000000001	0.8456706	1.6591957999999998	2.7752712488785236	0.6938178122196309
Parallel	4	8	0.813063	0.8094612	1.6225239999999999	2.8379970958827117	0.7094992739706779
Parallel	4	16	0.8173378	0.8175595999999998	1.6348976	2.8165179274836545	0.7041294818709136
Parallel	8	1	0.4639305999999997	3.143779400000006	3.6077099999999995	1.276354917662451	0.15954436470780636
Parallel	8	2	0.4641844	2.59724	3.0614244	1.5041097862811834	0.18801372328514793
Parallel	8	4	0.4650944	2.0274072	2.4925018000000003	1.8474283147960013	0.23092853934950017
Parallel	8	8	0.4675078000000003	2.0295244	2.4970322000000005	1.8440765000947923	0.23050956251184904
Parallel	8	16	0.4730484000000004	2.1033192	2.5763672	1.7872911904793698	0.22341139880992122

Tabella 28: Measurement with $O1$ optimization

5.6.2 10000 vertices - 50% of density

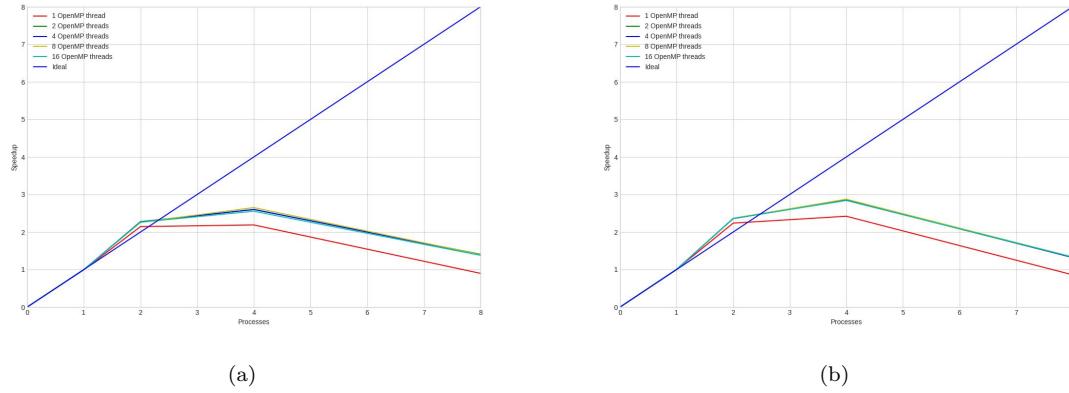


Figura 59: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

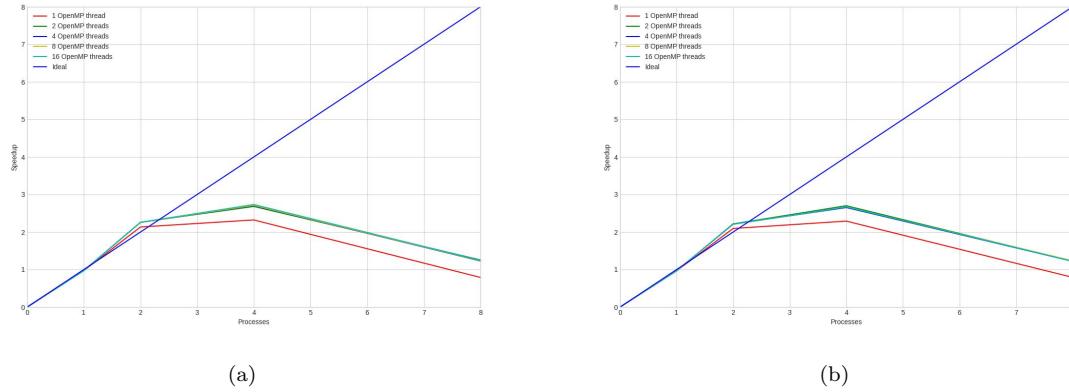


Figura 60: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Kosaraju	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	16.796227800000004	0.0	16.796227800000004	1.0	1.0
Parallel	1	1	16.68830779999997	0.0	16.68830779999997	1.0064668030631607	1.0064668030631607
Parallel	1	2	16.7258816	0.0	16.7258816	1.00420582912652	1.00420582912652
Parallel	1	4	16.7055284	0.0	16.7055284	1.0054293044690525	1.0054293044690525
Parallel	1	8	16.7320922	0.0	16.7320922	1.0038330890861338	1.0038330890861338
Parallel	1	16	16.737076000000002	0.0	16.737076000000002	1.0035341776544482	1.0035341776544482
Parallel	2	1	6.68968559999999	0.8183674	7.5080528	2.237095056124273	1.1185475280621364
Parallel	2	2	6.7069632	0.4113576	7.11832079999999	2.3595772474879197	1.1797886237439599
Parallel	2	4	6.714285000000001	0.4106067999999994	7.124891799999986	2.3574011046736185	1.178705523368093
Parallel	2	8	6.691253000000001	0.4109029999999996	7.10215619999999	2.3649476760311194	1.1824738380155597
Parallel	2	16	6.693527000000004	0.41639	7.10991679999999	2.3623662937940435	1.1811831468970218
Parallel	4	1	2.9452784	3.994055800000001	6.9393344	2.4204378736957834	0.6051094684239459
Parallel	4	2	2.953591999999996	2.9381614	5.891753400000001	2.850802920570301	0.7127007301425753
Parallel	4	4	2.9451928	2.91772	5.8629128	2.8648264732847473	0.7162066183211868
Parallel	4	8	2.9491666000000003	2.894377	5.843543599999999	2.874322320449531	0.7185805801123828
Parallel	4	16	2.931971199999996	2.9815496	5.91352079999999	2.8403092452131067	0.7100773113032767
Parallel	8	1	1.6777440000000001	17.8527728	19.5305168	0.8599991475904009	0.10749989344880011
Parallel	8	2	1.7062198	11.0147128	12.7209326	1.3203613546384174	0.16504516932980218
Parallel	8	4	1.711169	10.991712	12.70288079999999	1.3222376927287238	0.16527971159109048
Parallel	8	8	1.706790999999997	10.8748392	12.581630400000002	1.3349802264100845	0.16687252830126056
Parallel	8	16	1.705528199999998	10.878062400000001	12.5835908	1.3347722495871372	0.16684653119839216

Tabella 29: Measurement with O1 optimization

5.6.3 10000 vertices - 75% of density

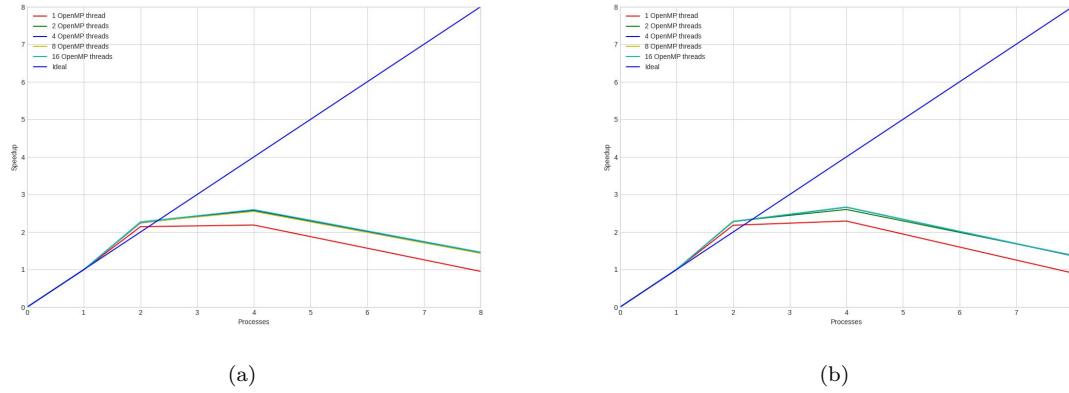


Figura 61: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

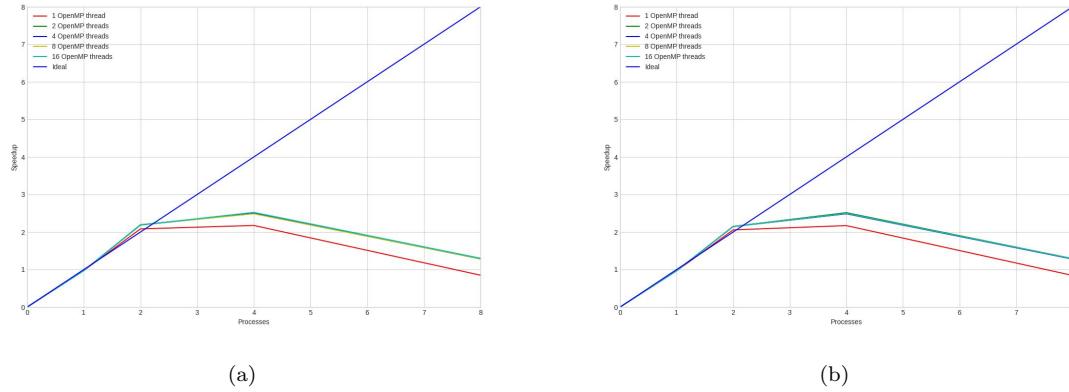


Figura 62: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Kosaraju	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	32.46680680000001	0.0	32.46680680000001	1.0	1.0
Parallel	1	1	32.27933899999999	0.0	32.27933899999999	1.0058076715883189	1.0058076715883189
Parallel	1	2	32.32407499999999	0.0	32.32407499999999	1.0044156499451264	1.0044156499451264
Parallel	1	4	32.2727116	0.0	32.2727116	1.0060142203854976	1.0060142203854976
Parallel	1	8	32.3266016	0.0	32.3266016	1.004337146283883	1.004337146283883
Parallel	1	16	32.285578400000006	0.0	32.285578400000006	1.0056132926520531	1.0056132926520531
Parallel	2	1	13.532039600000001	1.359787599999998	14.8918272	2.180176170725377	1.0900880853626884
Parallel	2	2	13.516655200000002	0.6807154	14.1973704	2.286818325173795	1.1434091625868974
Parallel	2	4	13.5817914	0.680527799999999	14.262319	2.276404475317093	1.1382022376585466
Parallel	2	8	13.51211319999998	0.6798088	14.1919218	2.2876962864888397	1.1438481432444199
Parallel	2	16	13.5506124	0.6889174	14.239530000000002	2.2800476420218927	1.1400238210109463
Parallel	4	1	6.3275904	7.834450800000001	14.1620408	2.292523179286421	0.5731307948216052
Parallel	4	2	6.3424792	6.1403266	12.4828056	2.600922247799806	0.6502305619499515
Parallel	4	4	6.3196682	5.8868562	12.206524400000001	2.659791250652807	0.6649478126632018
Parallel	4	8	6.339520200000001	5.852125	12.191645600000001	2.663037285138932	0.665759321284733
Parallel	4	16	6.310885600000001	5.867402	12.1782872	2.6659583787776007	0.6664895946944002
Parallel	8	1	4.2317986	31.5099036	35.74170200000004	0.9083732722073505	0.11354665902591882
Parallel	8	2	4.1967884	19.2950466	23.49183480000003	1.382046446197553	0.17275580577469413
Parallel	8	4	4.1966854	19.60025399999998	23.7969408	1.3643269138191076	0.17054086422738846
Parallel	8	8	4.2021982	19.4705906	23.6727888	1.3714821297269382	0.17143526621586727
Parallel	8	16	4.1932804	19.45027800000004	23.6435578	1.3731777203175406	0.17164721503969257

Tabella 30: Measurement with $O1$ optimization

5.6.4 10000 vertices - Special Cases - 0%

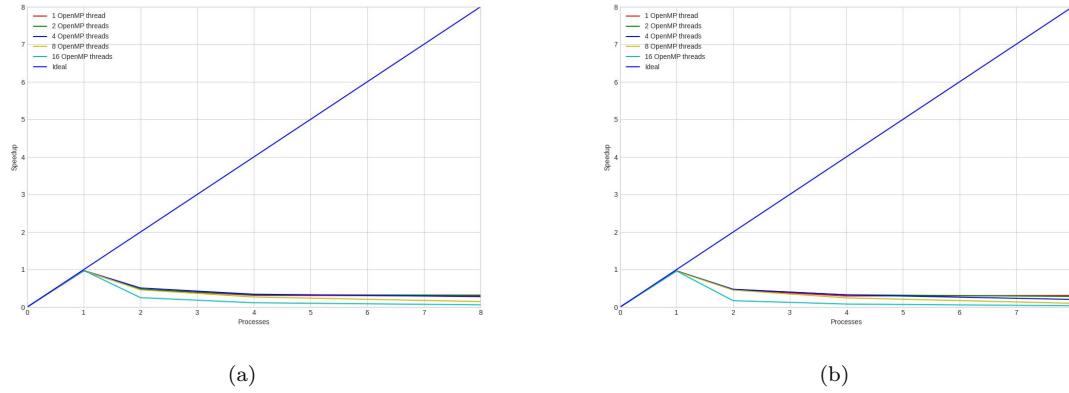


Figura 63: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

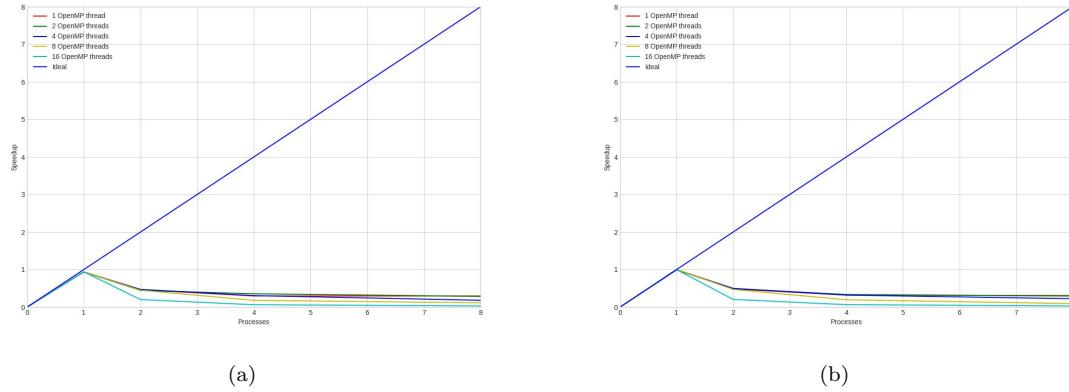


Figura 64: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Kosaraju	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.0021524	0.0	0.0021524	1.0	1.0
Parallel	1	1	0.0021578	0.0	0.0021578	0.9974974511076096	0.9974974511076096
Parallel	1	2	0.0021742000000000003	0.0	0.0021742000000000003	0.9899733235212951	0.9899733235212951
Parallel	1	4	0.0021578	0.0	0.0021578	0.9974974511076096	0.9974974511076096
Parallel	1	8	0.0021398	0.0	0.0021398	1.00588840078512	1.00588840078512
Parallel	1	16	0.002156	0.0	0.002156	0.9983302411873841	0.9983302411873841
Parallel	2	1	0.0036842	0.0007821999999999999	0.0044662000000000001	0.48193094800949343	0.24096547400474672
Parallel	2	2	0.0037302000000000004	0.0006566	0.0043866	0.4906761500934665	0.24533807504673324
Parallel	2	4	0.0036974	0.0006236	0.0043214	0.49807932614430506	0.24903966307215253
Parallel	2	8	0.0038788000000000004	0.0007015999999999999	0.0045806	0.46989477361044407	0.23494738680522204
Parallel	2	16	0.0038772	0.006632799999999995	0.0105102	0.20479153584137313	0.10239576792068657
Parallel	4	1	0.0051284	0.001595199999999997	0.0067234	0.3201356456554719	0.08003391141386798
Parallel	4	2	0.005094	0.0013160000000000001	0.0064098	0.3357983088395894	0.08394957720989735
Parallel	4	4	0.005466	0.0012482	0.00671419999999999	0.3205743052038963	0.08014357630097407
Parallel	4	8	0.005079800000000001	0.0058788	0.010959	0.19640478145816226	0.049101195364540565
Parallel	4	16	0.0048696	0.02848420000000005	0.0333538	0.06453237712044804	0.01613309428011201
Parallel	8	1	0.0051804	0.001831599999999999	0.0070118	0.30696825351550244	0.038371031689437805
Parallel	8	2	0.00522399999999995	0.002087600000000002	0.0073116	0.29438153071831064	0.03679769133978883
Parallel	8	4	0.005865400000000001	0.003777600000000003	0.009643200000000001	0.2232039157126265	0.02790048946407831
Parallel	8	8	0.0052266	0.0179794	0.02320599999999997	0.0927518745152116	0.01159398431440145
Parallel	8	16	0.0053796	0.0661244	0.0715038	0.03010189668241408	0.00376273708530176

Tabella 31: Measurement with O2 optimization

5.6.5 10000 vertices - Special Cases - 100%

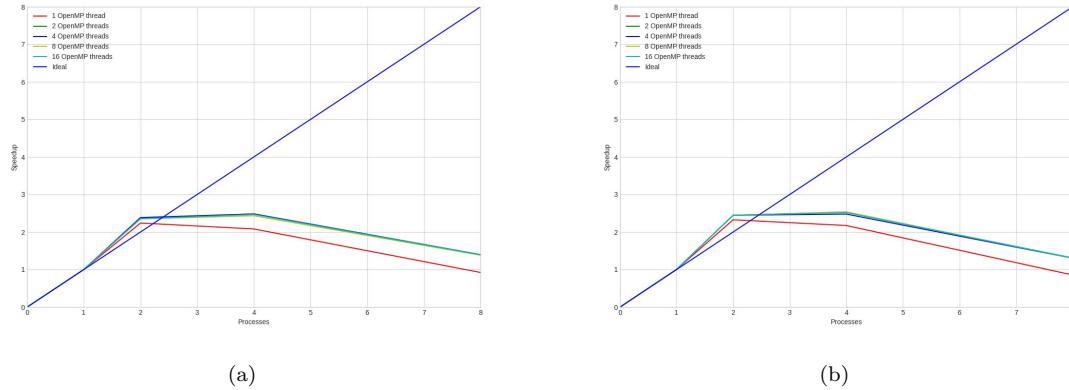


Figura 65: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

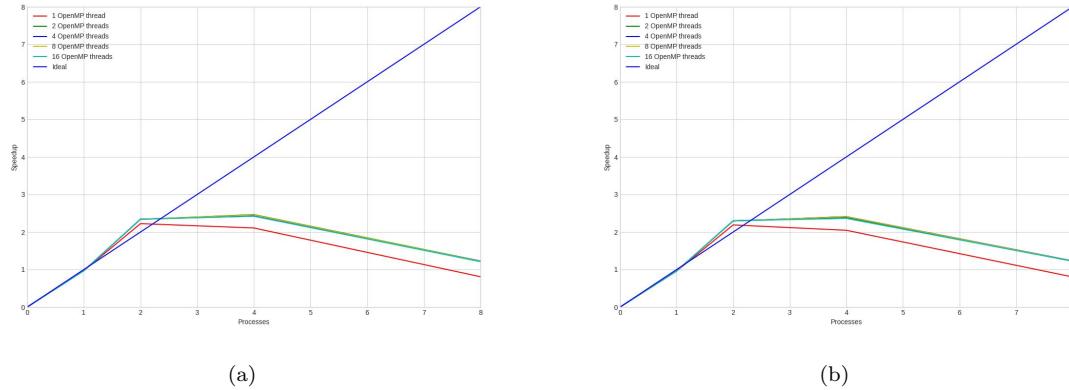


Figura 66: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Kosaraju	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	51.669504800000006	0.0	51.669504800000006	1.0	1.0
Parallel	1	1	51.39050119999996	0.0	51.39050119999996	1.005429088907193	1.005429088907193
Parallel	1	2	51.46892800000005	0.0	51.46892800000005	1.0038970463888426	1.0038970463888426
Parallel	1	4	51.35480279999995	0.0	51.35480279999995	1.006127995491008	1.006127995491008
Parallel	1	8	51.327782	0.0	51.327782	1.006657657640457	1.006657657640457
Parallel	1	16	51.410499800000004	0.0	51.410499800000004	1.005037978642643	1.005037978642643
Parallel	2	1	20.0413236	2.174503999999998	22.2158276	2.3257969826881446	1.1628984913440723
Parallel	2	2	20.016664400000003	1.0888214	21.1054856	2.448155222735079	1.2240776113675396
Parallel	2	4	20.03604739999998	1.0889616	21.1250092	2.4458926531497087	1.2229463265748544
Parallel	2	8	20.0166854	1.0897078	21.10639339999998	2.448049925952769	1.2240249629763844
Parallel	2	16	20.032993400000002	1.090061400000002	21.1230552	2.446118911813477	1.2230594559067385
Parallel	4	1	8.604989	15.163447	23.7684358	2.173870642341555	0.5434676605853888
Parallel	4	2	8.5861154	11.8555806	20.441696	2.527652539202227	0.6319131348005568
Parallel	4	4	8.610945000000001	12.2414702	20.8524156	2.47786663143238	0.619466657858095
Parallel	4	8	8.593002	11.75503799999999	20.348039800000002	2.53928659980014	0.6348216499950035
Parallel	4	16	8.602756600000001	11.9277536	20.53051	2.516718035742902	0.6291795089357255
Parallel	8	1	4.7579116	55.642517	60.4004288	0.8554493043599055	0.10693116304498819
Parallel	8	2	4.75253879999999	34.686370800000006	39.43891	1.3101149296468895	0.16376436620586118
Parallel	8	4	4.77487039999999	34.888955	39.66382559999996	1.3026858609422691	0.16283573261778364
Parallel	8	8	4.773937	34.8914938	39.6654312	1.3026331300792717	0.16282914125990897
Parallel	8	16	4.7803746	34.6970002	39.4773748	1.3088384185059845	0.16360480231324806

Tabella 32: Measurement with $O1$ optimization

6 Cluster

In addition to the measurements taken on our machine, we also tried running our program on a raspberry pi cluster, specifically, the front-end node we connected to is a raspberry *pi3*, while the worker nodes are raspberry *pi4*. In order to make a fair comparison with the measurements taken on our machine, on the cluster the program was launched using the same graphs except for the one with 10000 nodes.

Therefore, the following graphs were considered:

- graph with a number of vertices equal to 500:
 1. **low density 25%**: 0 - 125
 2. **medium density 50%**: 125 - 250
 3. **high density 75%**: 250 - 375
- graph with a number of vertices equal to 1000:
 1. **low density 25%**: 0 - 250
 2. **medium density 50%**: 250 - 500
 3. **high density 75%**: 500 - 750
- graph with a number of vertices equal to 2500:
 1. **low density 25%**: 0 - 625
 2. **medium density 50%**: 625 - 1350
 3. **high density 75%**: 1350 - 1875

For each density level of each graph on which the execution of the algorithm version has been defined, as regards the management of MPI processes and OpenMP threads, the latter has been structured by assigning the following resources to each of prefixed density level:

- 0 MPI process - {1, 2, 4, 8} OpenMP threads (*Sequential version*)
- 1 MPI process - {1, 2, 4, 8} OpenMP threads (*Parallel version*)
- 2 MPI process - {1, 2, 4, 8} OpenMP threads (*Parallel version*)
- 4 MPI process - {1, 2, 4, 8} OpenMP threads (*Parallel version*)
- 8 MPI process - {1, 2, 4, 8} OpenMP threads (*Parallel version*)

Therefore, from the considerations made, we can define the case studies analysed:

- **Case Study 1**: the sequential program and the parallel program are compiled on a graph with a number of vertices equal to 500
- **Case Study 2**: the sequential program and the parallel program are compiled on a graph with a number of vertices equal to 1000
- **Case Study 3**: the sequential program and the parallel program are compiled on a graph with a number of vertices equal to 2500

6.1 Tarjan

6.1.1 500 vertices - 25% of density

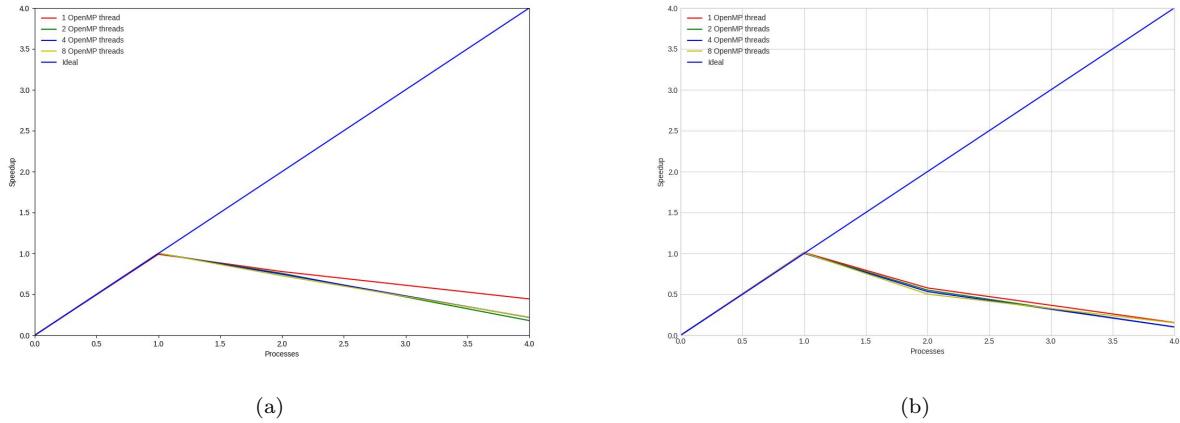


Figura 67: (a) Speedup with no optimization - (b) Speedup with *O1* optimization

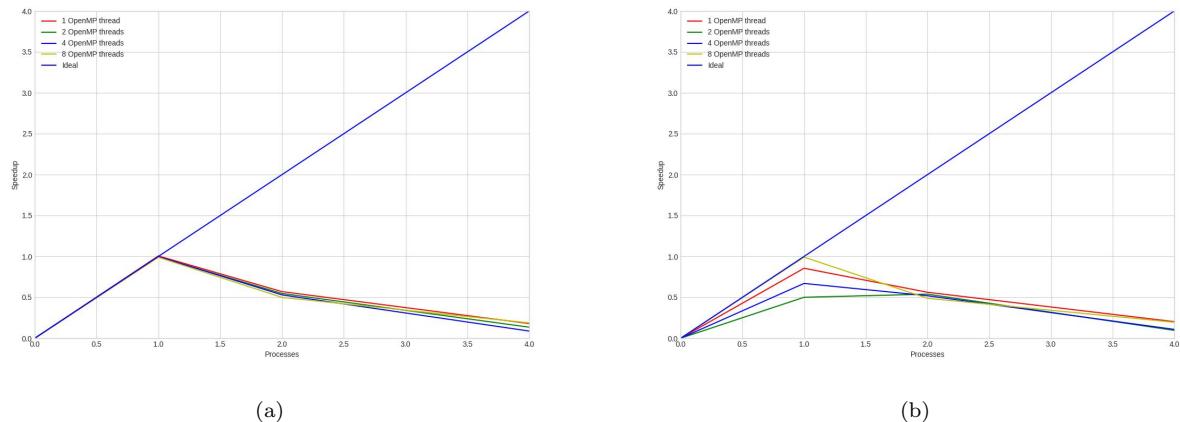


Figura 68: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Tarjan	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.006798	0.0	0.006798	1.0	1.0
Parallel	1	1	0.006722000000000005	0.0	0.006722000000000005	1.0113061588812853	1.0113061588812853
Parallel	1	2	0.006764	0.0	0.006764	1.0050266114725015	1.0050266114725015
Parallel	1	4	0.006796	0.0	0.006796	1.0002942907592702	1.0002942907592702
Parallel	1	8	0.0067475	0.0	0.0067475	1.0074842534271953	1.0074842534271953
Parallel	2	1	0.0026815	0.009071	0.011752499999999999	0.5784301212507977	0.28921506062539887
Parallel	2	2	0.002700499999999998	0.009611999999999999	0.0123125	0.5521218274111676	0.2760609137055838
Parallel	2	4	0.0026725	0.0100635	0.012736	0.5337625628140703	0.26688128140703515
Parallel	2	8	0.002679	0.0108295	0.0135085	0.5032387015582781	0.25161935077913905
Parallel	4	1	0.0032985	0.040594000000000005	0.0438925	0.15487839608133508	0.03871959082033377
Parallel	4	2	0.0034384999999999997	0.063927	0.06736500000000001	0.10091293698508126	0.02522834246270315
Parallel	4	4	0.00244	0.06442200000000001	0.066862	0.10167210074481768	0.02541802518620442
Parallel	4	8	0.003229	0.0415465	0.044775	0.15182579564489113	0.03795644891122278

Tabella 33: Measurement with O1 optimization

6.1.2 500 vertices - 50% of density

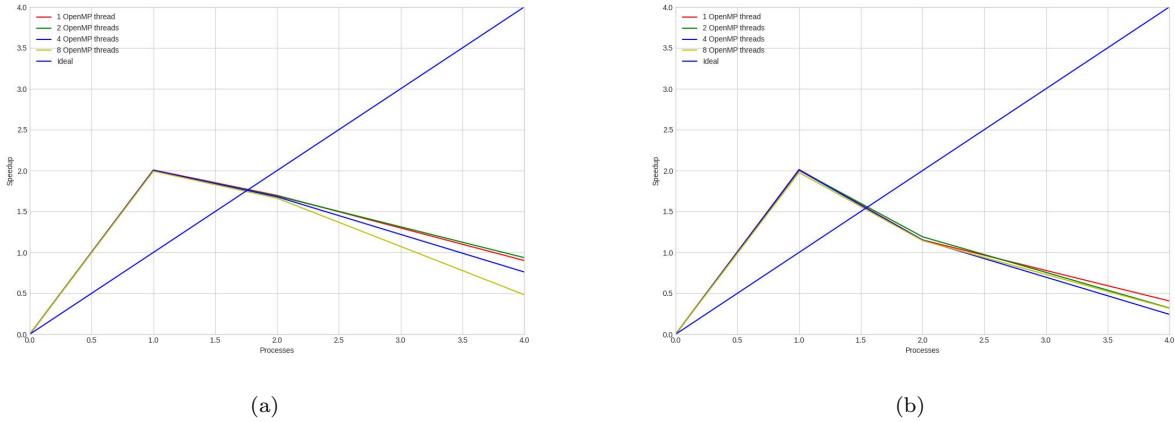


Figura 69: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

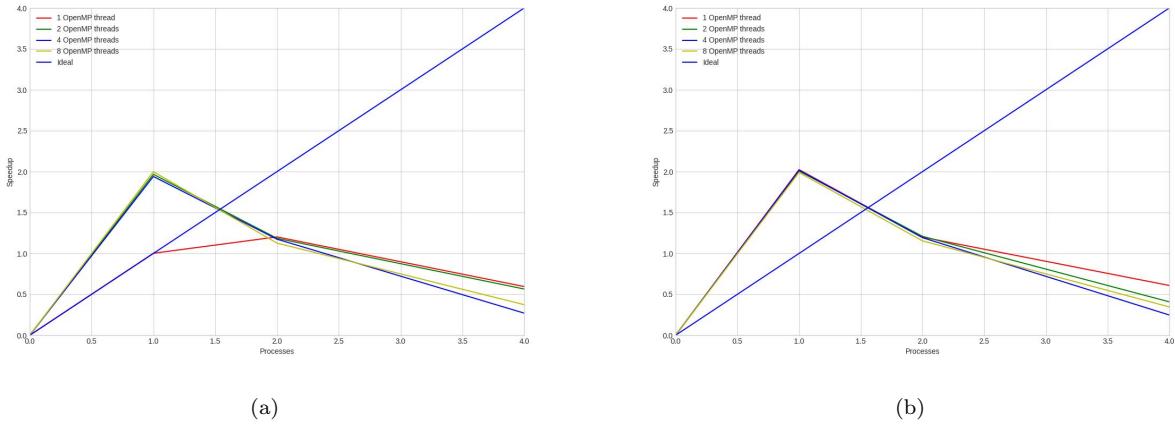


Figura 70: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Tarjan	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.1119405	0.0	0.1119405	1.0	1.0
Parallel	1	1	0.0557535	0.0	0.0557535	2.0077752966181497	2.0077752966181497
Parallel	1	2	0.055941	0.0	0.055941	2.0010457446238	2.0010457446238
Parallel	1	4	0.055836	0.0	0.055836	2.0048087255534064	2.0048087255534064
Parallel	1	8	0.056162000000000004	0.0	0.056162000000000004	1.9931715394750897	1.9931715394750897
Parallel	2	1	0.022352499999999997	0.043597	0.0659495	1.6973669246923784	0.8486834623461892
Parallel	2	2	0.022383	0.04384	0.0662225	1.6903695873758917	0.8451847936879459
Parallel	2	4	0.0223725	0.044256500000000004	0.066629	1.6800567320536104	0.8400283660268052
Parallel	2	8	0.022381	0.044908000000000003	0.0672885	1.6635903609086247	0.8317951804543123
Parallel	4	1	0.01176	0.112472	0.124232	0.9010601133363385	0.22526502833404863
Parallel	4	2	0.011737000000000001	0.1077975	0.119534	0.936474141248515	0.23411853531212876
Parallel	4	4	0.0117725	0.1353025	0.147075	0.7611116777154512	0.1902779194288628
Parallel	4	8	0.023448999999999998	0.208146	0.231595	0.48334592715732205	0.12083648178933051

Tabella 34: Measurement with O0 optimization

6.1.3 500 vertices - 75% of density

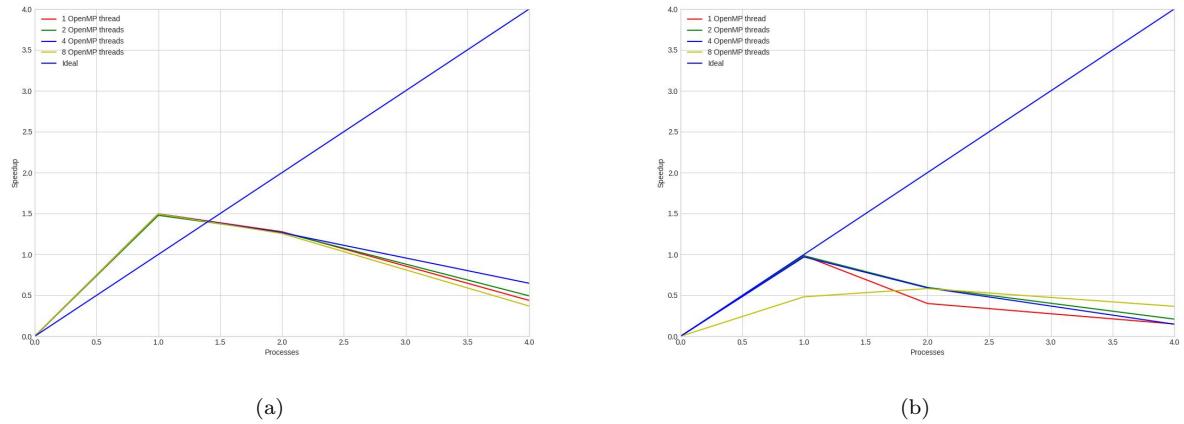


Figura 71: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

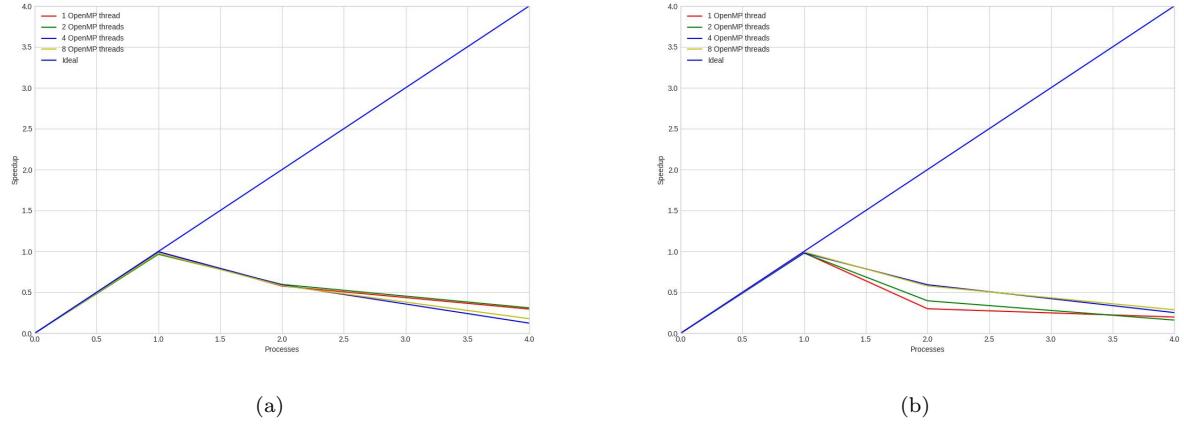


Figura 72: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Tarjan	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.1370005	0.0	0.1370005	1.0	1.0
Parallel	1	1	0.0915655	0.0	0.0915655	1.4962021722155179	1.4962021722155179
Parallel	1	2	0.0927035	0.0	0.0927035	1.4778352489388211	1.4778352489388211
Parallel	1	4	0.09162000000000001	0.0	0.09162000000000001	1.4953121589172669	1.4953121589172669
Parallel	1	8	0.091675	0.0	0.091675	1.4944150531769838	1.4944150531769838
Parallel	2	1	0.0363815	0.070842	0.107224	1.277703685741998	0.638851842870999
Parallel	2	2	0.03640699999999995	0.0714655	0.1078725	1.2700224802428794	0.6350112401214397
Parallel	2	4	0.03643100000000005	0.0717619999999999	0.10819300000000001	1.2662602941040546	0.6331301470520273
Parallel	2	8	0.03642099999999995	0.0726845	0.109106	1.2556642164500578	0.6278321082250289
Parallel	4	1	0.02866599999999997	0.2840405	0.312706	0.43811279604484726	0.10952819901121182
Parallel	4	2	0.0284525	0.248655	0.277107	0.4943956666558405	0.12359891666396013
Parallel	4	4	0.01892300000000002	0.1925965	0.2115195	0.647696784457225	0.16192419611430625
Parallel	4	8	0.03817900000000005	0.3329475	0.3711265	0.3691477164794214	0.09228692911985535

Tabella 35: Measurement with O0 optimization

6.1.4 1000 vertices - 25% of density

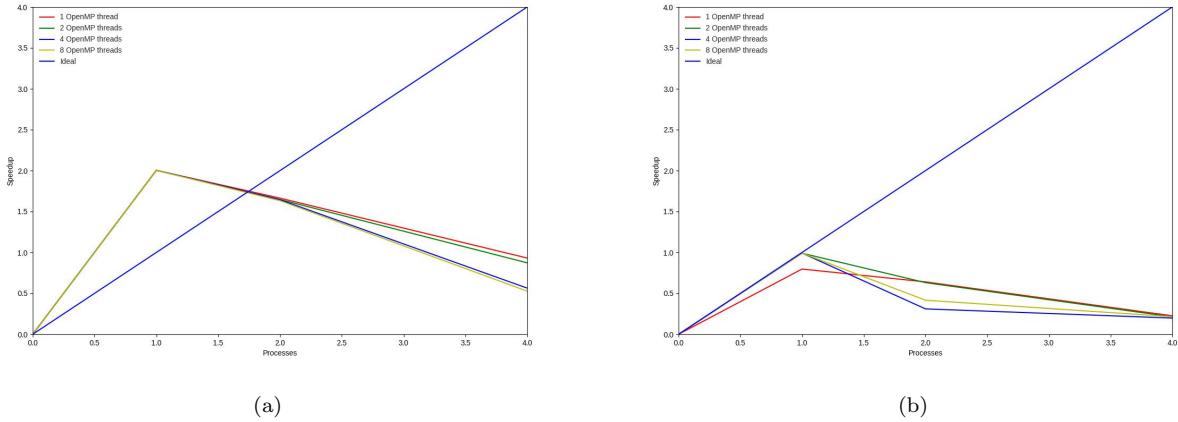


Figura 73: (a) Speedup with no optimization - (b) Speedup with *O1* optimization

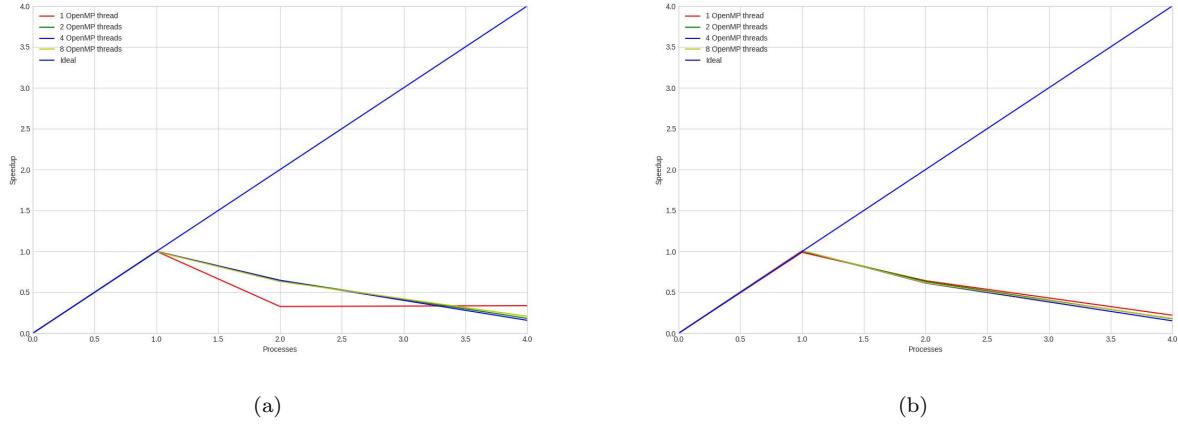


Figura 74: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Tarjan	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.1677215	0.0	0.1677215	1.0	1.0
Parallel	1	1	0.08372299999999999	0.0	0.08372299999999999	2.0032906130931765	2.0032906130931765
Parallel	1	2	0.0837155	0.0	0.0837155	2.0034700861847567	2.0034700861847567
Parallel	1	4	0.083701	0.0	0.083701	2.003817158695834	2.003817158695834
Parallel	1	8	0.083754500000000001	0.0	0.083754500000000001	2.0025371771069014	2.0025371771069014
Parallel	2	1	0.0333405	0.06747	0.10081	1.6637387163971828	0.8318693581985914
Parallel	2	2	0.0332655	0.0685635	0.1018285	1.6470978164266388	0.8235489082133194
Parallel	2	4	0.0333025	0.068686500000000001	0.101989	1.644505780035102	0.822252890017551
Parallel	2	8	0.033261	0.069466	0.102727	1.6326915027208038	0.8163457513604019
Parallel	4	1	0.0215405	0.15852	0.1800605	0.9314730326751286	0.23286825816878215
Parallel	4	2	0.021466	0.170658	0.1921245	0.8729834039906414	0.21824585099766036
Parallel	4	4	0.0321735	0.2653735	0.297547	0.5636806958228447	0.1409201739557112
Parallel	4	8	0.043039999999999995	0.27527749999999995	0.31831750000000003	0.5269000290590369	0.13172500726475922

Tabella 36: Measurement with O0 optimization

6.1.5 1000 vertices - 50% of density

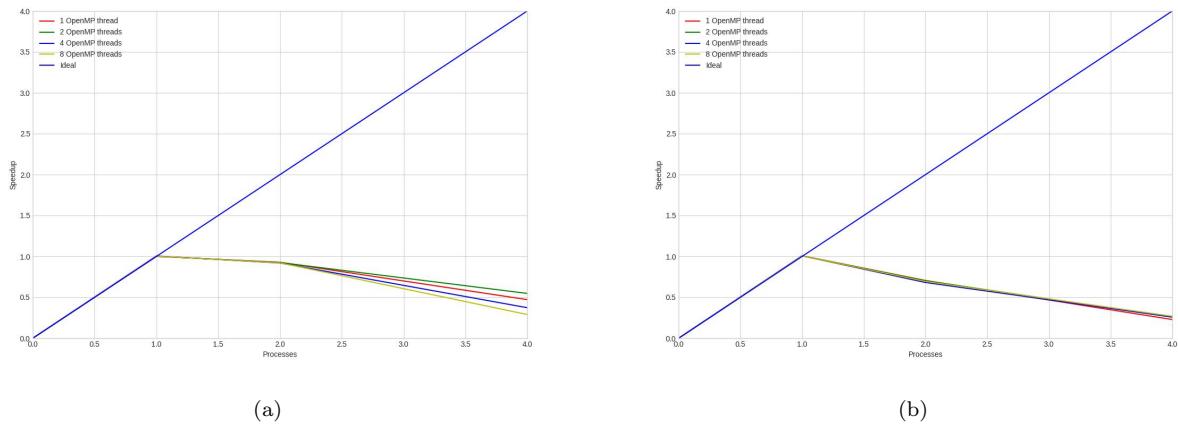


Figura 75: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

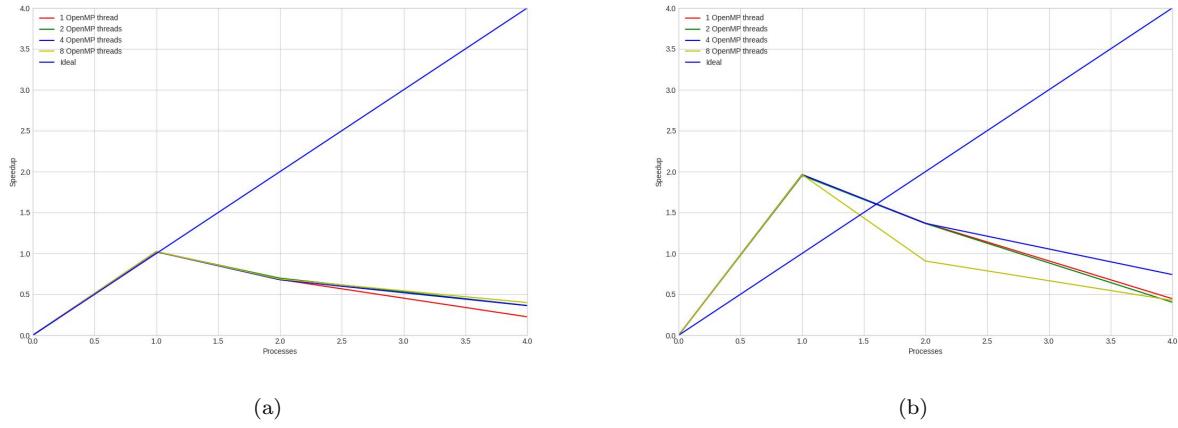


Figura 76: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Tarjan	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.173553	0.0	0.173553	1.0	1.0
Parallel	1	1	0.0885625	0.0	0.0885625	1.959666901905434	1.959666901905434
Parallel	1	2	0.088901000000000001	0.0	0.088901000000000001	1.952205262033048	1.952205262033048
Parallel	1	4	0.0882955	0.0	0.0882955	1.9655928105056317	1.9655928105056317
Parallel	1	8	0.0884035	0.0	0.0884035	1.9631915025988793	1.9631915025988793
Parallel	2	1	0.030814	0.095799	0.12661299999999998	1.370736022367372	0.68536801183686
Parallel	2	2	0.0307515	0.096322	0.1270735	1.365768629966122	0.682884314983061
Parallel	2	4	0.030677	0.0962355	0.12691249999999998	1.3675012311632033	0.6837506155816017
Parallel	2	8	0.045978	0.145363000000000002	0.191341	0.9070350839600504	0.4535175419800252
Parallel	4	1	0.02269299999999998	0.365568	0.3882605	0.447001433316034	0.1117503583290085
Parallel	4	2	0.029369	0.400739	0.4301075	0.40351075021942195	0.10087768755485549
Parallel	4	4	0.0151135	0.2189355	0.234049	0.741524210742195	0.18538105268554875
Parallel	4	8	0.0244915	0.38227900000000004	0.40677050000000003	0.4266607337552748	0.1066651834388187

Tabella 37: Measurement with $O3$ optimization

6.1.6 1000 vertices - 75% of density

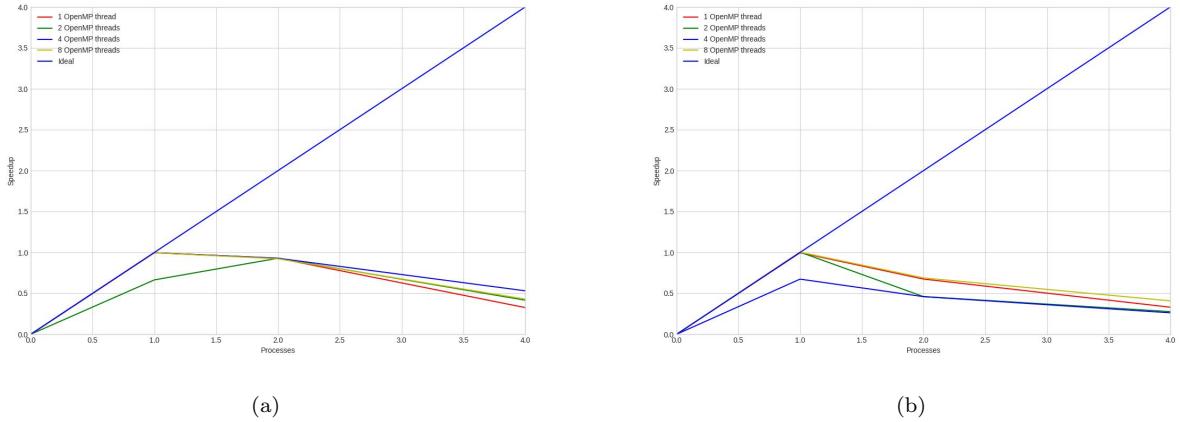


Figura 77: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

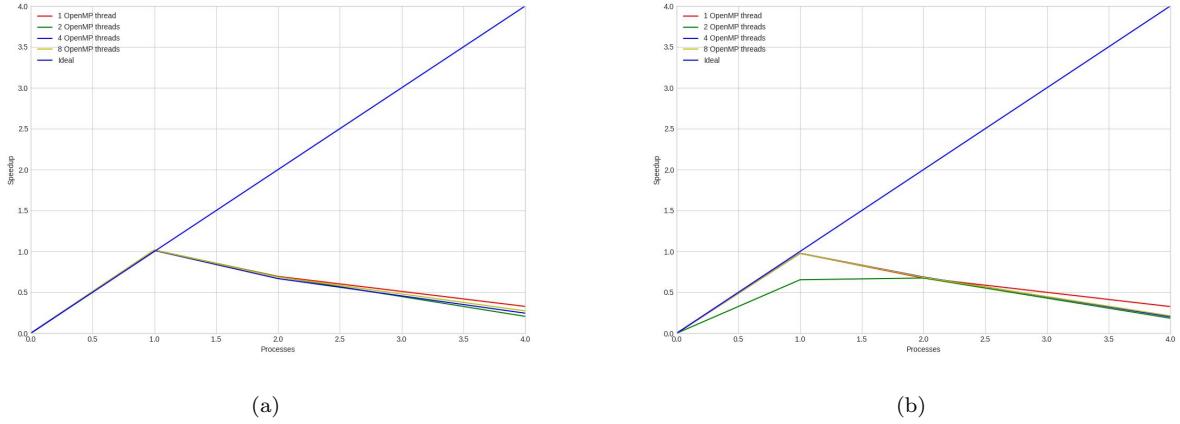


Figura 78: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Tarjan	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.1428705	0.0	0.1428705	1.0	1.0
Parallel	1	1	0.1408595	0.0	0.1408595	1.0142766373585028	1.0142766373585028
Parallel	1	2	0.14056	0.0	0.14056	1.0164378201479796	1.0164378201479796
Parallel	1	4	0.14120349999999998	0.0	0.14120349999999998	1.0118056563753735	1.0118056563753735
Parallel	1	8	0.140536	0.0	0.140536	1.016611402060682	1.016611402060682
Parallel	2	1	0.04922949999999995	0.156385	0.2056145	0.6948464237687517	0.34742321188437586
Parallel	2	2	0.049263	0.15758450000000002	0.20684750000000002	0.6907045045262814	0.3453522522631407
Parallel	2	4	0.049218	0.16457650000000001	0.2137935	0.668264002413544	0.334132001206772
Parallel	2	8	0.049678	0.158081	0.207759	0.6876741801799201	0.34383709008996005
Parallel	4	1	0.0248445	0.41063700000000003	0.4354814999999997	0.32807478618494706	0.08201869654623677
Parallel	4	2	0.049817	0.642997	0.692814	0.20621768613220864	0.05155442153305216
Parallel	4	4	0.04709149999999994	0.535216	0.5823075	0.24535232673458612	0.06133808168364653
Parallel	4	8	0.0421485	0.4804425	0.522591	0.2733887495192225	0.06834718737980562

Tabella 38: Measurement with $O2$ optimization

6.1.7 2500 vertices - 25% of density

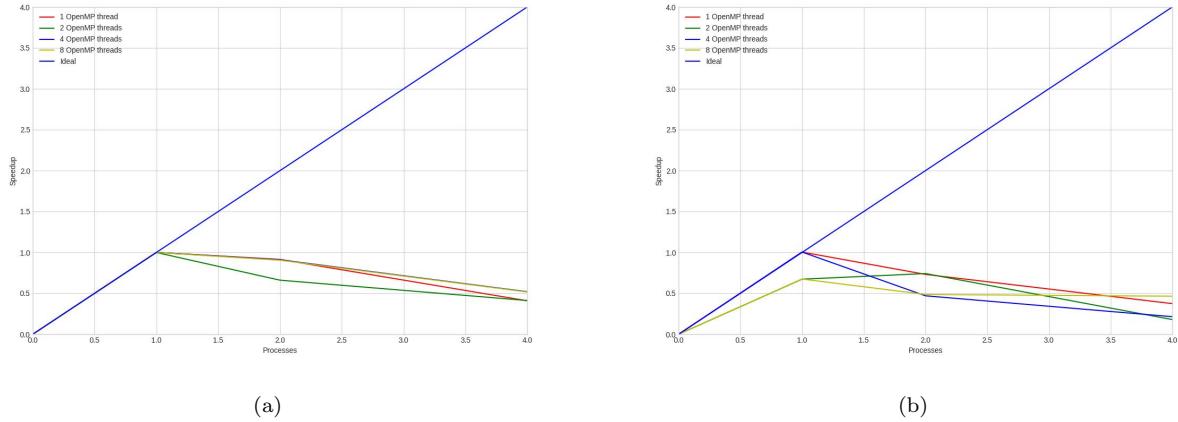


Figura 79: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

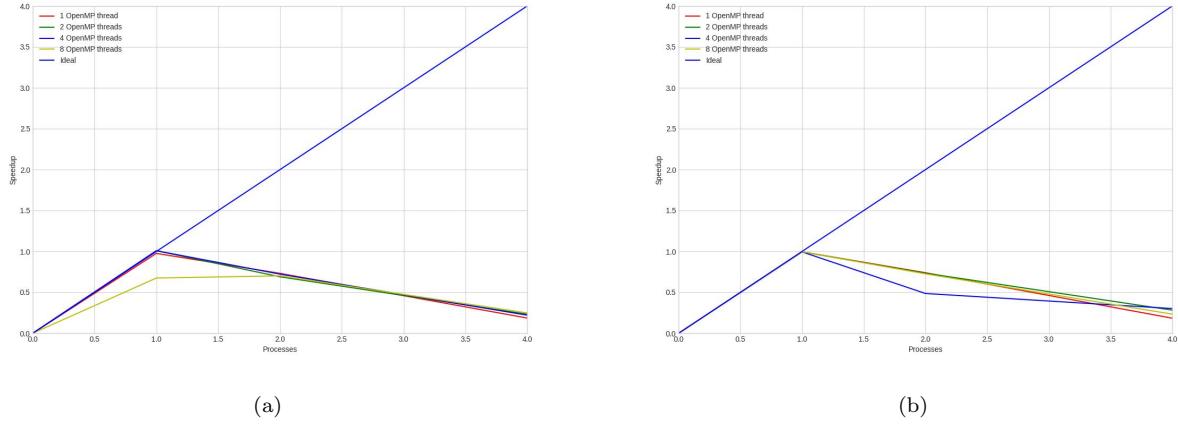


Figura 80: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Tarjan	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.6152565	0.0	0.6152565	1.0	1.0
Parallel	1	1	0.615105	0.0	0.615105	1.0002462994122954	1.0002462994122954
Parallel	1	2	0.6156505	0.0	0.6156505	0.9993600265085466	0.9993600265085466
Parallel	1	4	0.6148480000000001	0.0	0.6148480000000001	1.000664391849693	1.000664391849693
Parallel	1	8	0.6158144999999999	0.0	0.6158144999999999	0.9990938829793713	0.9990938829793713
Parallel	2	1	0.2494575	0.4232915000000004	0.6727485	0.9145416154774035	0.45727080773870177
Parallel	2	2	0.373166	0.557903	0.9310689999999999	0.6608065567643214	0.3304032783821607
Parallel	2	4	0.24908	0.426975	0.6760550000000001	0.91006870742765	0.455034353713825
Parallel	2	8	0.2491355	0.4307594999999996	0.679895	0.904928702226079	0.4524643511130395
Parallel	4	1	0.2107534999999998	1.293708	1.5044615000000001	0.40895463260442355	0.10223865815110589
Parallel	4	2	0.220294	1.271547	1.491841	0.4124142586240759	0.10310356465601897
Parallel	4	4	0.2185295	0.965278	1.1838085	0.5197263746627938	0.12993159366569845
Parallel	4	8	0.2252934999999998	0.9666995	1.191993	0.5161578130072911	0.12903945325182278

Tabella 39: Measurement with O0 optimization

6.1.8 2500 vertices - 50% of density

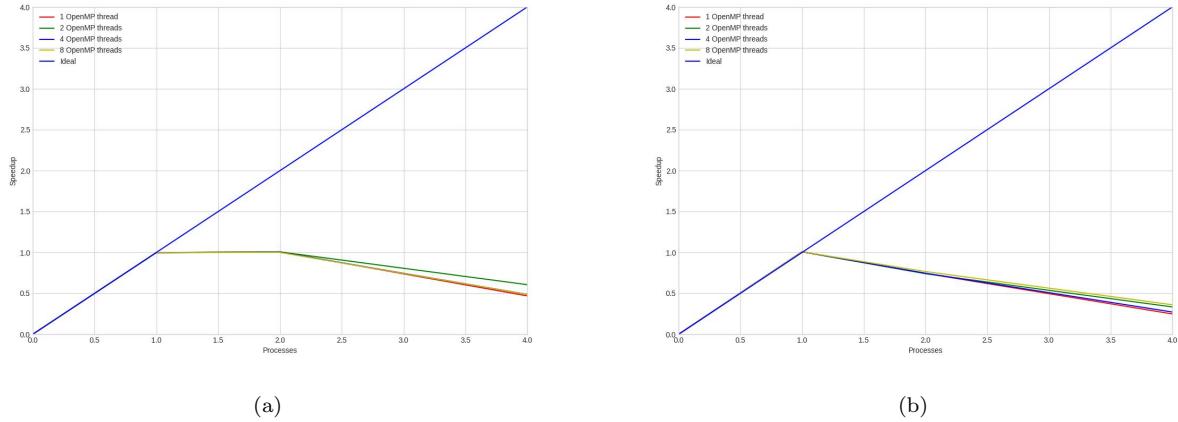


Figura 81: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

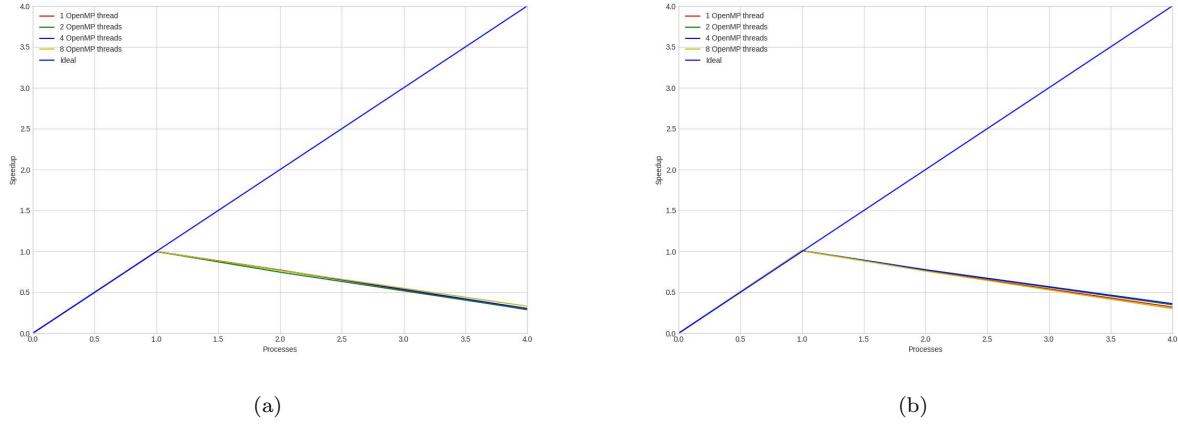


Figura 82: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Tarjan	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	1.905571	0.0	1.905571	1.0	1.0
Parallel	1	1	1.9094579999999999	0.0	1.9094579999999999	0.9979643438085573	0.9979643438085573
Parallel	1	2	1.9098525	0.0	1.9098525	0.9977582038403489	0.9977582038403489
Parallel	1	4	1.9089645000000002	0.0	1.9089645000000002	0.9982223346741125	0.9982223346741125
Parallel	1	8	1.909032	0.0	1.909032	0.9981870392953077	0.9981870392953077
Parallel	2	1	0.760822	1.1317865	1.892609	1.00684746885543	0.5034243734442772
Parallel	2	2	0.7601035	1.1329395	1.893043	1.0066179162332816	0.5033089581166408
Parallel	2	4	0.7624025000000001	1.137478	1.89988	1.0029954523443585	0.5014977261721792
Parallel	2	8	0.762183	1.1435655	1.905748	0.9999071230823802	0.4999535615411901
Parallel	4	1	0.3823815000000004	3.679025	4.0614065	0.46918992225969003	0.11729748056492251
Parallel	4	2	0.387339	2.7564960000000003	3.1438355	0.6061293601398674	0.15153234003496685
Parallel	4	4	0.578028	3.331068	3.9090955000000003	0.48747107866768663	0.12186776966692166
Parallel	4	8	0.5860529999999999	3.3269539999999997	3.9130075	0.48698373310043486	0.12174593327510871

Tabella 40: Measurement with O0 optimization

6.1.9 2500 vertices - 75% of density

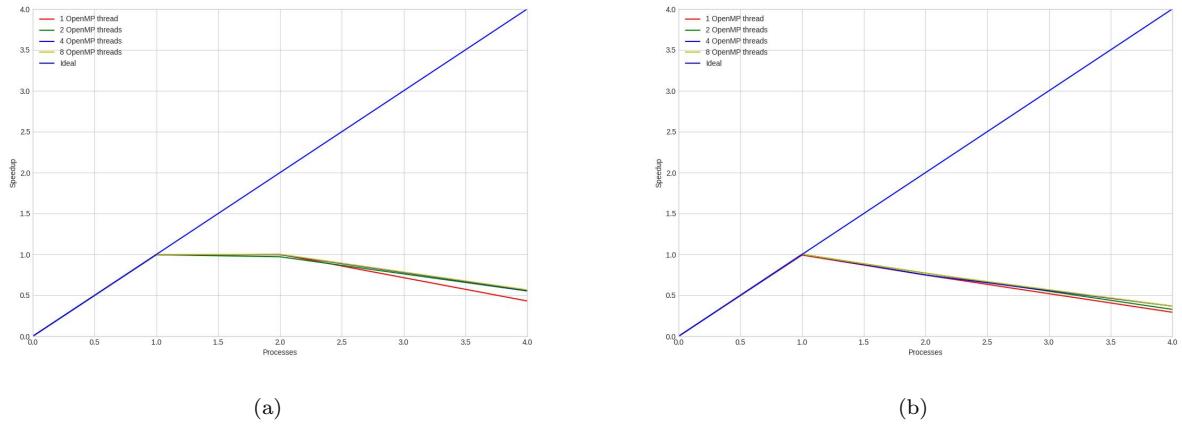


Figura 83: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

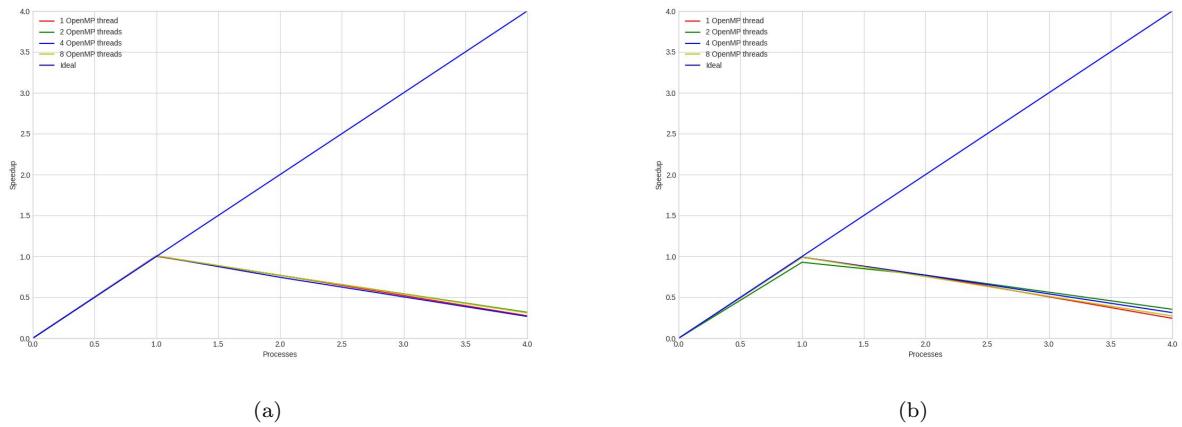


Figura 84: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Tarjan	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	3.085816	0.0	3.085816	1.0	1.0
Parallel	1	1	3.0966775	0.0	3.0966775	0.9964925311079373	0.9964925311079373
Parallel	1	2	3.097115	0.0	3.097115	0.9963517660790767	0.9963517660790767
Parallel	1	4	3.09511	0.0	3.09511	0.9969971988071505	0.9969971988071505
Parallel	1	8	3.095537	0.0	3.095537	0.9968596724897811	0.9968596724897811
Parallel	2	1	1.229075	1.8531795	3.082255	1.0011553229697088	0.5005776614848544
Parallel	2	2	1.228714	1.9479135	3.176627	0.9714127595087494	0.4857063797543747
Parallel	2	4	1.229393	1.8656139999999999	3.095007	0.997030378283474	0.498515189141737
Parallel	2	8	1.2298225	1.8503345	3.080157	1.0018372440106138	0.5009186220053069
Parallel	4	1	0.8727885	6.2820225	7.1548110000000005	0.4312924548251519	0.10782311370628797
Parallel	4	2	0.6274660000000001	4.9531155	5.580582	0.5529559461719226	0.13823898654298064
Parallel	4	4	0.6245455	4.8829815	5.5075275	0.5602906204281322	0.14007265510703304
Parallel	4	8	0.6212105	4.819912	5.4411225000000005	0.5671285658428017	0.14178214146070042

Tabella 41: Measurement with O0 optimization

6.2 Kosaraju

6.2.1 500 vertices - 25% of density

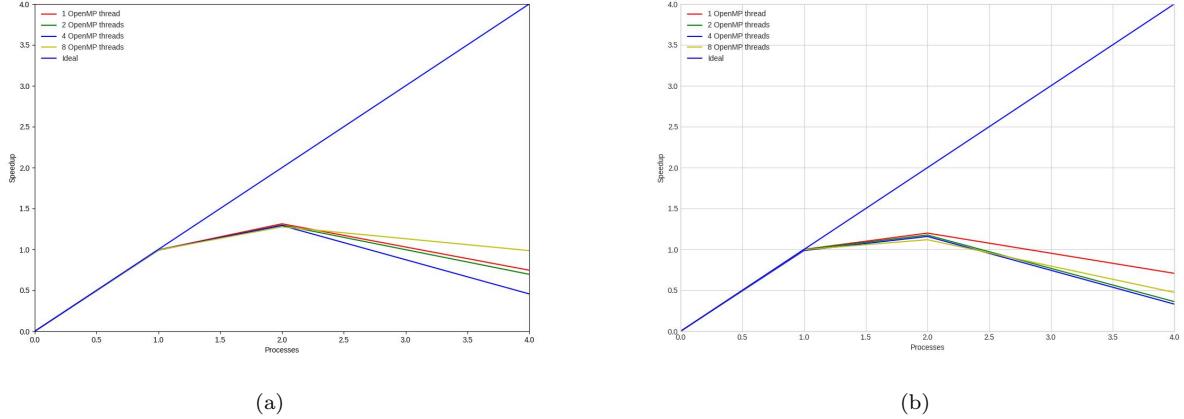
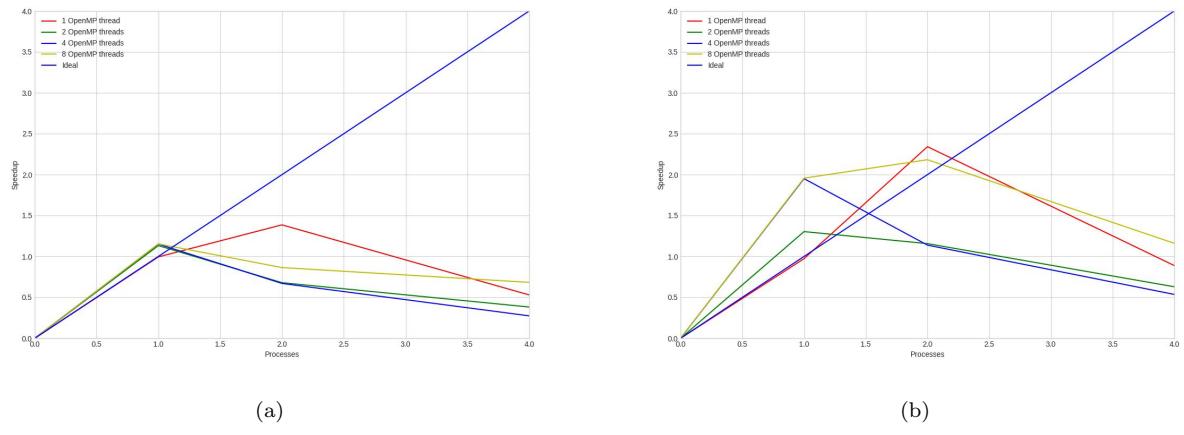


Figura 85: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization



caption(a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Kosaraju	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.05366450000000004	0.0	0.05366450000000004	1.0	1.0
Parallel	1	1	0.055174	0.0	0.055174	0.9726410990684018	0.9726410990684018
Parallel	1	2	0.0411665	0.0	0.0411665	1.3035963708355094	1.3035963708355094
Parallel	1	4	0.02751250000000002	0.0	0.02751250000000002	1.9505497501135847	1.9505497501135847
Parallel	1	8	0.0274215	0.0	0.0274215	1.9570227741006145	1.9570227741006145
Parallel	2	1	0.0139325	0.008985	0.0229175	2.3416384858732413	1.1708192429366207
Parallel	2	2	0.02756100000000002	0.01883549999999998	0.0463965	1.1566497472869721	0.5783248736434861
Parallel	2	4	0.0276155	0.019523	0.0471385	1.1384430985288034	0.5692215492644017
Parallel	2	8	0.0138595	0.010735	0.02459399999999998	2.1820159388468734	1.0910079694234367
Parallel	4	1	0.0182365	0.0421695	0.060406	0.8883968479952323	0.2220992119980808
Parallel	4	2	0.0183415	0.0669215	0.085263	0.6293996223449797	0.15734990558624493
Parallel	4	4	0.0181445	0.0820345000000001	0.1001789999999999	0.5356861218419031	0.13392153046047578
Parallel	4	8	0.0135335	0.0327075	0.0462410000000004	1.160539348197487	0.29013483704937176

Tabella 42: Measurement with O3 optimization

6.2.2 500 vertices - 50% of density

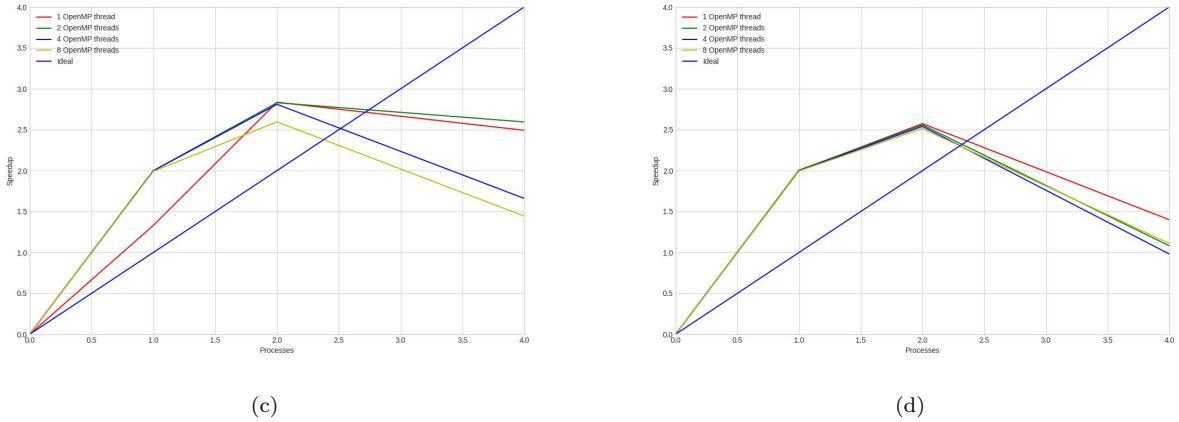


Figura 86: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

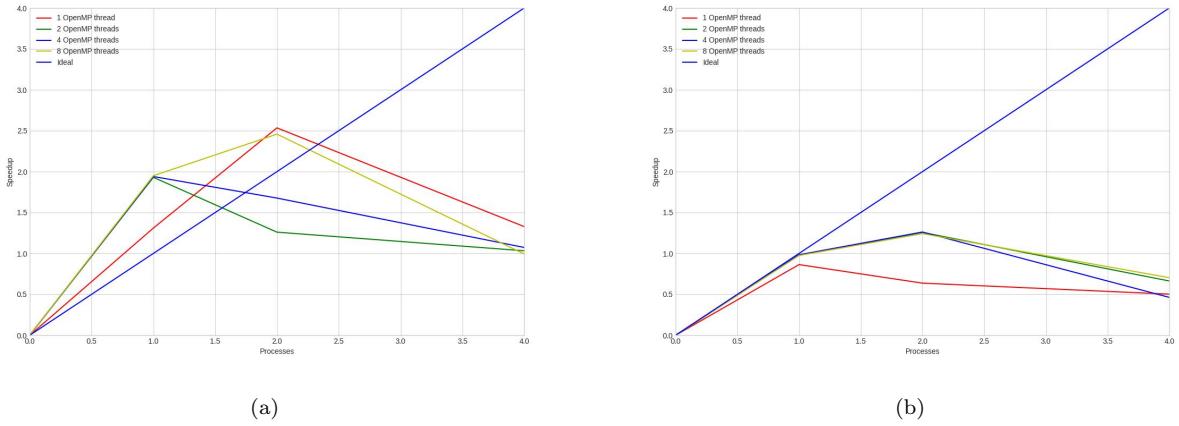


Figura 87: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Kosaraju	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.471223	0.0	0.471223	1.0	1.0
Parallel	1	1	0.3541485	0.0	0.3541485	1.3305802509399307	1.3305802509399307
Parallel	1	2	0.235648	0.0	0.235648	1.9996902159152634	1.9996902159152634
Parallel	1	4	0.2357785	0.0	0.2357785	1.9985834162147949	1.9985834162147949
Parallel	1	8	0.236559	0.0	0.236559	1.9919893134482307	1.9919893134482307
Parallel	2	1	0.123045	0.04308049999999994	0.16612500000000002	2.8365568096313014	1.4182784048156507
Parallel	2	2	0.1228259999999999	0.043676	0.1665025	2.830125673788682	1.415062836894341
Parallel	2	4	0.1234895	0.044186	0.1676754999999998	2.810327090123483	1.4051635450617415
Parallel	2	8	0.131428	0.050095	0.1815235	2.595933859803276	1.297966929901638
Parallel	4	1	0.0762455	0.1127954999999999	0.18904100000000001	2.492702641225977	0.6231756603064943
Parallel	4	2	0.0762235	0.1053574999999999	0.181581	2.595111823373591	0.6487779558433977
Parallel	4	4	0.11389300000000001	0.169828	0.2837215	1.6608646154767968	0.4152161538691992
Parallel	4	8	0.1516865	0.174523	0.32621	1.4445387940283867	0.3611346985070967

Tabella 43: Measurement with O0 optimization

6.2.3 500 vertices - 75% of density

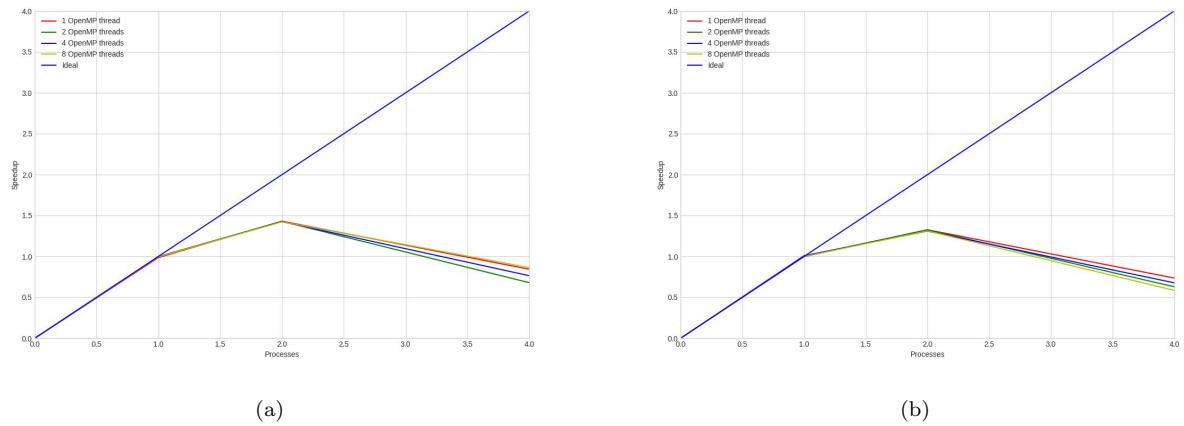


Figura 88: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

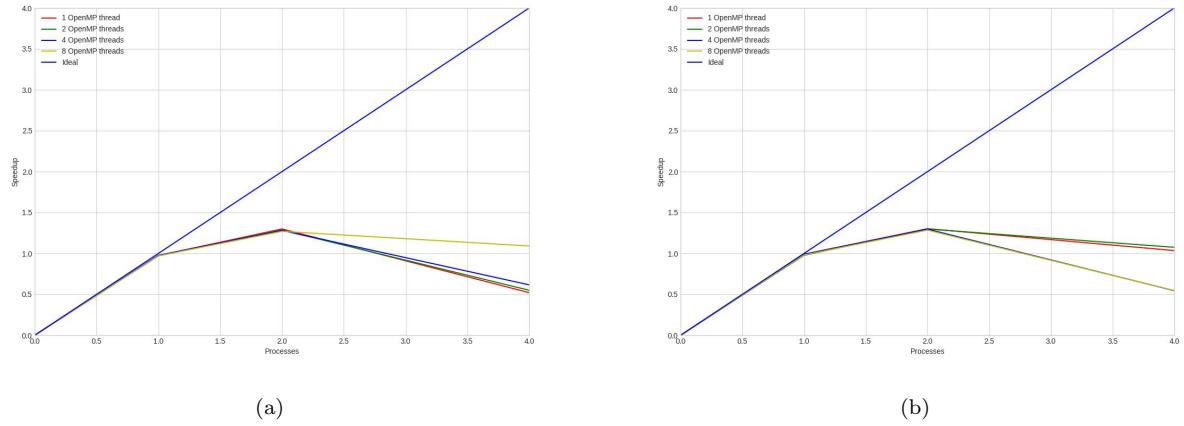


Figura 89: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Kosaraju	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.151996	0.0	0.151996	1.0	1.0
Parallel	1	1	0.152389	0.0	0.152389	0.9974210736995452	0.9974210736995452
Parallel	1	2	0.151681	0.0	0.151681	1.0020767268148283	1.0020767268148283
Parallel	1	4	0.150752	0.0	0.150752	1.008251963489705	1.008251963489705
Parallel	1	8	0.1521115	0.0	0.1521115	0.9992406885738421	0.9992406885738421
Parallel	2	1	0.07499349999999999	0.0396055	0.1145985	1.3263349869326386	0.6631674934663193
Parallel	2	2	0.074418	0.04029349999999996	0.1147115	1.3250284409148168	0.6625142204574084
Parallel	2	4	0.07475499999999999	0.041286	0.116041	1.3098473815289422	0.6549236907644711
Parallel	2	8	0.07450999999999999	0.0413035	0.1158135	1.312420400039719	0.6562102000198595
Parallel	4	1	0.064689	0.1418444999999998	0.2065329999999997	0.7359405034546538	0.18398512586366345
Parallel	4	2	0.0659215	0.175603	0.24152500000000002	0.6293178759962736	0.157329468990684
Parallel	4	4	0.065265	0.1591	0.2243655	0.6774481816500308	0.1693620454125077
Parallel	4	8	0.085883	0.174058	0.2599405	0.584733814084377	0.14618345352109424

Tabella 44: Measurement with $O1$ optimization

6.2.4 1000 vertices - 25% of density

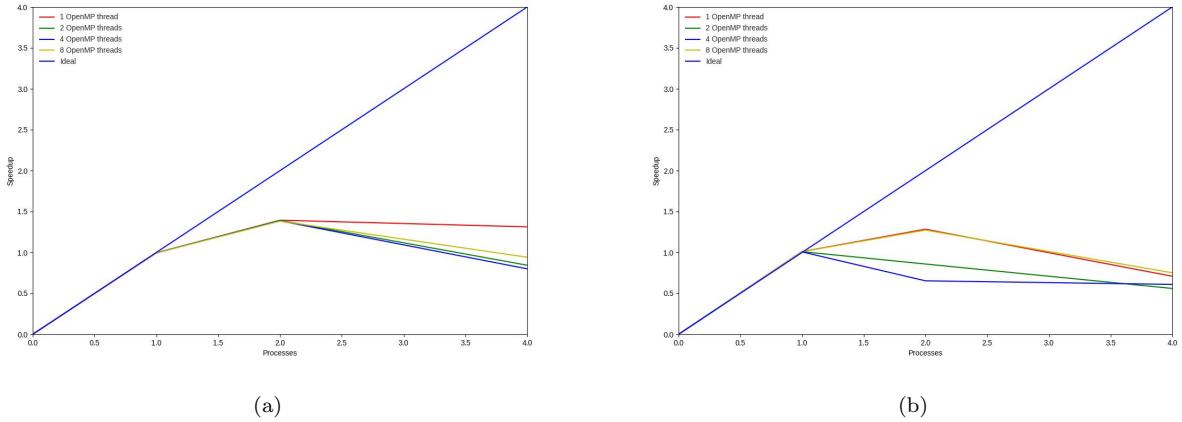


Figura 90: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

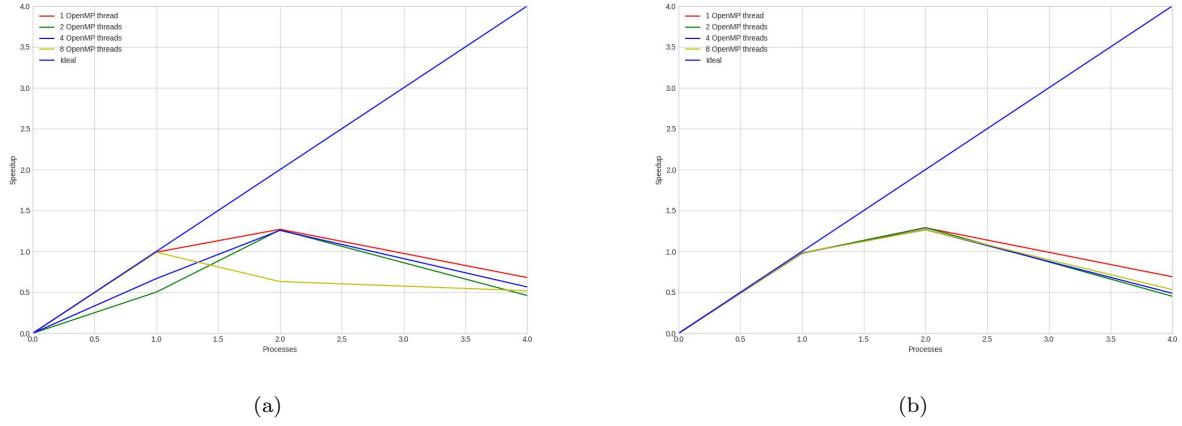


Figura 91: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Kosaraju	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.1289945	0.0	0.1289945	1.0	1.0
Parallel	1	1	0.1276515	0.0	0.1276515	1.0105208321092978	1.0105208321092978
Parallel	1	2	0.12778099999999998	0.0	0.12778099999999998	1.0094967170393097	1.0094967170393097
Parallel	1	4	0.1281655	0.0	0.1281655	1.00646819932041	1.00646819932041
Parallel	1	8	0.1277915	0.0	0.1277915	1.0094137716514793	1.0094137716514793
Parallel	2	1	0.0637290000000001	0.0367405	0.10047	1.2839106200855979	0.6419553100427989
Parallel	2	2	0.0943875	0.0558795	0.1502665	0.8584381748426962	0.4292190874213481
Parallel	2	4	0.1256485	0.071918	0.197566	0.6529185183685453	0.32645925918427265
Parallel	2	8	0.0639415	0.0374	0.1013419999999999	1.272863176175722	0.636431588087861
Parallel	4	1	0.0601815	0.1217439999999999	0.181926	0.7090492837747217	0.1772623209436804
Parallel	4	2	0.0802545	0.15014850000000002	0.2304025	0.5598658868718873	0.13996647171797183
Parallel	4	4	0.0698685	0.1418319999999999	0.2117005	0.6093254385322662	0.15233135963306654
Parallel	4	8	0.0661435	0.1055565	0.1717005	0.7512762047868237	0.18781905119670592

Tabella 45: Measurement with $O1$ optimization

6.2.5 1000 vertices - 50% of density

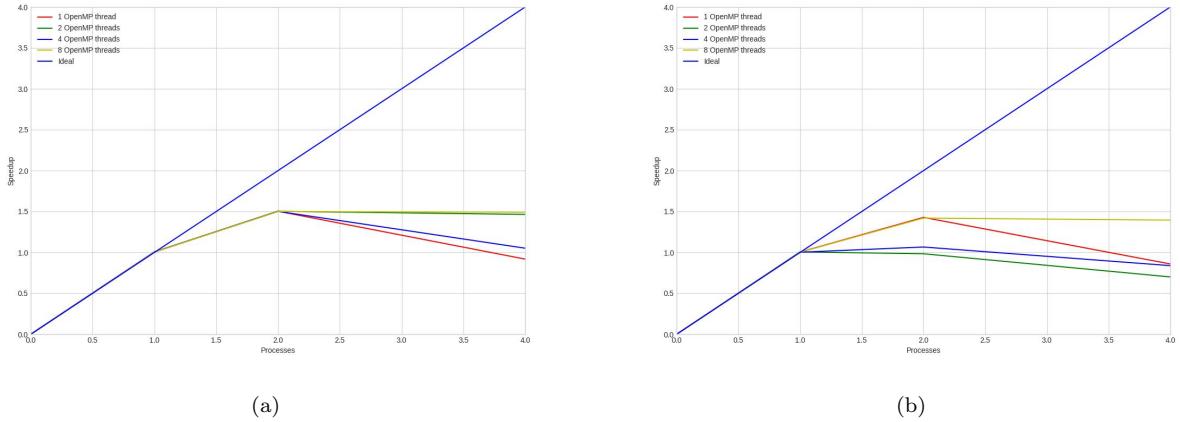


Figura 92: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

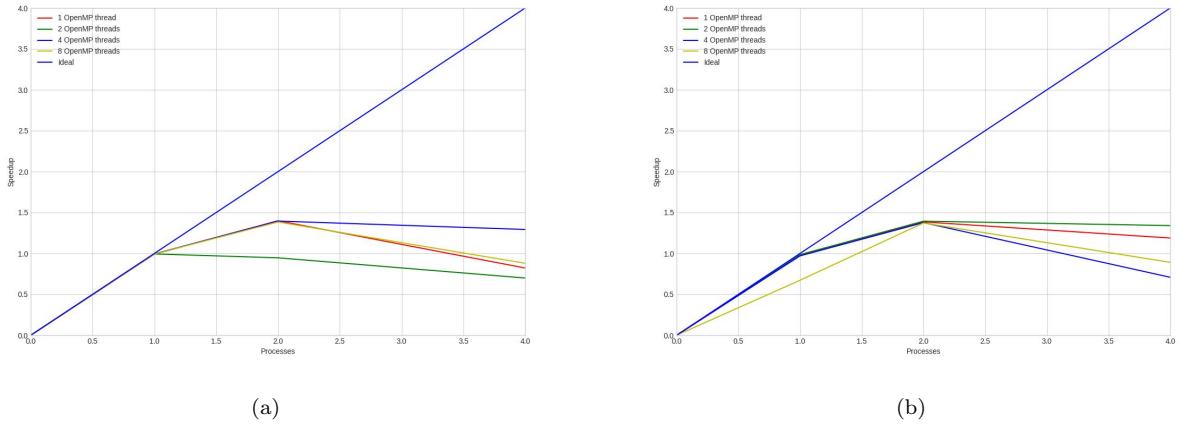


Figura 93: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Kosaraju	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	1.178648	0.0	1.178648	1.0	1.0
Parallel	1	1	1.1710034999999999	0.0	1.1710034999999999	1.0065281615298332	1.0065281615298332
Parallel	1	2	1.172132	0.0	1.172132	1.0055591008521225	1.0055591008521225
Parallel	1	4	1.1711385	0.0	1.1711385	1.0064121365662557	1.0064121365662557
Parallel	1	8	1.172515	0.0	1.172515	1.0052306367082724	1.0052306367082724
Parallel	2	1	0.6073105000000001	0.1765045	0.7838145000000001	1.5037333450708041	0.7518666725354021
Parallel	2	2	0.6066925	0.1774705	0.7841629999999999	1.5030650515262771	0.7515325257631386
Parallel	2	4	0.6069125	0.1776935	0.784606	1.5022163990588906	0.7511081995294453
Parallel	2	8	0.6070655	0.1782845	0.7853495	1.500794232376795	0.7503971161883976
Parallel	4	1	0.552286	0.731128	1.2834135	0.9183696447014154	0.22959241117535384
Parallel	4	2	0.381012	0.4236565	0.8046679999999999	1.464763107269085	0.36619077681727125
Parallel	4	4	0.555158	0.5655195	1.1206775	1.0517280841276817	0.2629320210319204
Parallel	4	8	0.369267	0.4210115	0.790278	1.4914346597020287	0.3728586649255072

Tabella 46: Measurement with O0 optimization

6.2.6 1000 vertices - 75% of density

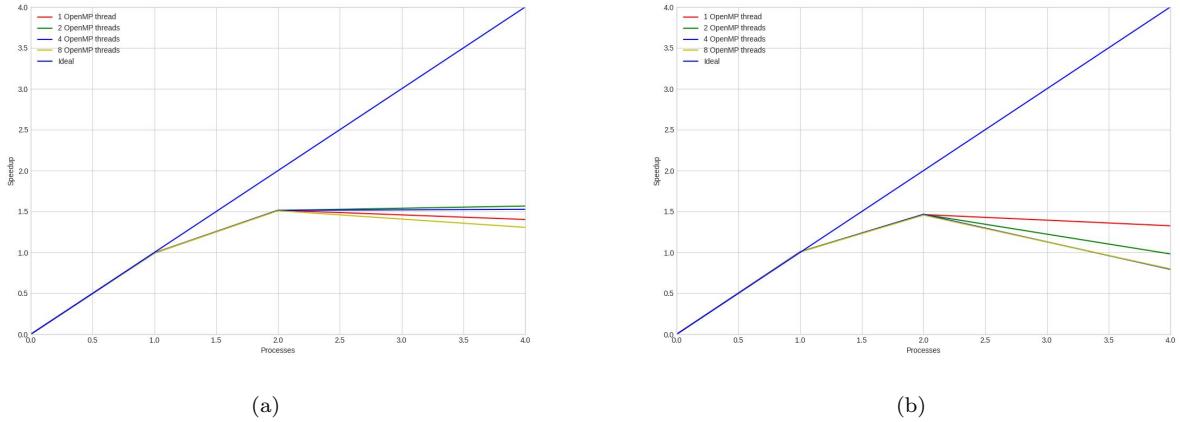


Figura 94: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

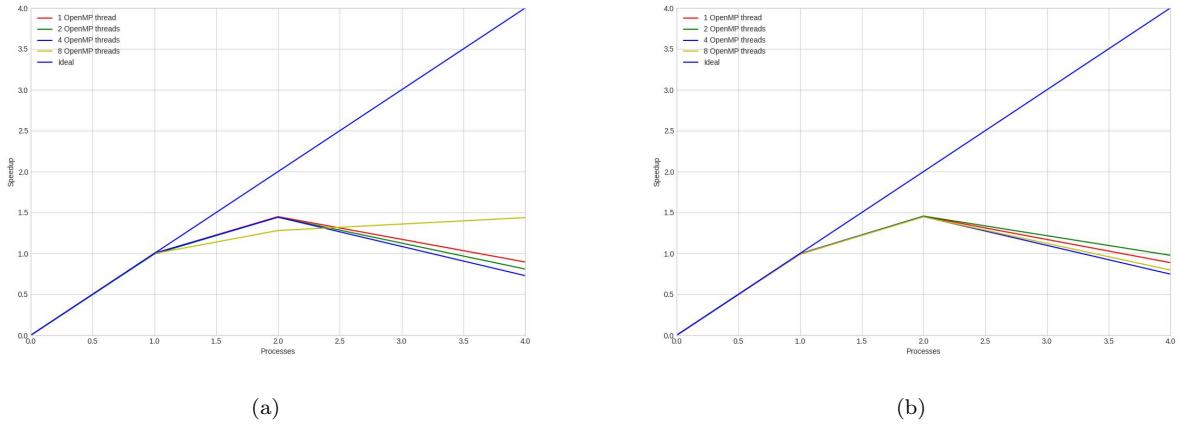


Figura 95: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Kosaraju	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	0.916936	0.0	0.916936	1.0	1.0
Parallel	1	1	0.9102635	0.0	0.9102635	1.0073302950189698	1.0073302950189698
Parallel	1	2	0.9110935	0.0	0.9110935	1.0064126239513287	1.0064126239513287
Parallel	1	4	0.9105989999999999	0.0	0.9105989999999999	1.0069591554570123	1.0069591554570123
Parallel	1	8	0.911638	0.0	0.911638	1.0058115172908546	1.0058115172908546
Parallel	2	1	0.463241	0.163566	0.6268075	1.46286698866877	0.731433494334385
Parallel	2	2	0.4628645	0.1628425	0.625707	1.4654398943914644	0.7327199471957322
Parallel	2	4	0.462677	0.1630565	0.6257334999999999	1.4653778325756892	0.7326889162878446
Parallel	2	8	0.4664355	0.1632974999999998	0.6297335	1.4560699089376699	0.7280349544688349
Parallel	4	1	0.2868435000000003	0.4046619999999997	0.691505	1.3260005350648223	0.3315001337662056
Parallel	4	2	0.3973905	0.5364445	0.9338345	0.9819041810941874	0.24547604527354686
Parallel	4	4	0.4989765000000004	0.6567805	1.155757	0.793364002986787	0.19834100074669675
Parallel	4	8	0.4876810000000003	0.6594599999999999	1.1471405	0.7993231866541196	0.1998307966635299

Tabella 47: Measurement with O1 optimization

6.2.7 2500 vertices - 25% of density

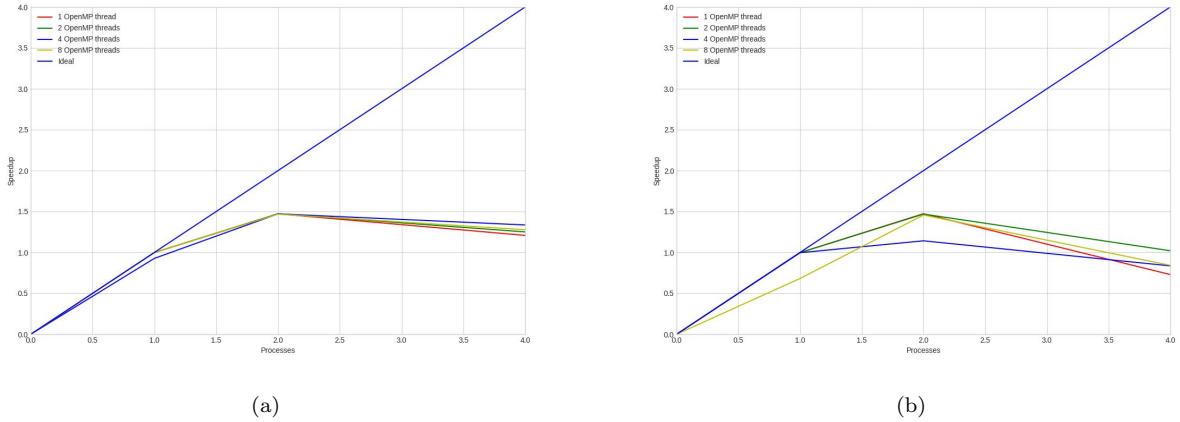


Figura 96: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

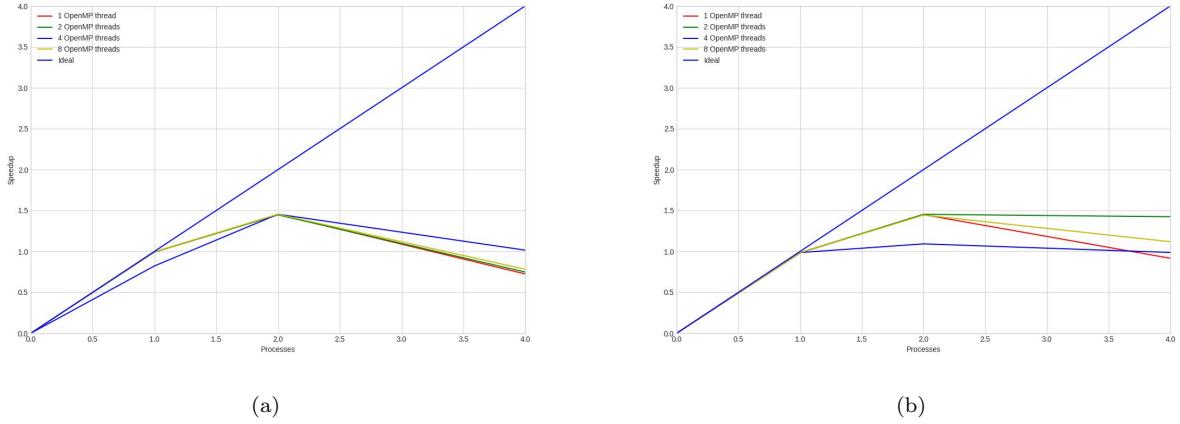


Figura 97: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Kosaraju	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	1.262461	0.0	1.262461	1.0	1.0
Parallel	1	1	1.266	0.0	1.266	0.9972045813586098	0.9972045813586098
Parallel	1	2	1.265112	0.0	1.265112	0.9979045333535688	0.9979045333535688
Parallel	1	4	1.266654	0.0	1.266654	0.9966897037391428	0.9966897037391428
Parallel	1	8	1.8521855	0.0	1.8521855	0.681606135022653	0.681606135022653
Parallel	2	1	0.6389245	0.217843	0.8567675	1.473516444076135	0.7367582220380675
Parallel	2	2	0.6406229999999999	0.2188435	0.8594665	1.4688891306409266	0.7344445653204633
Parallel	2	4	0.8784595	0.22688950000000002	1.105349	1.1421379130030427	0.5710689565015213
Parallel	2	8	0.6418265000000001	0.2243399999999998	0.8661665000000001	1.4575269304458207	0.7287634652229104
Parallel	4	1	0.6949375	1.035975	1.730913	0.7293613254970066	0.18234033137425165
Parallel	4	2	0.563077	0.673924	1.2370014999999999	1.0205816241936652	0.2551454060484163
Parallel	4	4	0.7204385	0.7900445	1.510483	0.8357995422656197	0.20894988556640492
Parallel	4	8	0.7134395	0.7856515	1.4990915	0.8421507292917078	0.21053768232292694

Tabella 48: Measurement with $O1$ optimization

6.2.8 2500 vertices - 50% of density

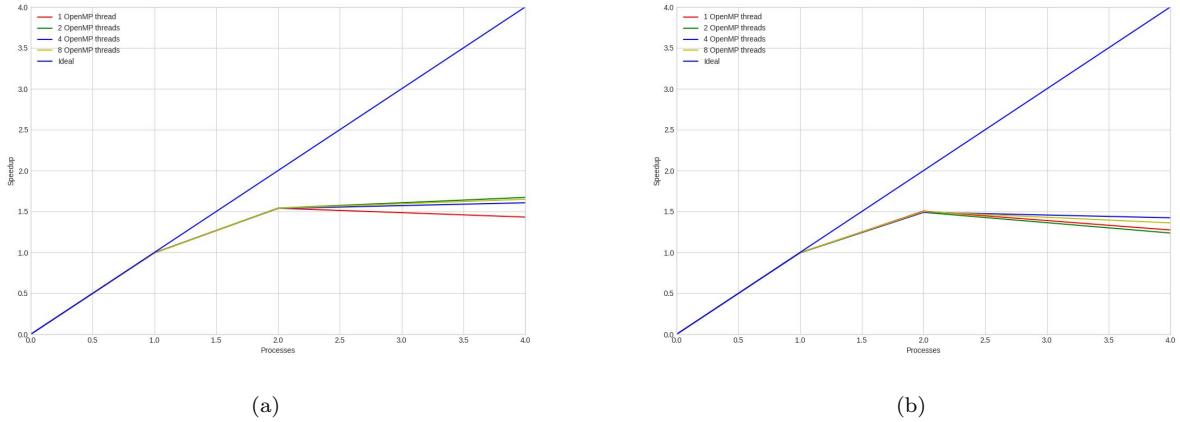


Figura 98: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

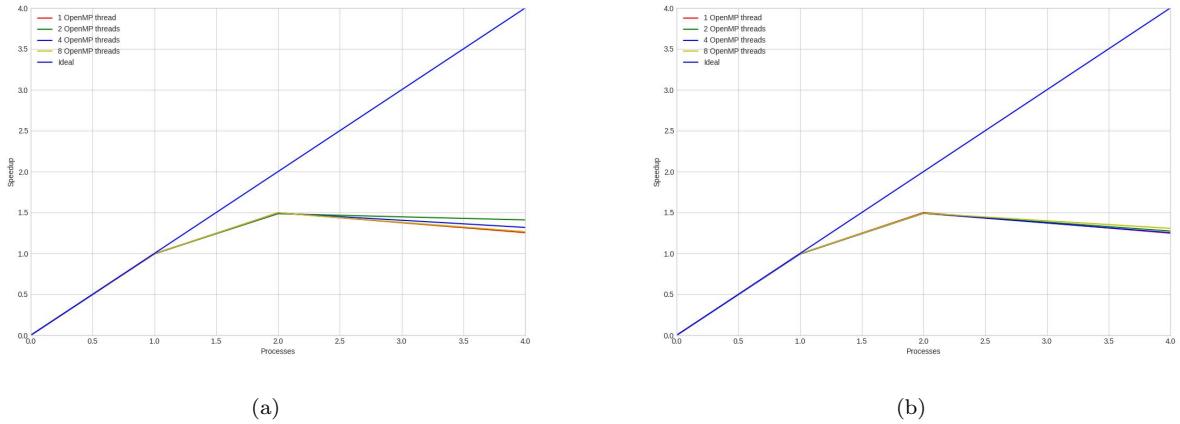


Figura 99: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Kosaraju	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	10.194613499999999	0.0	10.194613499999999	1.0	1.0
Parallel	1	1	10.258692	0.0	10.258692	0.9937537358563839	0.9937537358563839
Parallel	1	2	10.247829	0.0	10.247829	0.9948071440302136	0.9948071440302136
Parallel	1	4	10.2489585	0.0	10.2489585	0.9946975099957717	0.9946975099957717
Parallel	1	8	10.2670505	0.0	10.2670505	0.9929447118235173	0.9929447118235173
Parallel	2	1	5.4724725	1.146126	6.618598	1.5402980359284546	0.7701490179642273
Parallel	2	2	5.4775955	1.135443	6.6130385	1.541592945512112	0.770796472756056
Parallel	2	4	5.482314000000001	1.1479835	6.6302975	1.5375801010437313	0.7687900505218657
Parallel	2	8	5.478479	1.1426150000000002	6.6210945	1.5397172627576905	0.7698586313788452
Parallel	4	1	3.4221494999999997	3.6953110000000002	7.1174605	1.432338612908354	0.3580846532270885
Parallel	4	2	3.4129625	2.680319	6.0932815	1.6730908460408402	0.41827271151021006
Parallel	4	4	3.4194295	2.9244665000000003	6.343896	1.6069956853012721	0.40174892132531803
Parallel	4	8	3.4281699999999997	2.749977	6.178146	1.6501088676117397	0.4125272169029349

Tabella 49: Measurement with O0 optimization

6.2.9 2500 vertices - 75% of density

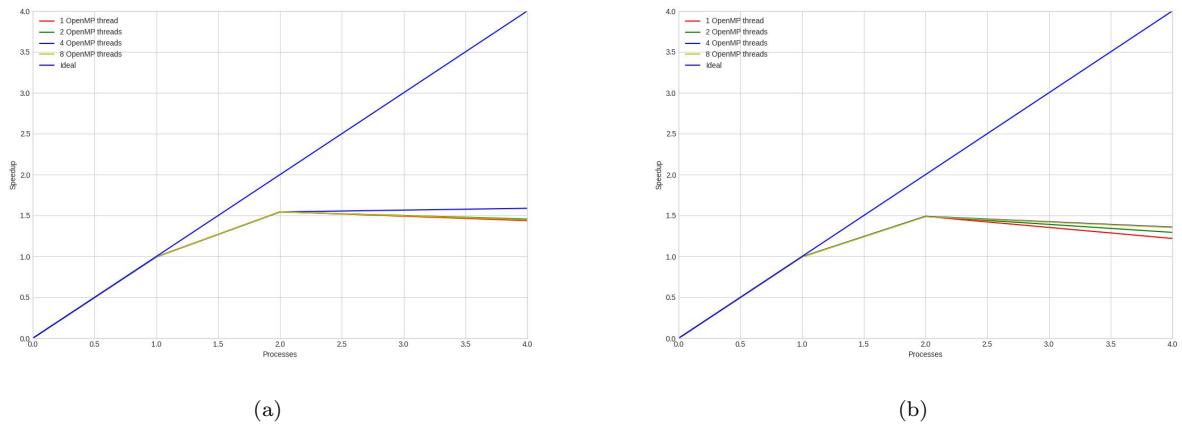


Figura 100: (a) Speedup with no optimization - (b) Speedup with $O1$ optimization

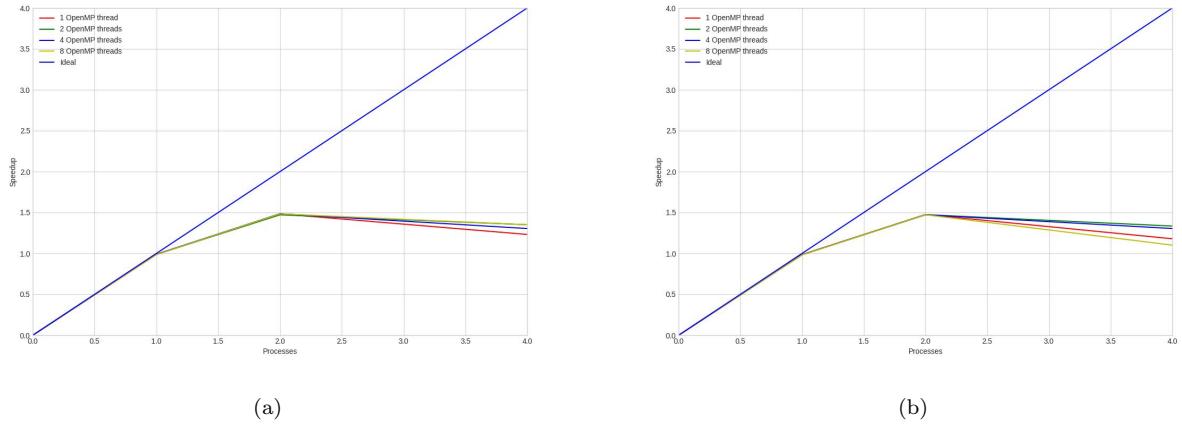


Figura 101: (a) Speedup with $O2$ optimization - (b) Speedup with $O3$ optimization

Version	Processes	OMP	Kosaraju	Communication	Elapsed	Speedup	Efficiency
Serial	1	0	18.0724475	0.0	18.0724475	1.0	1.0
Parallel	1	1	18.170600999999998	0.0	18.170600999999998	0.9945982249018621	0.9945982249018621
Parallel	1	2	18.165741	0.0	18.165741	0.9948643162973643	0.9948643162973643
Parallel	1	4	18.18484	0.0	18.18484	0.9938194397091202	0.9938194397091202
Parallel	1	8	18.171785999999997	0.0	18.171785999999997	0.9945333661754547	0.9945333661754547
Parallel	2	1	9.83179099999999	1.861803	11.6935945	1.5454997605740475	0.7727498802870237
Parallel	2	2	9.825023	1.8783055000000002	11.703328500000001	1.5442143233012726	0.7721071616506363
Parallel	2	4	9.834575000000001	1.869847	11.704422000000001	1.5440700531816092	0.7720350265908046
Parallel	2	8	9.8256115	1.868599	11.69421	1.545418416464216	0.772709208232108
Parallel	4	1	6.213698	6.3526495	12.5663475	1.4381623220271444	0.3595405805067861
Parallel	4	2	6.7782915	5.627601500000001	12.40589299999999	1.4567631286196003	0.3641907821549001
Parallel	4	4	6.5868285	4.7915735	11.37840149999999	1.5883116358655476	0.3970779089663869
Parallel	4	8	6.222115	6.2490665	12.471181000000001	1.4491368139071992	0.3622842034767998

Tabella 50: Measurement with O0 optimization

7 Final Considerations

In this section we will analyze what were the results following the various measurements made. In particular, as per trace, both MPI and OpenMP have been used in the development. As far as MPI is concerned, this has been used in order to divide the work among several processes with the aim of improving performance compared to the serial algorithm, thus exploiting data-parallelism. With regard to OpenMP, on the other hand, during the writing and testing of the code's performance, following careful analyses, some portions of the code have been identified where this has brought considerable benefits, but also others where the use of this approach does not it did nothing but slow down the execution. In particular, for the two algorithms that compute the strictly connected components used (Tarjan and Kosaraju) the use of OpenMP was not advantageous given their recursive nature. The problem here is that both algorithms perform one or more dfs and the entire body of the dfs is one critical section, which means that even if there are 1000 recursive calls in parallel, they will execute one after another sequentially and not in parallel. It will even be slower than the sequential version because of the constant cache invalidation and the added OpenMP overhead. Conversely, advantages have been obtained by parallelizing the code snippet for the creation of the new graph by a process following receipt of the necessary information from the previous process, as well as the functions for serializing the data structures before being transmitted to the other processes.

7.1 Tarjan Implementation

Considering the results previously obtained, when it comes to parallelizing Tarjan's algorithm, the number of processes used plays a crucial role in the overall performance of the algorithm. As the number of processes increases, the overall execution time of the algorithm decreases. However, as the number of processes grows up, the communication overhead also increases. This is because each process has to communicate with the other processes to share information about his graph and the strongly connected components found.

This phenomenon can be observed particularly when the number of vertices and the density of the graph increase. In sparse graphs, the algorithm would have to spend most of its time waiting for the next non-trivial component to be processed, which could lead to increased overhead and synchronization costs when using MPI. Additionally, if the graph is very sparse, the overhead of parallelism and communication might be higher than the overhead of the serial version, leading to slower performance.

In particular, the best performance is achieved when using two processes. This is because the communication overhead is minimized while still achieving a significant speedup. Consequently, as the number of nodes in the graph increases, better and better performance will be obtained using two MPI processes.

However, when the graph has a density of 0%, the parallel version of the algorithm may not be as efficient as the serial version. This is because, in this type of graph, the number of strongly connected components is equal to the number of vertices, and there are no macro-components to be found. This means that the graph never decreases in size and the last iteration of the parallel algorithm will perform Tarjan on the exact same graph as the initial one, leading to no speedup compared to the serial version.

Let's consider an example of a graph with 100 vertices and a density of 0%. In this case, there are no edges between the vertices, meaning that each vertex is a separate strongly connected component.

- **Serial version:** The algorithm will start from one vertex and perform the depth-first search to find the strongly connected components. Since there are no edges between the vertices, the algorithm will find that each vertex is a separate strongly connected component. The total execution time of the algorithm will be the time it takes to perform the Tarjan's algorithm.
- **Parallel version with 4 processes:** The graph is divided among the 4 processes and each process performs the Tarjan's algorithm on its part of the graph. However, since the graph is sparse and there are no edges between the vertices, each process will find that each vertex of its part of the graph is a separate strongly connected component. At this point, when the processes will exchange the various information, the remaining processes will run Tarjan on a graph that never tends to get smaller, up to the last iteration, in which the last process will find itself at run Tarjan on the starting graph. Therefore, in this case, the execution time will be greater since the time required for the last process to execute Tarjan on the final graph (which is therefore

equivalent to the initial one) will be added to the Tarjan execution times of the other processes up to that moment.

7.2 Kosaraju Implementation

Same as Tarjan's algorithm, when it comes to parallelizing Kosaraju's algorithm, the number of processes used plays a crucial role in the overall performance of the algorithm. Kosaraju's algorithm benefits a lot from parallelism and it scales well with the number of processes used. As the number of processes increases, the overall execution time of the algorithm decreases significantly.

The communication overhead does increase as the number of processes increases, but the speedup that the algorithm achieves by using more processes compensates for this overhead. This is not the case with Tarjan's algorithm, where the parallel version does not have the same scaling as the serial version, and the communication overhead is not compensated in the same way.

In contrast to Tarjan, as the number of nodes increases in Kosaraju, performance is also observed to improve with an increase in the number of processes.

Regarding optimization levels, as the algorithm is already sufficiently optimized, no significant benefits are observed even when compiled with different levels of optimization.

7.3 Cluster execution

As already discussed in chapter 6, in addition to the measurements taken on our machine, the same measurements were taken on a Raspberry Pi cluster, where the front-end node is represented by a Raspberry Pi 3, while the worker nodes are Raspberry Pi 4s.

Initially, we were unaware that the sequential program was run on the cluster frontend, which is a Raspberry Pi 3. In contrast, the parallel algorithm was executed on the cluster workers, which are Raspberry Pi 4. This is the reason why in some measurements taken on the cluster, an abnormal behavior is observed in which some speedup curves are above the ideal line, as the sequential algorithm was run on an older device. Upon examining a helpful table, we discovered that the Raspberry Pi 3 contains a Broadcom BCM2837B0 quad core Cortex A53 processor with a clocking speed of 1.4GHz whereas on the other hand the Raspberry Pi 4 model has a Broadcom BCM2711 quad core Cortex-A72 processor with the clock speeds at 1.5GHz. Therefore, Raspberry Pi 4 has 10% faster processing performance in comparison to Raspberry Pi 3.

Despite the presence of this inconsistency, since the time available for using the cluster had expired, it was not possible to repeat the measurements in a uniform manner.

8 Test Case

The code provided in the *test* folder is a test program that compares the correctness of two different algorithms for finding strongly connected components (SCCs) in a directed graph. Six test cases have been carried out, which consider two different levels of number of vertices (100 and 1000), each with an increasing level of edges.

The program first includes the "TestUtil.h" header file, which likely contains utility functions used in the testing process. Then, the program defines some macro constants, such as the number of vertices in the test graph, and the lower and upper bounds for the edges.

The main function first allocates memory for a *Results* struct and uses the *generateGraph* function to create a graph with the specified number of vertices, lower and upper bounds, and writes it to a file. Next, the program runs the Tarjan algorithm in both sequential and parallel versions using the "system" function, which allows for

execution of command-line commands. The program then reads the results of the sequential and parallel versions of the algorithm and stores them in the Results struct.

The program then calls the compareResults function to compare the results of the two versions of the algorithm. The function begins by asserting that the number of SCCs found by the sequential and parallel versions of the algorithm is the same. This ensures that both versions of the algorithm have found the same number of SCCs in the graph, which is an important condition for the validity of the comparison. Then the function iterates over the SCCs found by the sequential version of the algorithm and for each SCC it asserts that the length of the SCC found by the parallel version is the same. It also iterates over the items of each SCC, and for each item, it asserts that the item found by the parallel version is the same as the item found by the sequential version.

This comparison ensures that the SCCs found by the parallel version of the algorithm are the same as the SCCs found by the sequential version of the algorithm. This is important because it shows that the parallel version of the algorithm is producing the same results as the sequential version, which means that it is working correctly. After that, the program frees the memory allocated for the Results struct, removes the output files, and repeats the process for the Kosaraju algorithm.

9 Api Documentation

9.1 include/tarjan.h File Reference

```
#include <stdlib.h>
#include <stdio.h>
#include <omp.h>

#include "../include/Tarjan.h"

#define NOT_INITIALIZED -1
#define PARAMETERS_NUM 3
```

Functions:

- void sccUtil(TGraph, int, THTAuxiliarySCC*, TStack*, int*, int, int, int*, TArray*);
- TArray* scc(TGraph, int, int, int*);

Function documentation:

```
void sccUtil(TGraph graph, int u, THTAuxiliarySCC *auxiliaryHT, TStack *st, int *time, int start, int stop, int *sccCount, TArray* sccs)
```

This is a helper function for the SCC() function. It performs the recursive DFS traversal of the graph to find the SCCs.

Parameters:

graph	Pointer to the graph structure
u	The current node being visited
auxiliaryHT	Pointer to a hash table that stores auxiliary information for each node
st	Pointer to a stack used to track nodes during the SCC search
time	Pointer to a variable that tracks the current time
start	The index of the first node in the subgraph being searched
stop	The index of the last node in the subgraph being searched
sccCount	Pointer to a variable that tracks the number of SCCs found
sccs	Pointer to an array of nodes representing the SCCs found

Function documentation:

```
TArray* scc(TGraph graph, int start, int stop, int *sccCount)
```

This function finds the strongly connected components (SCCs) in a graph using the Tarjan's algorithm. It returns an array of nodes, where each node represents a SCC and its children represent the nodes in the SCC.

Parameters:

graph	a graph data structure represented as an adjacency list
start	the index of the first vertex in the graph
stop	the index of the last vertex in the graph
sccCount	a pointer to a variable that stores the number of SCCs found

Returns:

an array of nodes representing the SCCs found in the graph.

9.2 include/kosaraju.h File Reference

```
#include <stdbool.h>

#include "TGraph.h"
#include "TArray.h"
#include "THTAuxiliarySCC.h"
#include "Utils.h"

#define NOT_INITIALIZED -1
#define DELTA 10
```

Functions:

- TGraph transposeGraph(TGraph*, int, int);
- void dfs1(int, TGraph*, THTAuxiliarySCC*, TStack*, int, int);
- void dfs2(int, TGraph*, THTAuxiliarySCC*, int, int, int *, TArray *);
- TArray* kosaraju(TGraph , int , int , int *);

Function documentation:

```
TGraph transposeGraph(TGraph* graph, int start, int stop)
```

This function finds the strongly connected components (SCCs) in a graph using the Tarjan's algorithm. It returns an array of nodes, where each node represents a SCC and its children represent the nodes in the SCC.

Parameters:

graph	The original graph
start	The start index of the cut
stop	The stop index of the cut
sccCount	a pointer to a variable that stores the number of SCCs found

Returns:

The transposed graph.

Function documentation:

```
void dfs1(int node, TGraph* graph, THTAuxiliarySCC* auxiliaryHT, TStack* stack, int start, int stop)
```

This function finds the strongly connected components (SCCs) in a graph using the Tarjan's algorithm. It returns an array of nodes, where each node represents a SCC and its children represent the nodes in the SCC.

Parameters:

node	The starting node of the search
graph	The graph to be searched
auxiliaryHT	A hash table used for auxiliary data
stack	A stack used for storing visited nodes
start	The start of the cut range
stop	The stop of the cut range

Function documentation:

```
void dfs2(int node, TGraph* transpose, THTAuxiliarySCC* auxiliaryHT, int start,
int stop, int *sccCount, TArray *sccs)
```

This function finds the strongly connected components (SCCs) in a graph using the Tarjan's algorithm. It returns an array of nodes, where each node represents a SCC and its children represent the nodes in the SCC.

Parameters:

node	The starting node of the dfs
transpose	A pointer to the transposed graph on which to perform the DFS
auxiliaryHT	A pointer to the auxiliary hash table used to store information of the nodes
start	The lower bound of the range in which to search for SCCs
stop	The upper bound of the range in which to search for SCCs
sccCount	A pointer to the variable that keeps track of the number of SCCs found so far
sccs	A pointer to the array that stores the SCCs found

Function documentation:

```
TArray* kosaraju(TGraph graph, int start, int stop, int *sccCount)
```

The code is an implementation of the Kosaraju's algorithm to find the strongly connected components (SCCs) in a given graph.

Parameters: Returns:

graph	The graph on which to perform the Kosaraju algorithm
start	The lower bound of the range in which to search for SCCs
stop	The upper bound of the range in which to search for SCCs
sccCount	A pointer to the variable that keeps track of the number of SCCs found so far

An array of SCCs.

9.3 include/Utils.h File Reference

```
#include <stdio.h>

#include "TGraph.h"
#include "THTSCCs.h"
#include "THTAuxiliaryGraph.h"
```

Functions:

- int calculateId(int, int, int, int*);
- bool isInCut(TGraph*, int, int, int);
- int min(int, int);
- void createNewGraph(TGraph *, TGraph*, TGraph*, TGraph*, THTSCCs*,
THTAuxiliaryGraph*, int*, int, int);
- int* serializeSCCs(TArray*, int, int*);
- void deserializeSCCs(TArray*, int*, int, int*);
- int* serializeSCCsHT(THTSCCs*, int*);
- void deserializeSCCsHT(THTSCCs*, int*);
- int *serializeAuxiliaryGraphHT(THTAuxiliaryGraph*, int*);
- void deserializeAuxiliaryGraphHT(THTAuxiliaryGraph*, int*);
- int* serializeGraph(TGraph*, int*);
- TGraph deserializeGraph(int*, int);
- TArray* getAdjacencyList(int, TGraph*, TGraph*, TGraph*, int);

Function documentation:

```
int calculateId(int rank, int numItem, int vertex, int* offset)
```

This function calculates an id for a given vertex, rank, number of items, and offset.

Parameters:

rank	The rank of the item
numItem	The number of items
vertex	The vertex of the item
offset	Pointer to the current offset value

Returns:

The calculated ID

Function documentation:

```
bool isInCut(TGraph* graph, int vertex, int start, int stop)
```

This function checks if a given vertex is in a cut of a given graph.

Parameters:

graph	Pointer to the graph data structure
vertex	The vertex to check for
start	The starting index of the cut
stop	The ending index of the cut

Returns:

True if the vertex is within the cut, false otherwise.

Function documentation:

```
int minParallel(int a, int b)
```

This function is used to compare 2 int values and after return the minimum value between.

Parameters:

a	the first int value
b	the second int value

Returns:

The minimum of a and b

Function documentation:

```
void createNewGraph(TGraph* graph, TGraph* oldGraph, TGraph* newGraph, TGraph* rvdGraph,
THTSCCs* sccsHT, THTAuxiliaryGraph* auxiliaryGraphHT, int* nodes, int sccCount, int numIteration)
```

The function uses OpenMP to parallelize the creation of the new graph.

Parameters:

graph	Pointer to the original graph
oldGraph	Pointer to the previous iteration graph
newGraph	Pointer to the new graph to be created
rvdGraph	Pointer to the received graph
sccsHT	Pointer to the hash table for strongly connected components
auxiliaryGraphHT	Pointer to the hash table for auxiliary graph
nodes	Array containing the ids of the new graph's vertices
sccCount	Number of strongly connected components
numIteration	The number of iteration

Function documentation:

```
int *serializeSCCs(TArray *sccs, int sccCount, int *size)
```

This function serializes an array of strongly connected components (sccs) into a one-dimensional array of integers (buf).

Parameters: Returns:

graph	Pointer to the graph data structure
vertex	The vertex to check for
start	The starting index of the cut
stop	The ending index of the cut

True if the vertex is within the cut, false otherwise.

Function documentation:

```
void deserializeSCCs(TArray *sccs, int *buf, int bufSize, int* sccCount)
```

This function deserializes an array of strongly connected components (SCCs) that were previously serialized.

Parameters:

sccs	Array of strongly connected components
buf	Buffer containing the serialized SCCs
bufSize	Size of the buffer
sccCount	Pointer to the number of strongly connected components

Function documentation:

```
int *serializeSCCsHT(THTSCCs* sccsHT, int *size)
```

This function serializes an hash table of strongly connected components (sccs) into a one-dimensional array of integers.

Parameters:

sccsHT	Pointer to a THTSCCs struct
size	Pointer to an integer which will be used to store the size of the serialized data

Returns:

Pointer to the serialized data

Function documentation:

```
void deserializeSCCsHT(THTSCCs* sccsHT, int *buf)
```

Deserialization function for a hash table of strongly connected components (THTSCCs).

Parameters:

sccsHT	Pointer to a THTSCCs struct
buf	Pointer to an integer array which contains the serialized data

Function documentation:

```
int *serializeAuxiliaryGraphHT(THTAuxiliaryGraph* auxGraphHT, int* size)
```

This function serializes a THTAuxiliaryGraph data structure into an array of integers.

Parameters: Returns:

auxGraphHT	Pointer to a THTAuxiliaryGraph struct
size	Pointer to an integer variable that stores the size of the serialized data array

Pointer to an integer array which contains the serialized data.

Function documentation:

```
void deserializeAuxiliaryGraphHT(THTAuxiliaryGraph* auxGraphHT, int* buf)
```

This function is used to deserialize an auxiliary graph hash table from a buffer of integers.

Parameters:

auxGraphHT	Pointer to a THTAuxiliaryGraph struct
buf	Pointer to an integer array which contains the serialized data

Function documentation:

```
int* serializeGraph(TGraph* graph, int* size)
```

This function serializes a graph into an array of integers.

Parameters:

graph	Pointer to a TGraph struct
size	Pointer to an integer variable which will store the size of the serialized data array

Returns:

Pointer to an integer array which contains the serialized data.

Function documentation:

```
int* serializeGraph(TGraph* graph, int* size)
```

This function deserializes a graph from an array of integers.

Parameters:

buf	Pointer to an integer array which contains the serialized data
bufSize	Size of the serialized data array

Returns:

Pointer to a TGraph struct.

Function documentation:

```
TArray* getAdjacencyList(int numIteration, TGraph* graph, TGraph* oldGraph, TGraph* rvdGraph, int vertex)
```

This function returns the adjacency list of a given vertex in a given graph.

Parameters:

numIteration	The number of iteration
graph	Pointer to the original graph
oldGraph	Pointer to the previous iteration graph
rvdGraph	Pointer to the received graph
vertex	The vertex

Returns:

The adjacency list of the vertex.

10 How To Run

Starting from the main directory of our project, we can compile and build the program with few and simple terminal commands:

1. Create a build directory and launch cmake

```
mkdir build  
cd build  
cmake ..
```

2. Generate executables with `make`
3. To generate measures (It can take a long time!) run `make generate_measures`
4. To extract mean times and speedup curves from them run `make extract_measures`

Results can be found in the *measures*, divided into *Tarjan* and *Kosaraju* directories with all the optimization version (O0, O1, O2, O3).

Finally, if it is necessary to remove all executable files from a binary program and coding directory there is a simple command: `make clean`